

Topic 7

Introduction to Multicore I

Xinfei Guo
xinfei.guo@sjtu.edu.cn

July 6th, 2022



T7 learning goals

- Multicore/Multiprocessing

- Section I


- Motivation
 - SMT vs. Multicore
 - Communication
 - Cache coherence protocol

- Section II


- Memory consistency
 - Parallel programming
 - Dark silicon
 - Roofline model

Spectrum of Parallelism


Bit-level




Pipelining




ILP



Multithreading
Multiprocessing



Distributed




Why Multiprocessing?

- A single core can only be so fast
 - Limited clock frequency
 - Limited instruction-level parallelism
- What if we need even more computing power?
 - Use multiple cores! But how?
- Old-school (2000s): Ultra Enterprise 25k
 - 72 dual-core UltraSPARC IV+ processors
 - Up to 1TB of memory
 - Niche: large database servers
 - \$\$\$, weighs more than 1 ton
- Today: multicore is everywhere
 - Can't buy a single-core smartphone...
 - ...or even smart watch!




Recall: Moore's Law



Dennard Scaling

- In 1974, [Robert Dennard](#) observed that power density remained **constant** for a given area of silicon (square nanometers) when the dimension of the transistor shrank, thanks to technology improvements ([Dennard Scaling](#)).
- Capacitance is related to area
 - So, as the size of the transistors shrunk, and the voltage was reduced, circuits could operate at higher frequencies at the same power




- $\text{Power} = \alpha * CFV^2$
 - ◆ Alpha – percent time switched
 - ◆ C = capacitance
 - ◆ F = frequency
 - ◆ V = voltage

End of Dennard Scaling

- Dennard scaling ignored the “leakage current” and “threshold voltage”, which establish a baseline of power per transistor
- As transistors get smaller, power density increases because these don’t scale with size
- These created a “Power Wall” that has limited practical processor frequency to around 4 GHz since 2006

End of Dennard Scaling



source: <https://silvanogai.github.io/posts/dennard/>

Power Wall

Power = $\alpha * \text{capacitance} * \text{voltage}^2 * \text{frequency}$

In practice: Power $\sim \text{voltage}^3$

Reducing voltage helps (a lot)


... so does reducing clock speed

Better cooling helps

The power wall

- We can't reduce voltage further (leakage)
- We can't remove more heat

Why Multicore?



Power Efficiency


Q: Does multiple issue / ILP cost much?

A: Yes. → Dynamic issue and speculation requires power

CPU	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
Core i5 Nehal	2010	3300MHz	14	4	Yes	1	87W
Core i5 Ivy Br	2012	3400MHz	14	4	Yes	8	77W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

Why Multicore?


- Difficult to make single-core clock frequencies even higher
- Deeply pipelined circuits:
 - heat problems
 - speed of light problems
 - difficult design and verification
 - large design teams necessary
 - server farms need expensive air-conditioning
- Many new applications are multithreaded
- General trend in computer architecture (shift towards more parallelism)



source: Intel

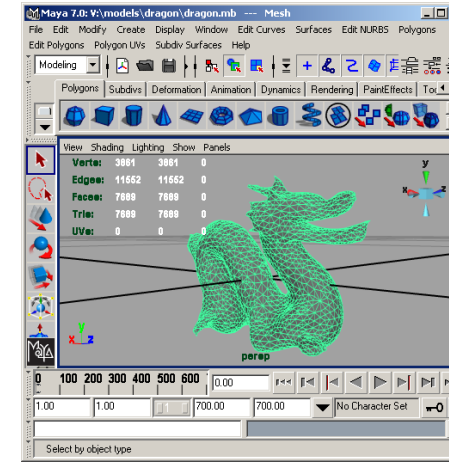
Multicore/Multiprocessing Overview

- A **multi-core** microprocessor is one that combines two or more independent processors into a single package, often a single integrated circuit (IC).
- A **dual-core** device contains two independent microprocessors.
- In general, multi-core microprocessors allow a computing device to exhibit some form of thread-level parallelism (**TLP**) without including multiple microprocessors in separate physical packages.
- Multi-core processors are MIMD: Different cores execute different threads (Multiple Instructions), operating on different parts of memory (Multiple Data).
- Multi-core is **usually** a shared memory multiprocessor:
All cores share the same memory



What applications benefit from multi-core?

- Database servers
- Web servers (Web commerce)
- Compilers
- Multimedia applications
- Scientific applications, CAD/CAM
- In general, applications with *Thread-level parallelism* (as opposed to instruction-level parallelism)



Each can
run on its
own core



More examples

- Editing a photo while recording a TV show through a digital video recorder
- Downloading software while running an anti-virus program
- “Anything that can be threaded today will map efficiently to multi-core”
- BUT: some applications difficult to parallelize

Multicore Design Considerations

- Cores in a multicore processor are connected together and can collaborate.
- Why not staying with multi-threading style?
- There are a number of challenges to consider when creating a system like this.
 - How do cores communicate with each other?
 - How is data synchronized?
 - How do we ensure that cores don't get stale data when it's been modified by other cores?
 - How do cores see the ordering of events coming from different cores?
 - How to program multicore?
- We'll explore each of these by considering the concepts of
 - Shared memory and message passing
 - Cache coherence
 - Memory consistency
 - Parallel programming

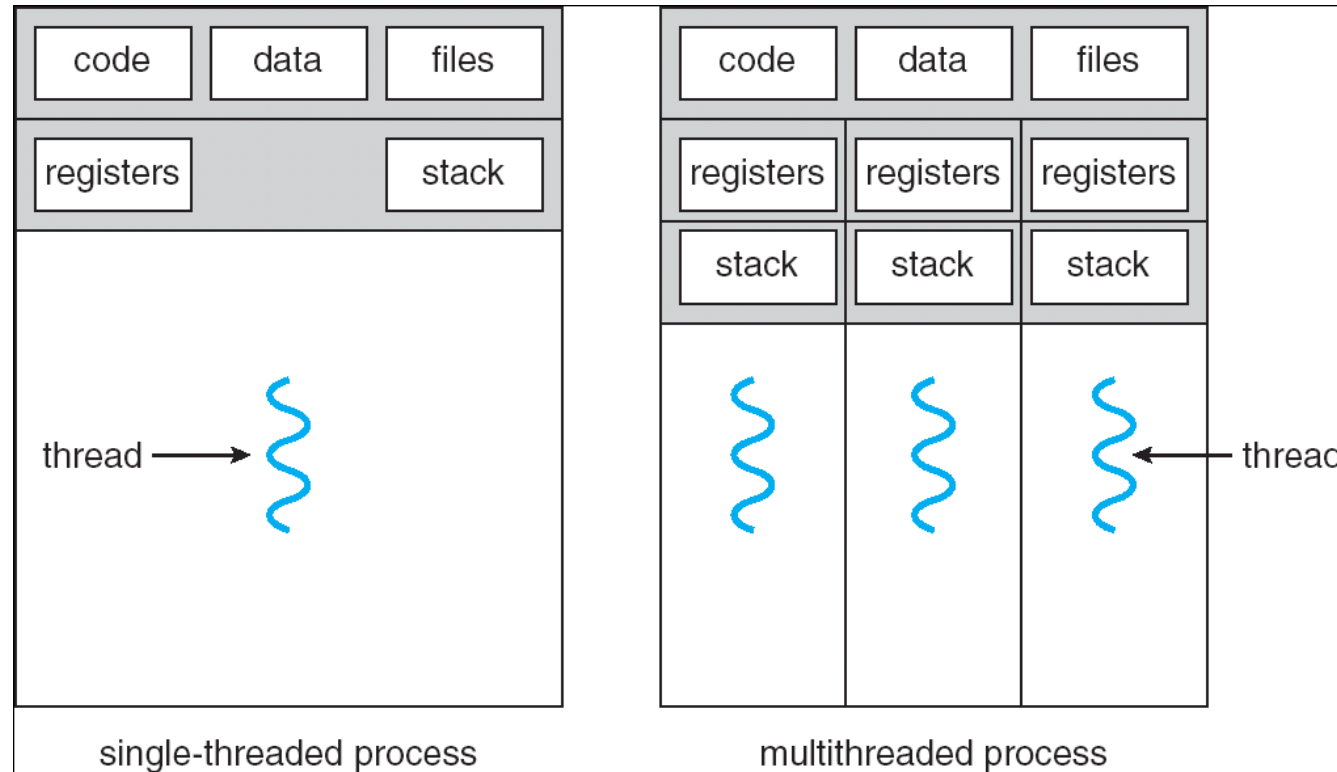
MULTICORE VS SMT



Recall thread

Process: multiple threads, code, data and OS state

Threads: share code, data, files, **not** regs or stack

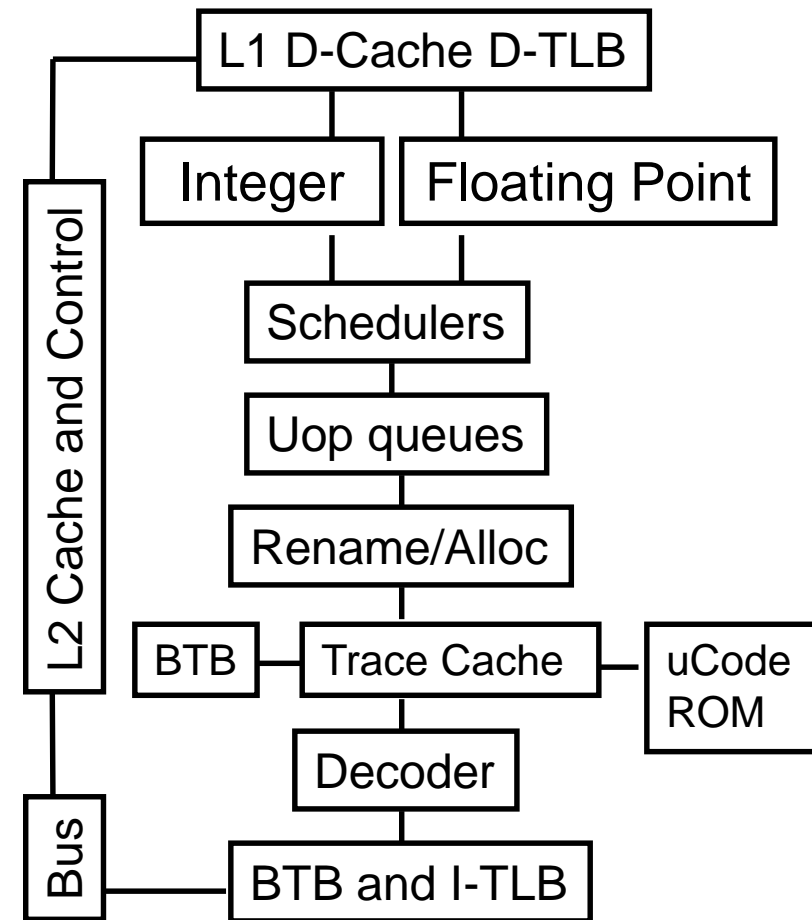


A technique complementary to multi-core: Simultaneous multithreading

- Problem addressed:
The processor pipeline can get stalled:

- Waiting for the result of a long floating point (or integer) operation
- Waiting for data to arrive from memory

Other execution units wait unused



Source: Intel





Simultaneous multithreading (SMT)

- Permits multiple independent threads to execute SIMULTANEOUSLY on the SAME core
- Weaving together multiple “threads” on the same core
- Example: if one thread is waiting for a floating point operation to complete, another thread can use the integer units

Standard Multithreading Picture

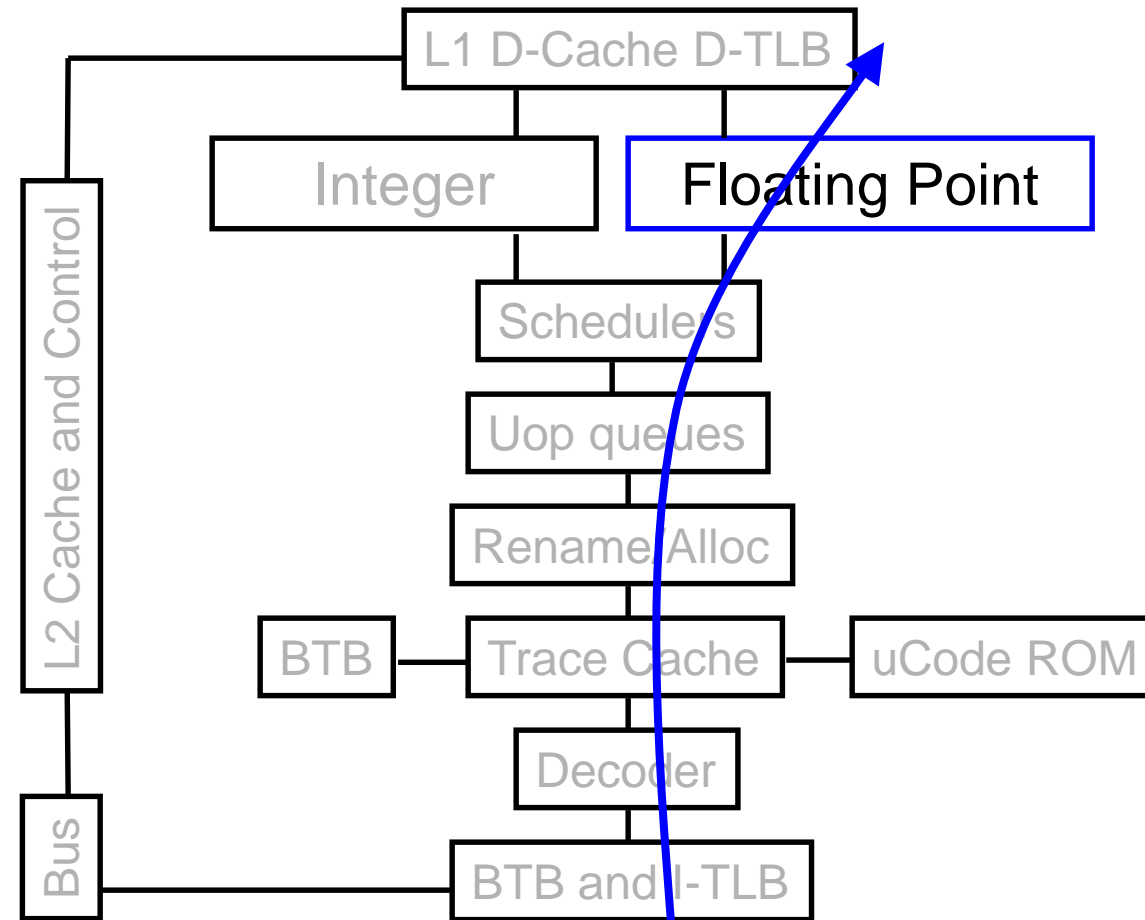
Time evolution of issue slots

- Color = thread, white = no instruction (bubble)




CG: coarse grain
FG: fine grain

Without SMT, only a single thread can run at any given time



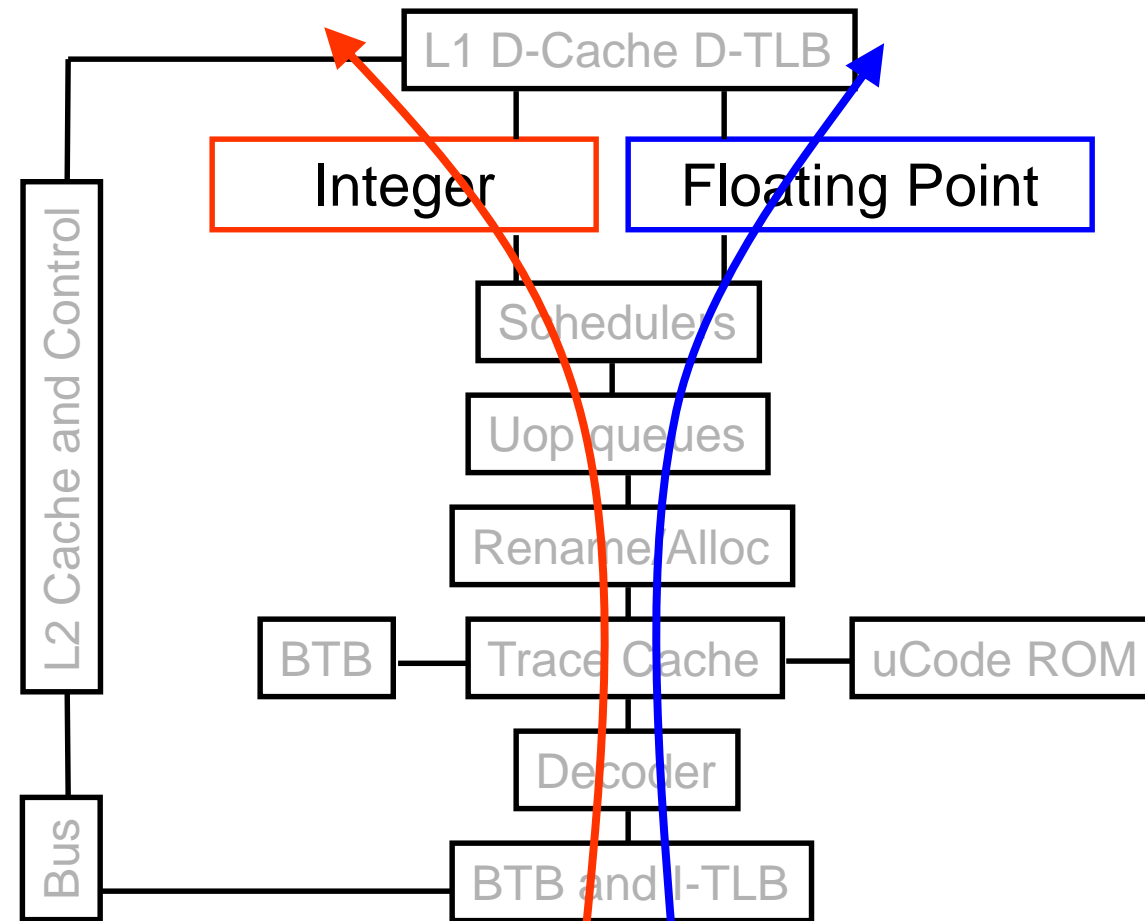
Thread 1: floating point

Without SMT, only a single thread can run at any given time



Thread 2:
integer operation

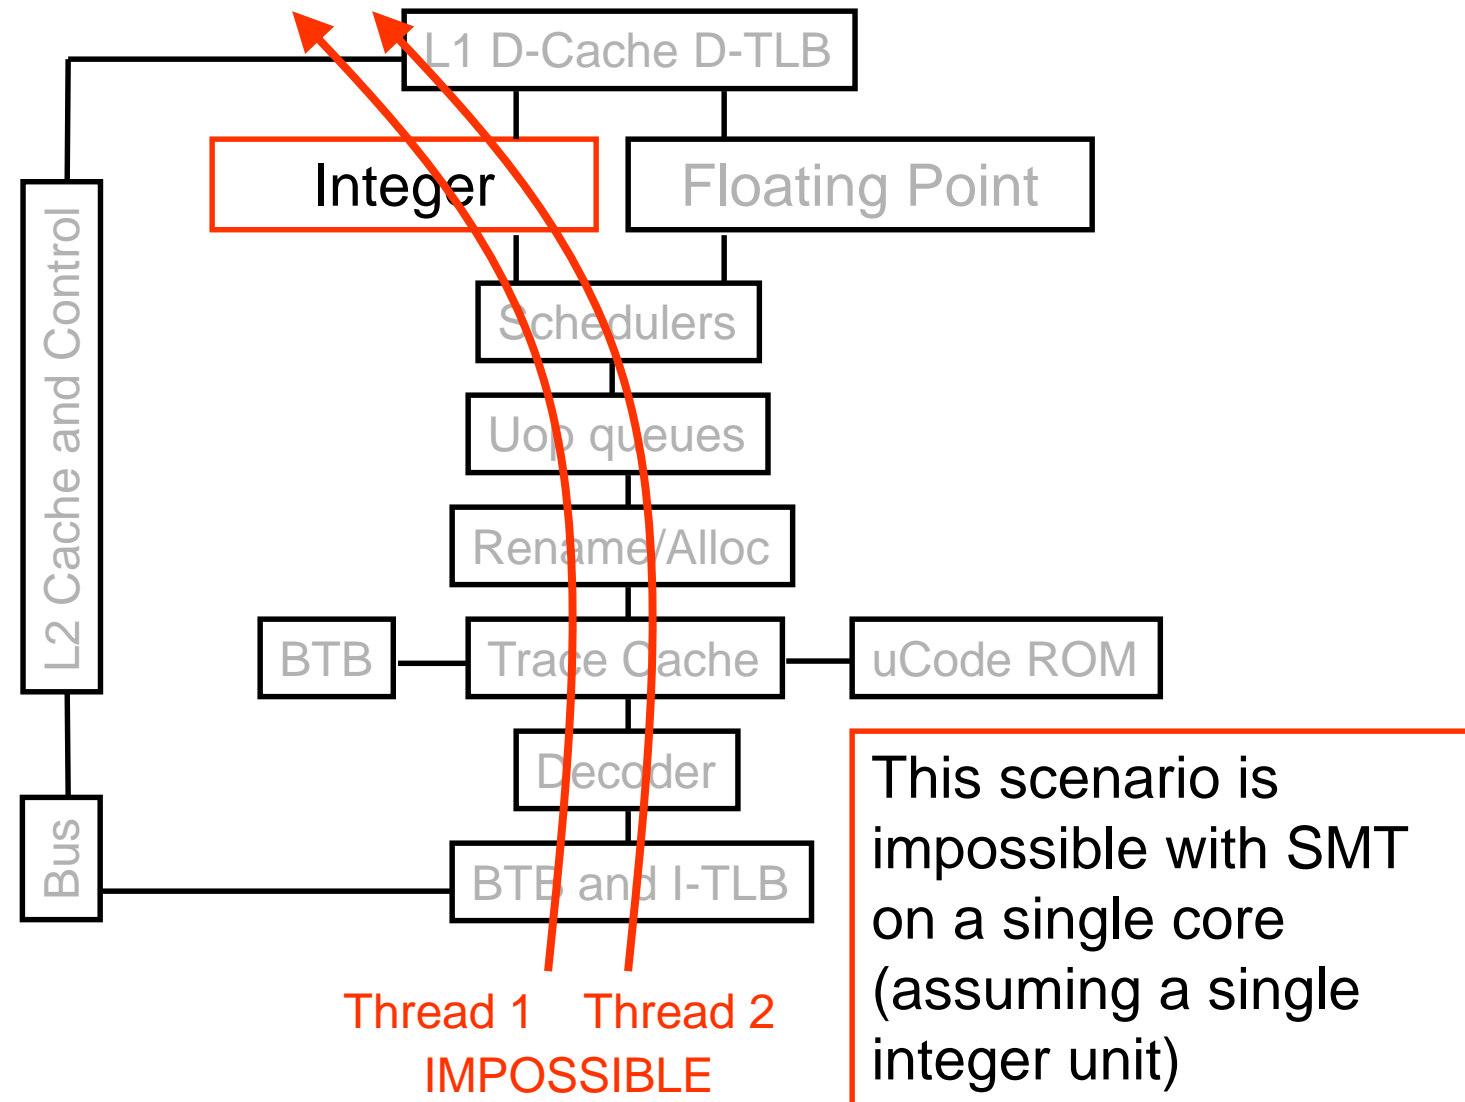
SMT processor: both threads can run concurrently



Thread 2:
integer operation

Thread 1: floating point

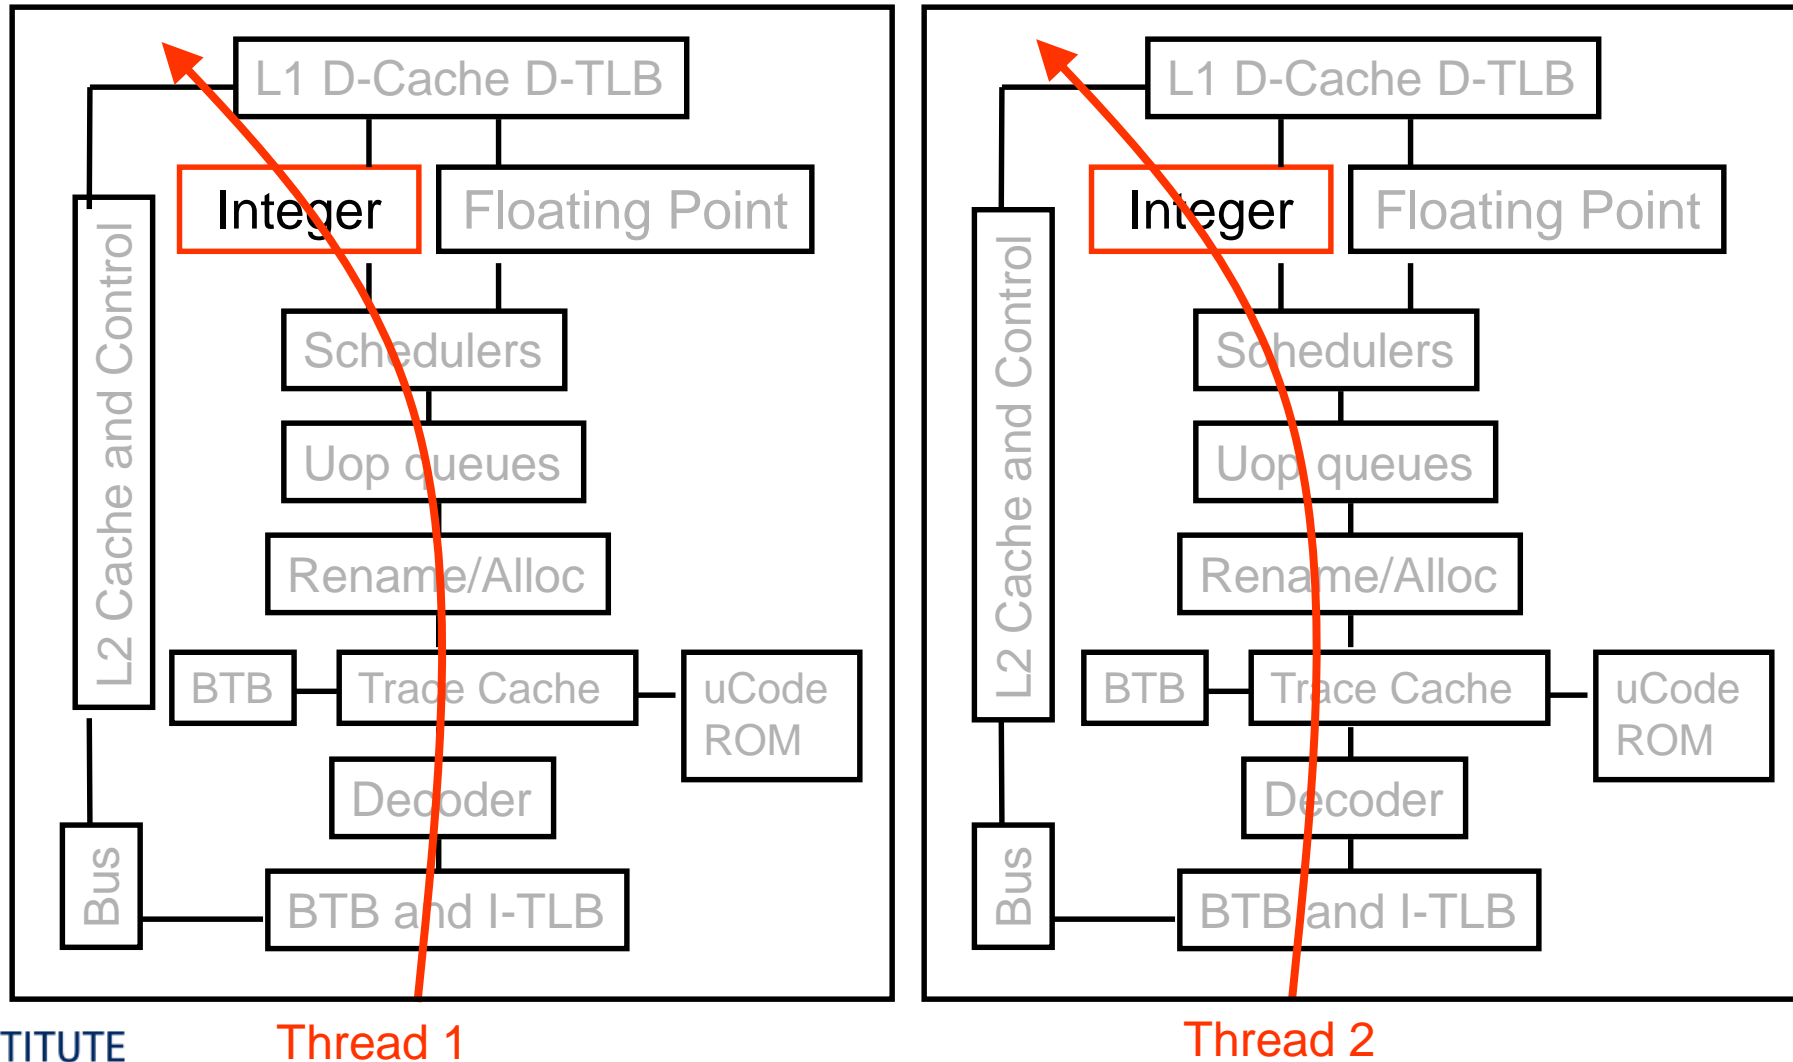
But: Can't simultaneously use the same functional unit



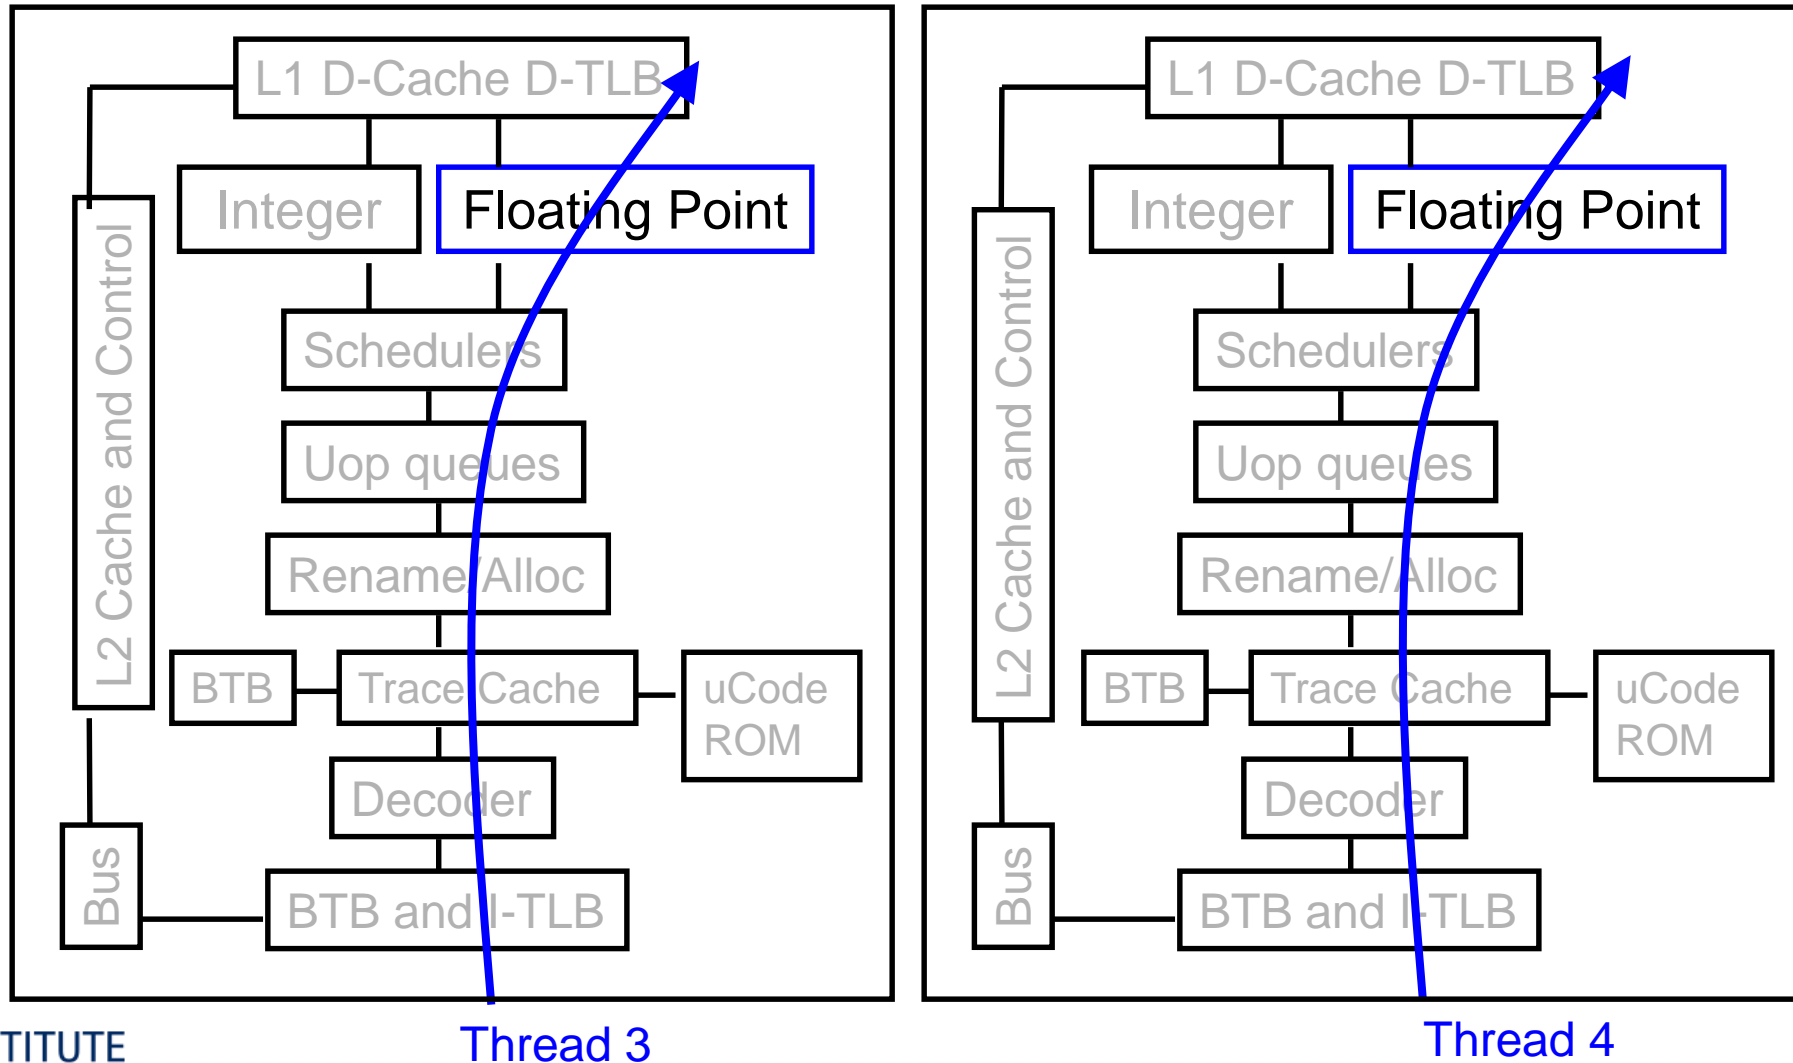
SMT not a “true” parallel processor

- Enables better threading (e.g. up to 30%)
- OS and applications perceive each simultaneous thread as a separate “virtual processor”
- The chip has only a single copy of each resource
- Recall GPU (SIMT)
- Compare to multi-core:
each core has its own copy of resources

Multi-core: threads can run on separate cores



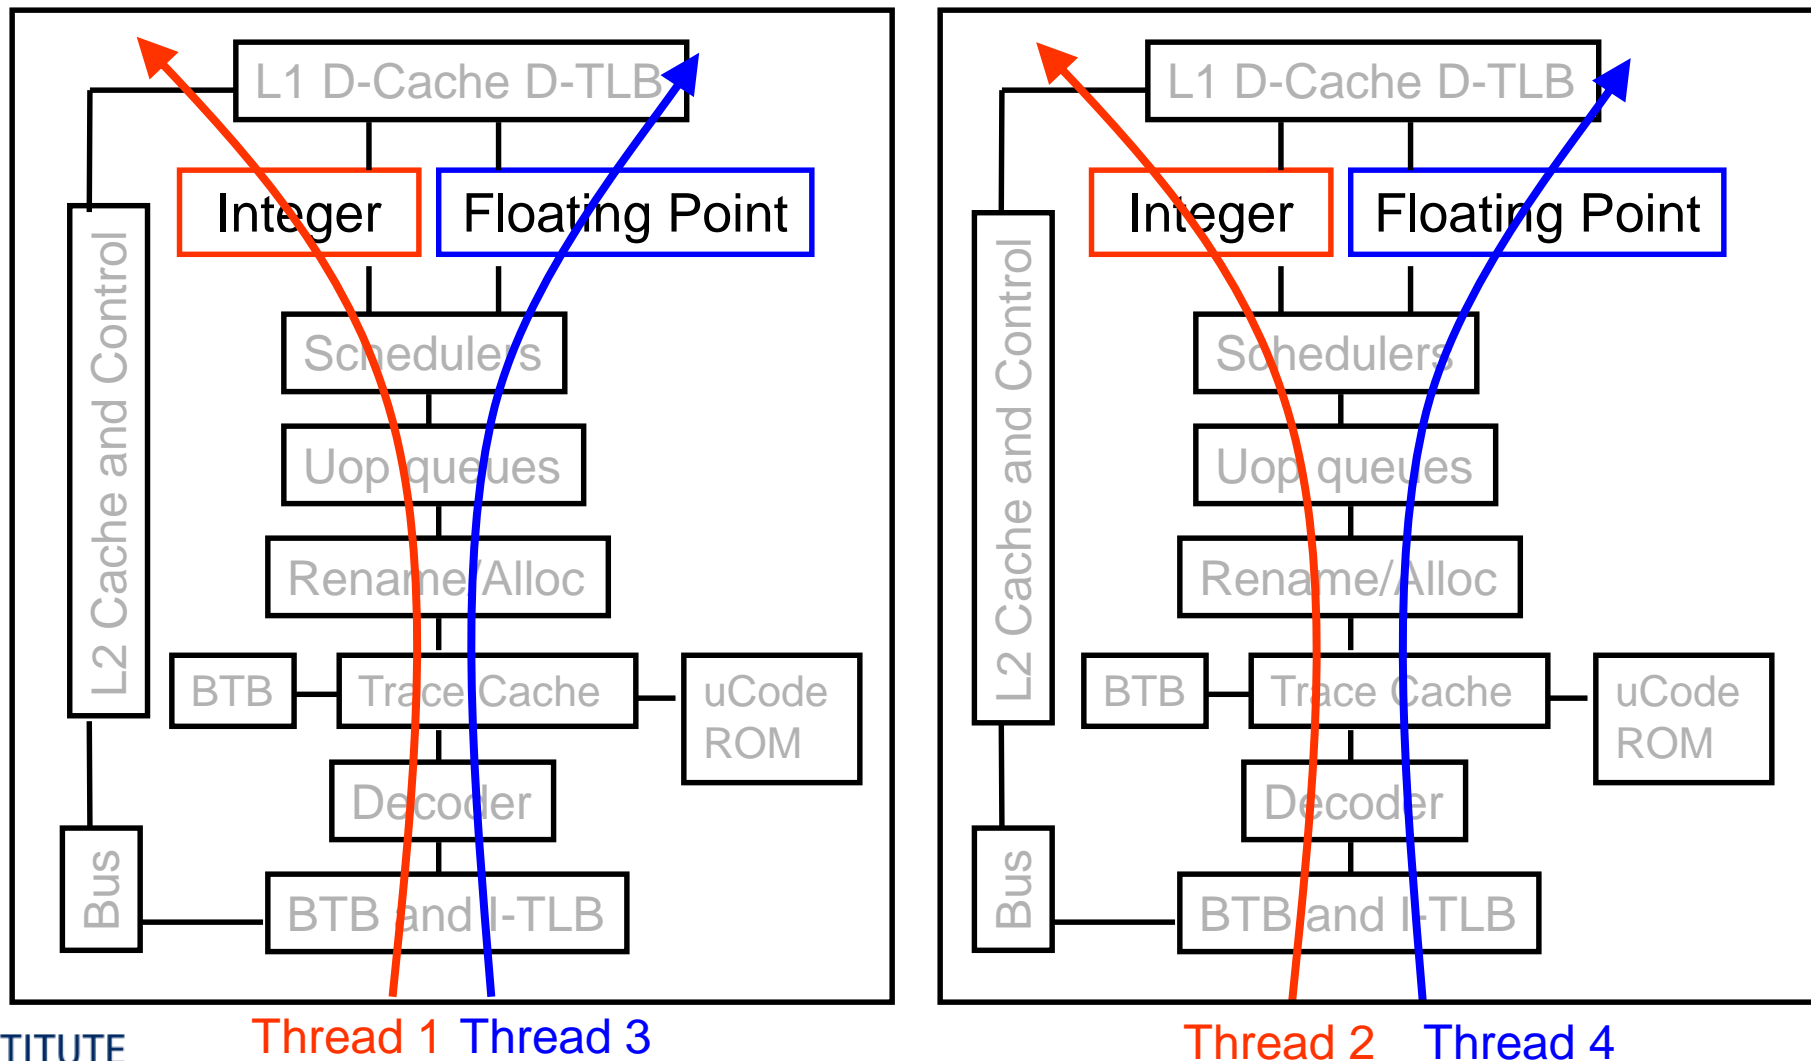
Multi-core: threads can run on separate cores



Combining Multi-core and SMT

- Cores can be SMT-enabled (or not)
- The different combinations:
 - Single-core, non-SMT: standard uniprocessor
 - Single-core, with SMT
 - Multi-core, non-SMT
 - Multi-core, with SMT: our fish machines
- The number of SMT threads:
2, 4, or sometimes 8 simultaneous threads
- Intel calls them “hyper-threads”

SMT Dual-core: all four threads can run concurrently



Comparison: multi-core vs SMT

- Advantages/disadvantages?

Comparison: multi-core vs SMT

- Multi-core:
 - Since there are several cores, each is smaller and not as powerful (but also easier to design and manufacture)
 - However, great with thread-level parallelism
- SMT
 - Can have one large and fast superscalar core
 - Great performance on a single thread
 - Mostly still only exploits instruction-level parallelism

Hyperthreading (HT)

Multi-Core vs. Multi-Issue vs. HT

Programs:
Num. Pipelines:
Pipeline Width:

N	1	N
N	1	1
1	N	N

Hyperthreads

- HT = MultiIssue + extra PCs and registers – dependency logic
- HT = MultiCore – redundant functional units + hazard avoidance


Hyperthreads (Intel)

- Illusion of multiple cores on a single core
- Easy to keep HT pipelines full + share functional units

Example: All of the above

8 die (aka 8 sockets)
4 core per socket
2 HT per core

Note: a socket is a processor, where each processor may have multiple processing cores, so this is an example of a multiprocessor multicore hyperthreaded system



COMMUNICATION




Communication

Message Passing

- In this paradigm, applications running on each core wrap up the data they want to communicate into messages sent to other cores.
- Explicit communication via *send* and *receive* operations
- Synchronization is implicit via blocks of messages.
- Distributed memory: In this model, each processor has its own (small) local memory, and its content is not replicated anywhere else


Shared Memory

- In this paradigm, there is a shared memory address space accessible to all cores, where they can read and write data.
- Implicit communication via memory accesses
- Synchronization is performed using atomic memory operations.




Message Passing

- Cores do not rely on a shared memory space.
- Each processing element (PE) has its core, data, I/O.
 - Explicit I/O to communicate with other cores
 - Synchronization via sending and receiving messages
- Advantages
 - Less hardware, easier to design
 - Focuses attention on costly non-local operations



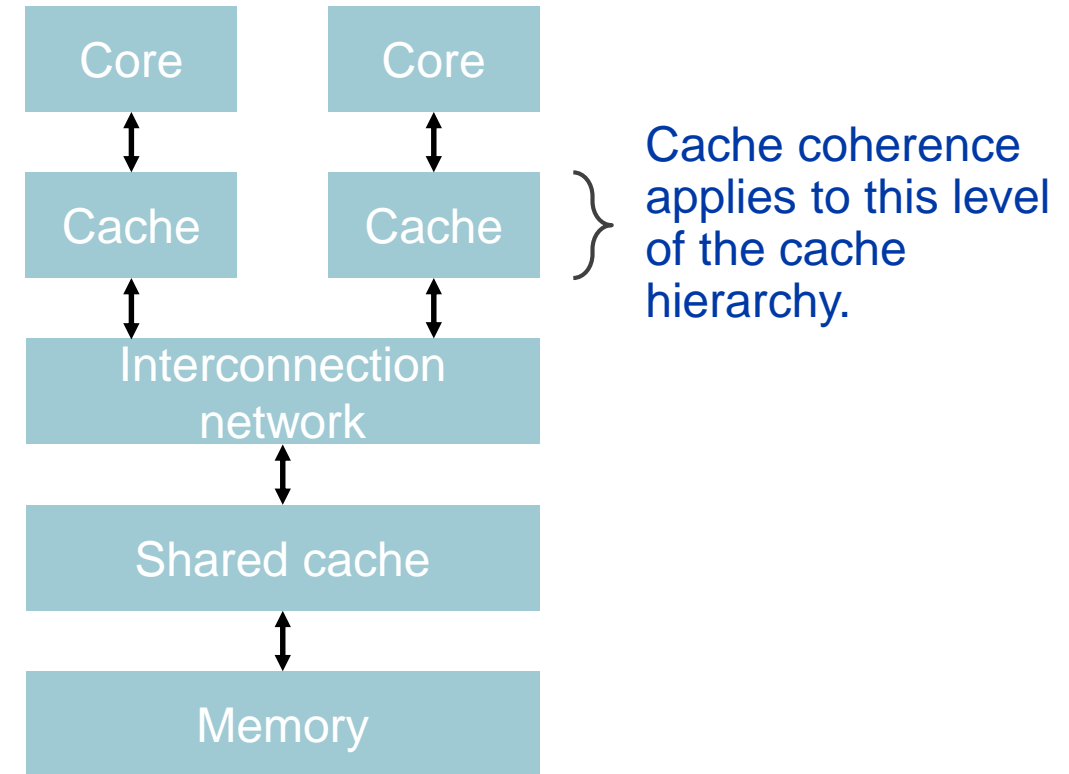
Shared Memory

- Cores share a memory that they see as a single address space.
- Cores may have caches holding data.
 - Communication involves reading and writing to locations in memory.
 - Synchronization via atomic operations to modify memory
 - Specific instructions provided in the ISA
 - The hardware guarantees correct operation



Shared Memory with Caches

- Caching shared-memory data has to be handled carefully.
 - The cache stores a copy of some of the data in memory.
- In particular, writes to the data must be dealt with correctly.
 - A core may write to data in its own cache.
 - This makes copies in other caches become stale.
 - The version in memory won't get updated immediately either, if it's a write-back cache.
 - In these situations, there is a danger of one core subsequently reading a stale (old) value.



Distributed vs shared caches

- Advantages of distributed:
 - They are closer to core, so faster access
 - Reduces contention
- Advantages of shared:
 - Threads on different cores can share the same cache data
 - More cache space available if a single (or a few) high-performance thread runs on the system
 - Matches the programmer's view of the system
 - Hardware handles communication implicitly

CACHE COHERENCE



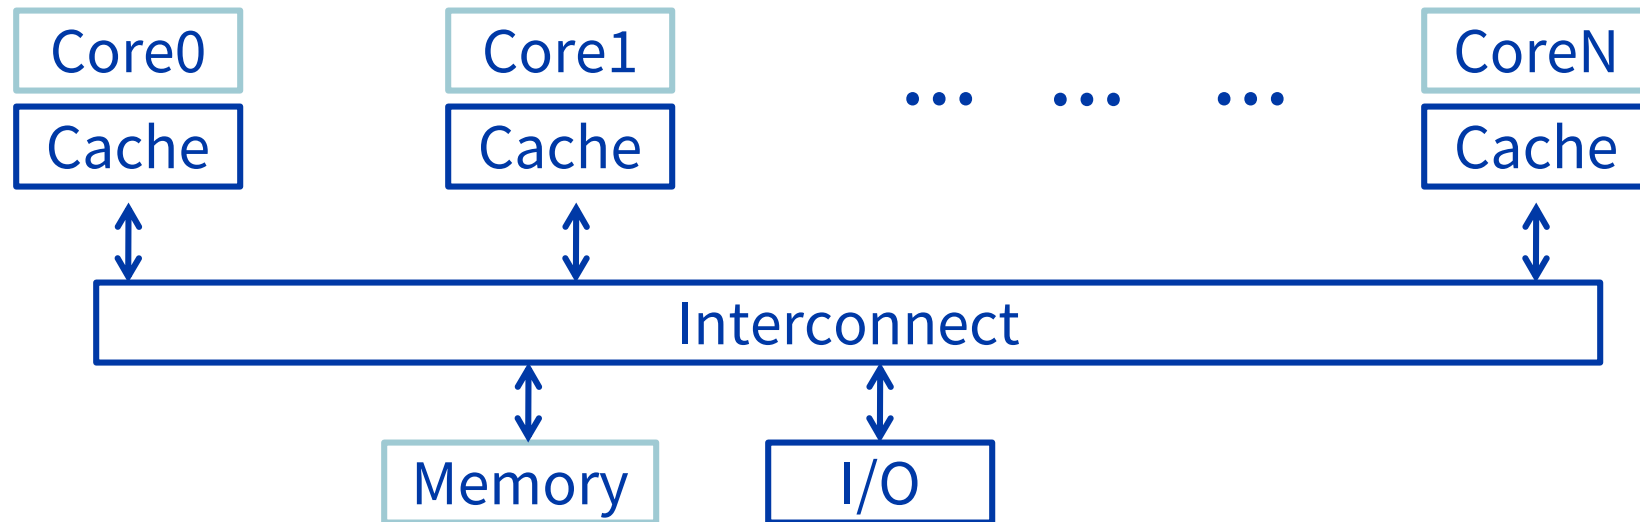
Recall: 3C Model -> 4C Model

- Divide cache misses into three categories
 - **Compulsory (cold)**: never seen this address before
 - Would miss even in infinite cache
 - **Capacity**: miss caused because cache is too small
 - Would miss even in fully associative cache
 - Identify? Consecutive accesses to block separated by access to at least N other distinct blocks (N is number of frames in cache)
 - **Conflict**: miss caused because cache associativity is too low
 - Identify? **All other misses**
 - **COHERENCE: miss due to external invalidations**
 - **ONLY IN SHARED MEMORY MULTIPROCESSORS**

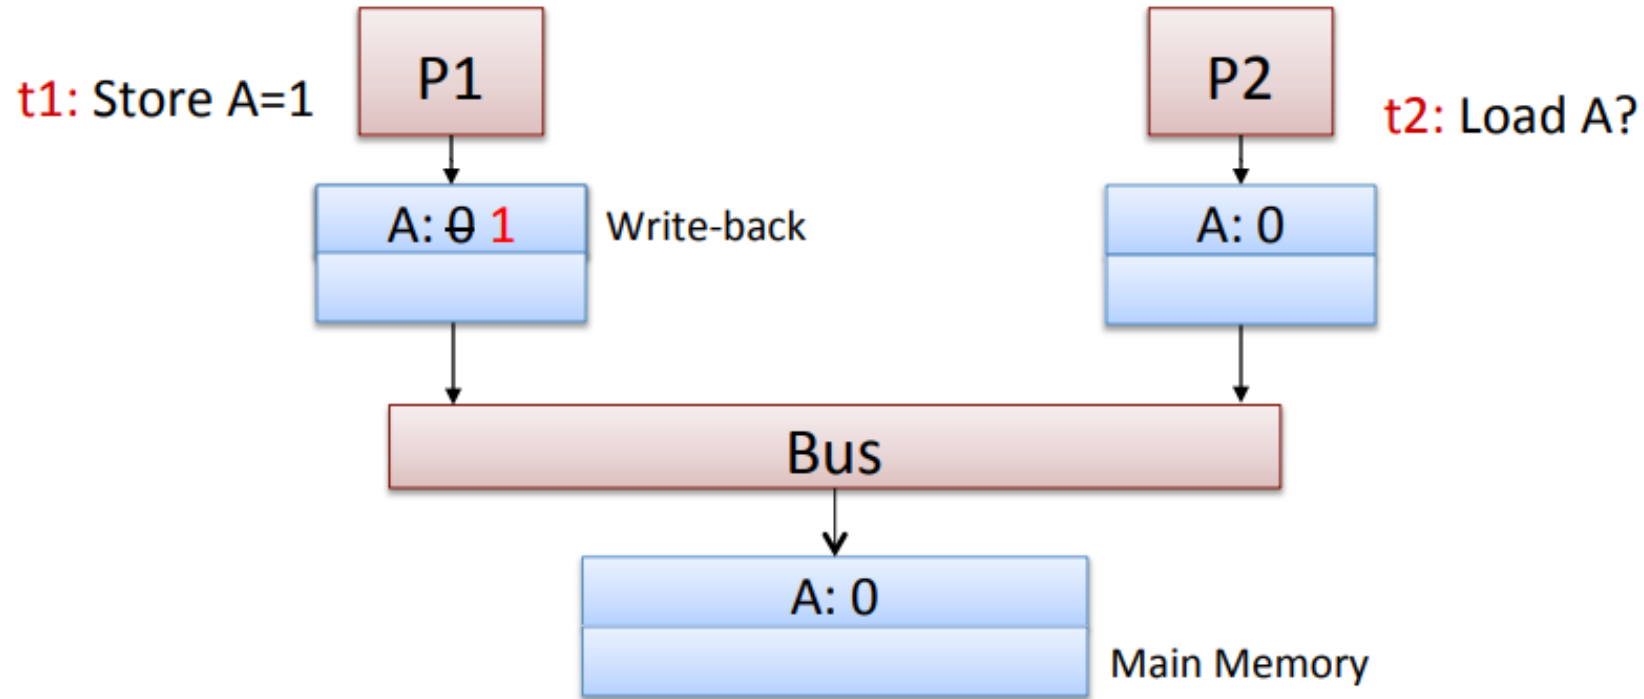
Shared Memory Multiprocessors

Shared Memory Multiprocessor (SMP)

- Typical (today): 2 – 4 processor dies, 2 – 8 **cores** each
- HW provides *single physical address space* for all processors



Cache Coherency Problem




Cache Coherency Problem

Thread A (on Core0)
`for(int i = 0, i < 5; i++) {
 x = x + 1;
}`

Thread B (on Core1)
`for(int j = 0; j < 5; j++) {
 x = x + 1;
}`

What will the value of `x` be after both loops finish?



Cache Coherency Problem

Thread A (on Core0)

```
for(int i = 0, i < 5; i++) {  
    x = x + 1;  
}
```

Thread B (on Core1)

```
for(int j = 0; j < 5; j++) {  
    x = x + 1;  
}
```

What will the value of **x** be after both loops finish?

(x starts as 0)

- A. 6
- B. 8
- C. 10
- D. Could be any of the above
- E. Couldn't be any of the above

Cache Coherency Problem

Thread A (on Core0)

```
for(int i = 0, i < 5; i++) {  
    x = x + 1;  
}
```

Thread B (on Core1)

```
for(int j = 0; j < 5; j++) {  
    x = x + 1;  
}
```

What will the value of **x** be after both loops finish?

(x starts as 0)

A. 6

B. 8

C. 10

D. Could be any of the above

E. Couldn't be any of the above

Cache Coherency Problem, WB \$

Thread A (on Core0)

```
for(int i = 0, i < 5; i++) {
```

```
t0=0 LW t0, addr(x)
```

```
t0=1 ADDIU t0, t0, 1
```

```
x=1 SW t0, addr(x)
```

```
}
```

Thread B (on Core1)

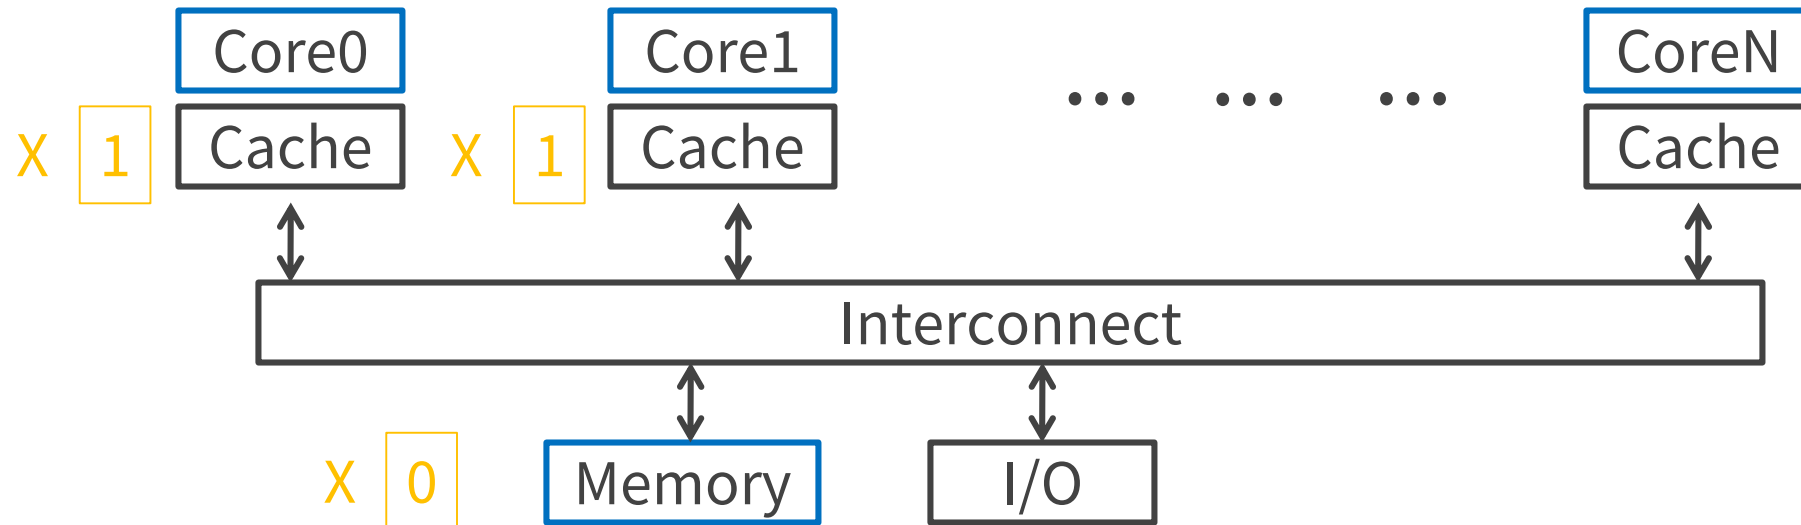
```
for(int j = 0; j < 5; j++) {
```

```
t0=0 LW t0, addr(x)
```

```
t0=1 ADDIU t0, t0, 1
```

```
x=1 SW t0, addr(x)
```


Problem!



Not just a problem for Write-Back Caches

Executing on a write-thru cache

Time step	Event	CPU 0's cache	CPU 1's cache	Memory
0				0
1	CPU 0 reads X	0		0
2	CPU 2 reads X	0	0	0
3	CPU 1 writes 1 to X	1	0	1



Two issues here

Coherence

- What values can be returned by a read
- Need a globally uniform (consistent) view of a single memory location

Solution: Cache Coherence Protocols

Consistency

- When a written value will be returned by a read
- Need a globally uniform (consistent) view of *all memory locations relative to each other*

Solution: Memory Consistency Models

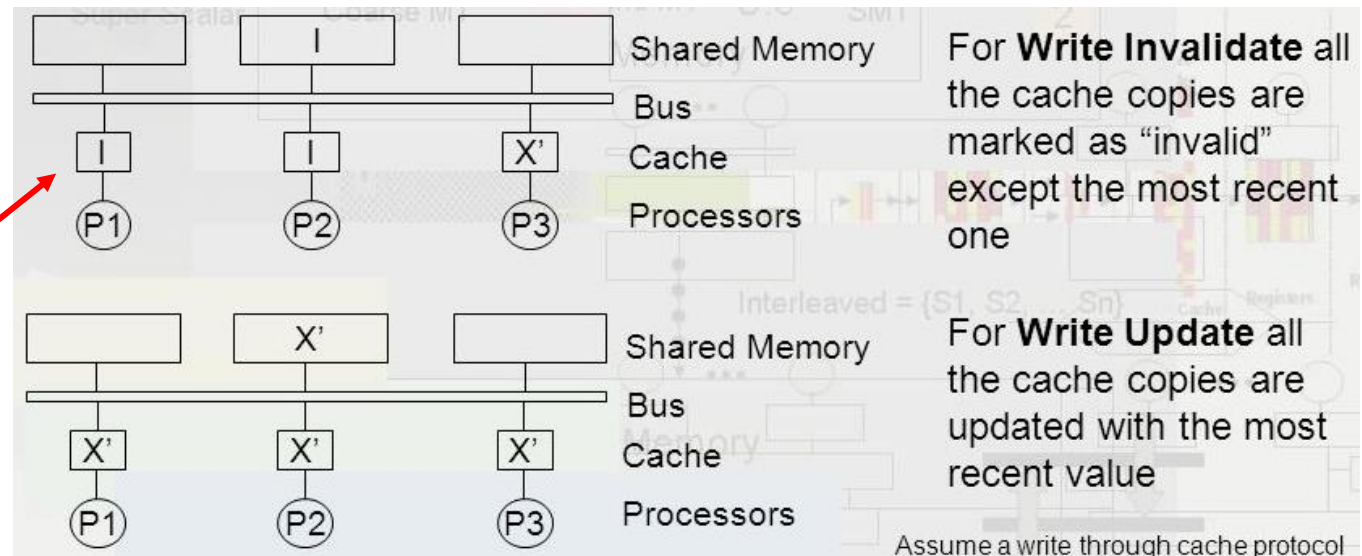
Approaches to Cache Coherence

- Software-based solutions
 - Mechanisms
 - Mark cache blocks/memory pages as cacheable/non-cacheable
 - Add “Flush” and “Invalidate” instructions
 - Could be done by compiler or runtime system
 - Difficult to get perfect
- Hardware solutions are far more common
 - Simple schemes rely on broadcast over a bus (Snooping Protocol)

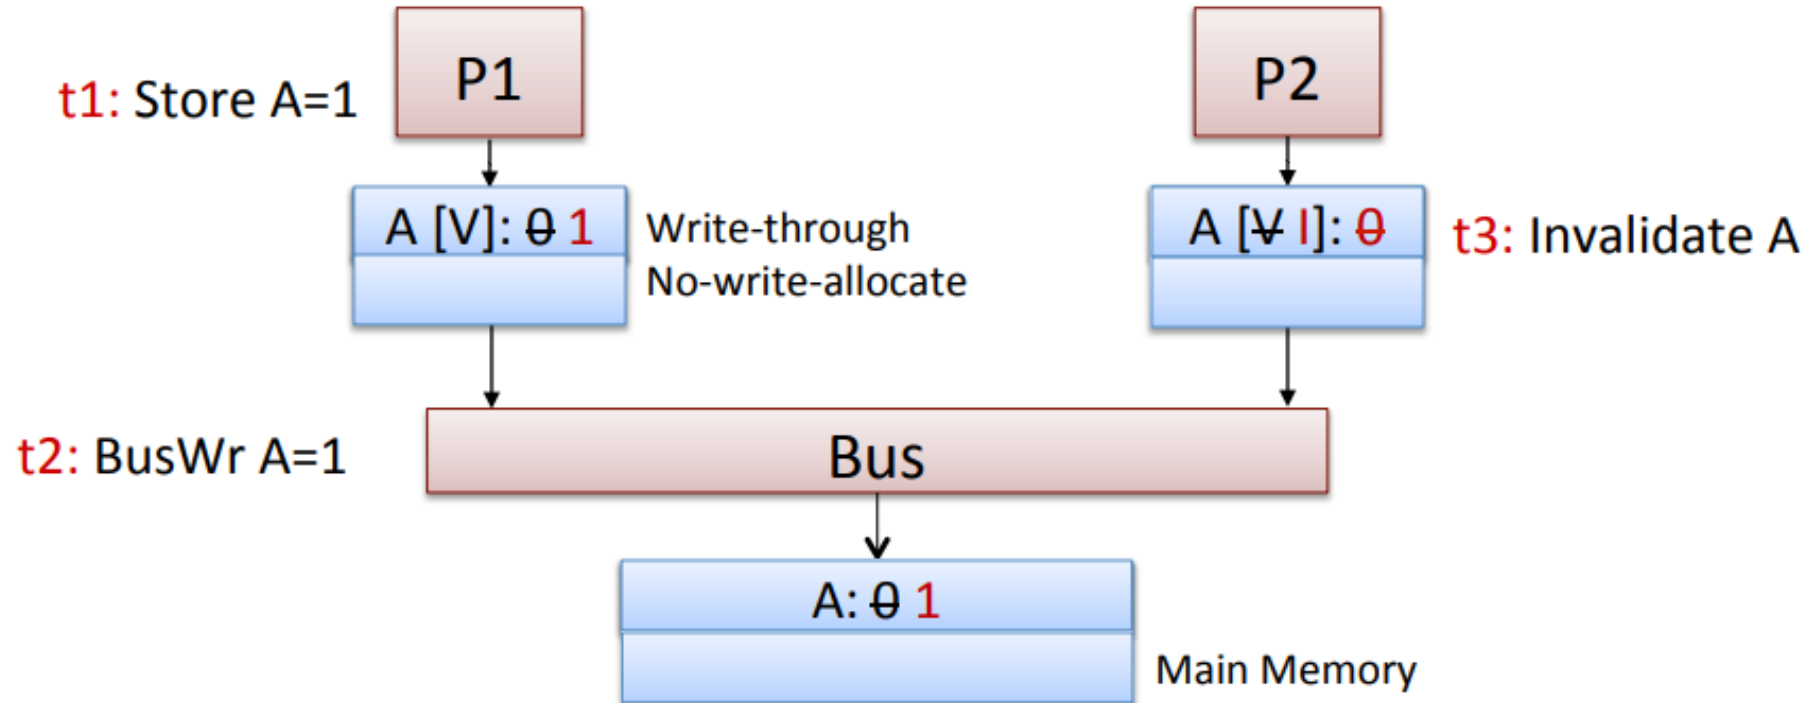
Snooping-based Protocol

- broadcasting coherence information to all processors over the shared interconnect
- Every time a cache miss occurred, the triggering cache communicated with all other caches!
- Two types
 - Write update
 - Write invalidate

Experience has shown that invalidate protocols use significantly less bandwidth.




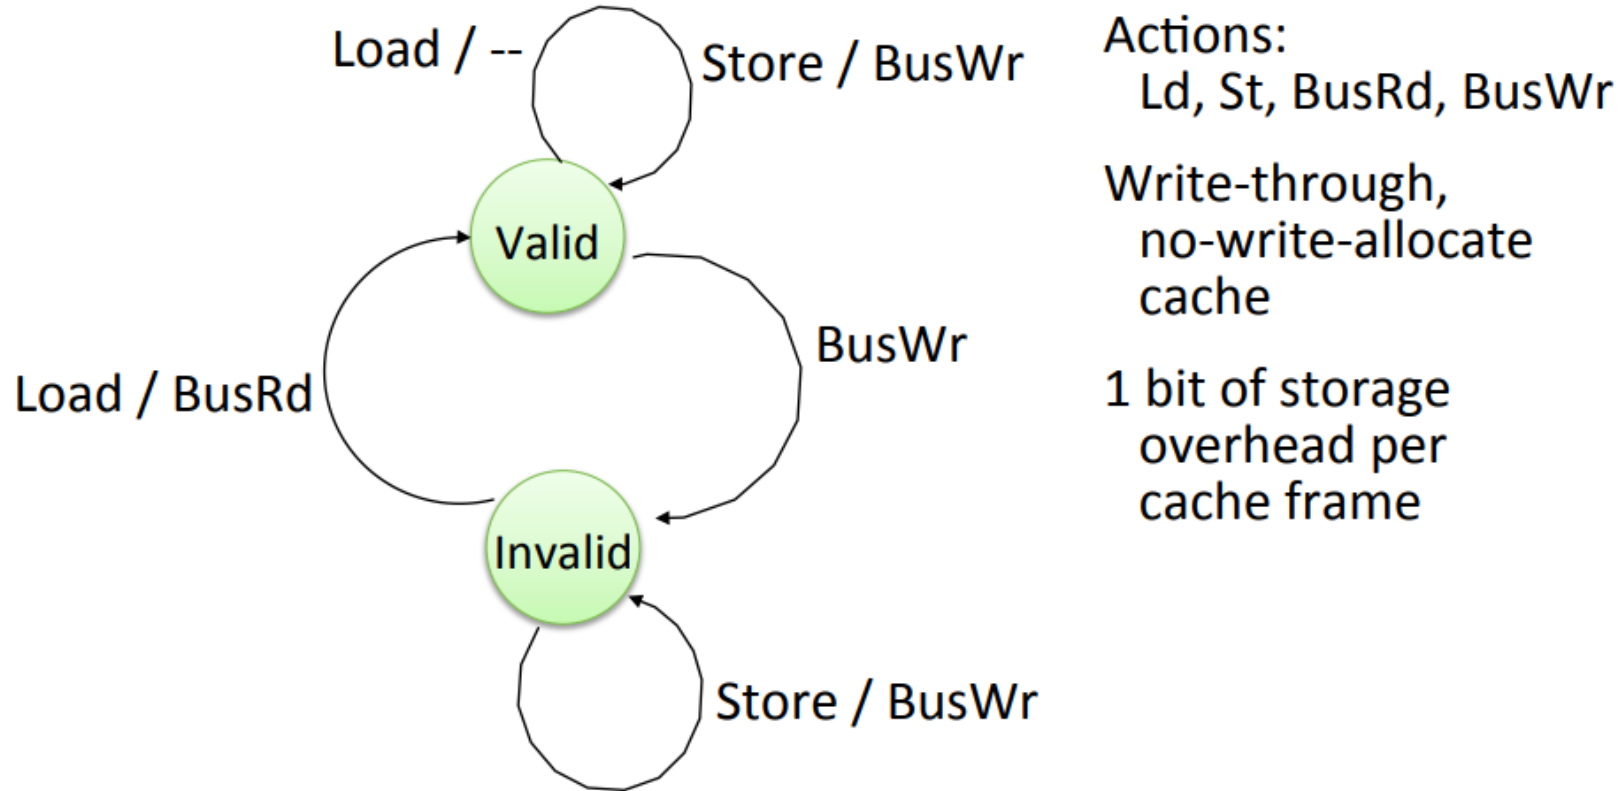
Simple Write-Through Scheme: Valid-Invalid Coherence



Valid-Invalid Coherence

- Allows multiple readers, but must write through to bus
 - Write-through, no-write-allocate cache
- All caches must monitor (aka “snoop”) all bus traffic
 - simple state machine for each cache frame

Valid-Invalid (VI) Snooping Protocol

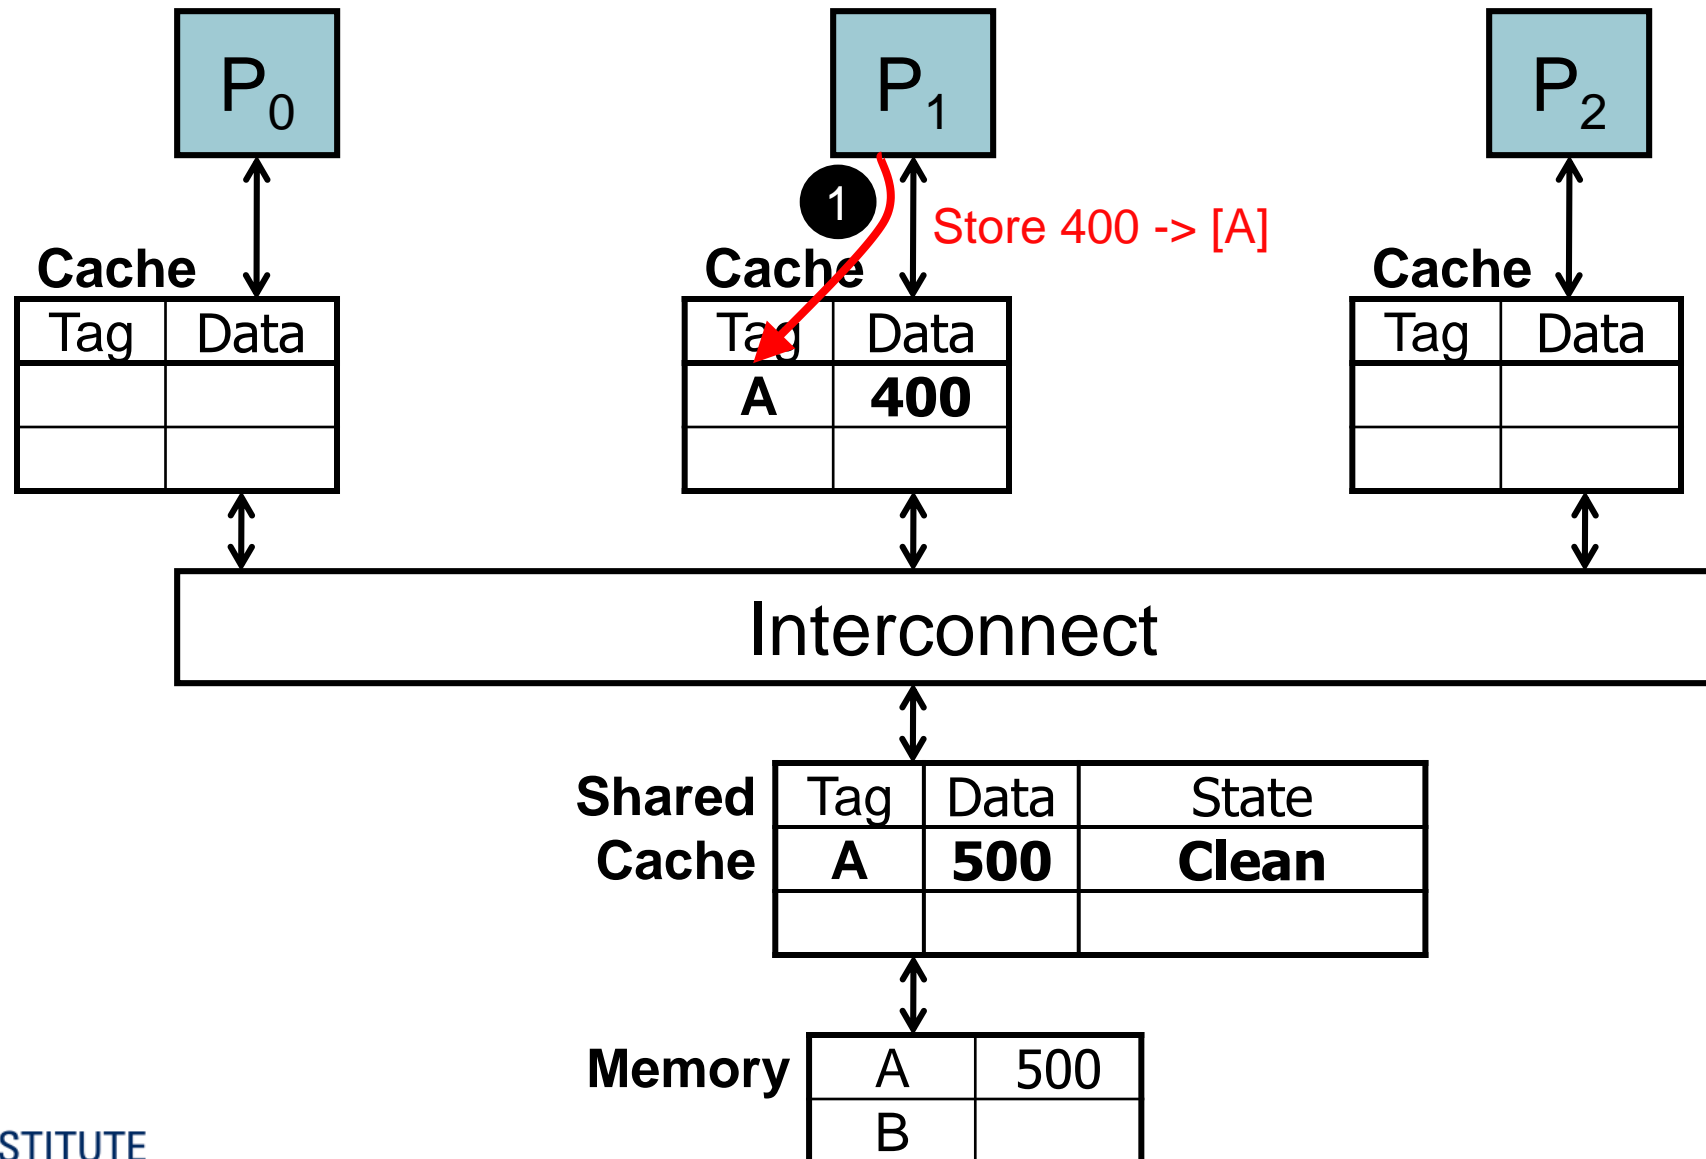


- If *you* load/store a block: transition to **V**
- If anyone *else* wants to read/write block:
 - Give it up: transition to **I** state
 - Write-back if your own copy is dirty

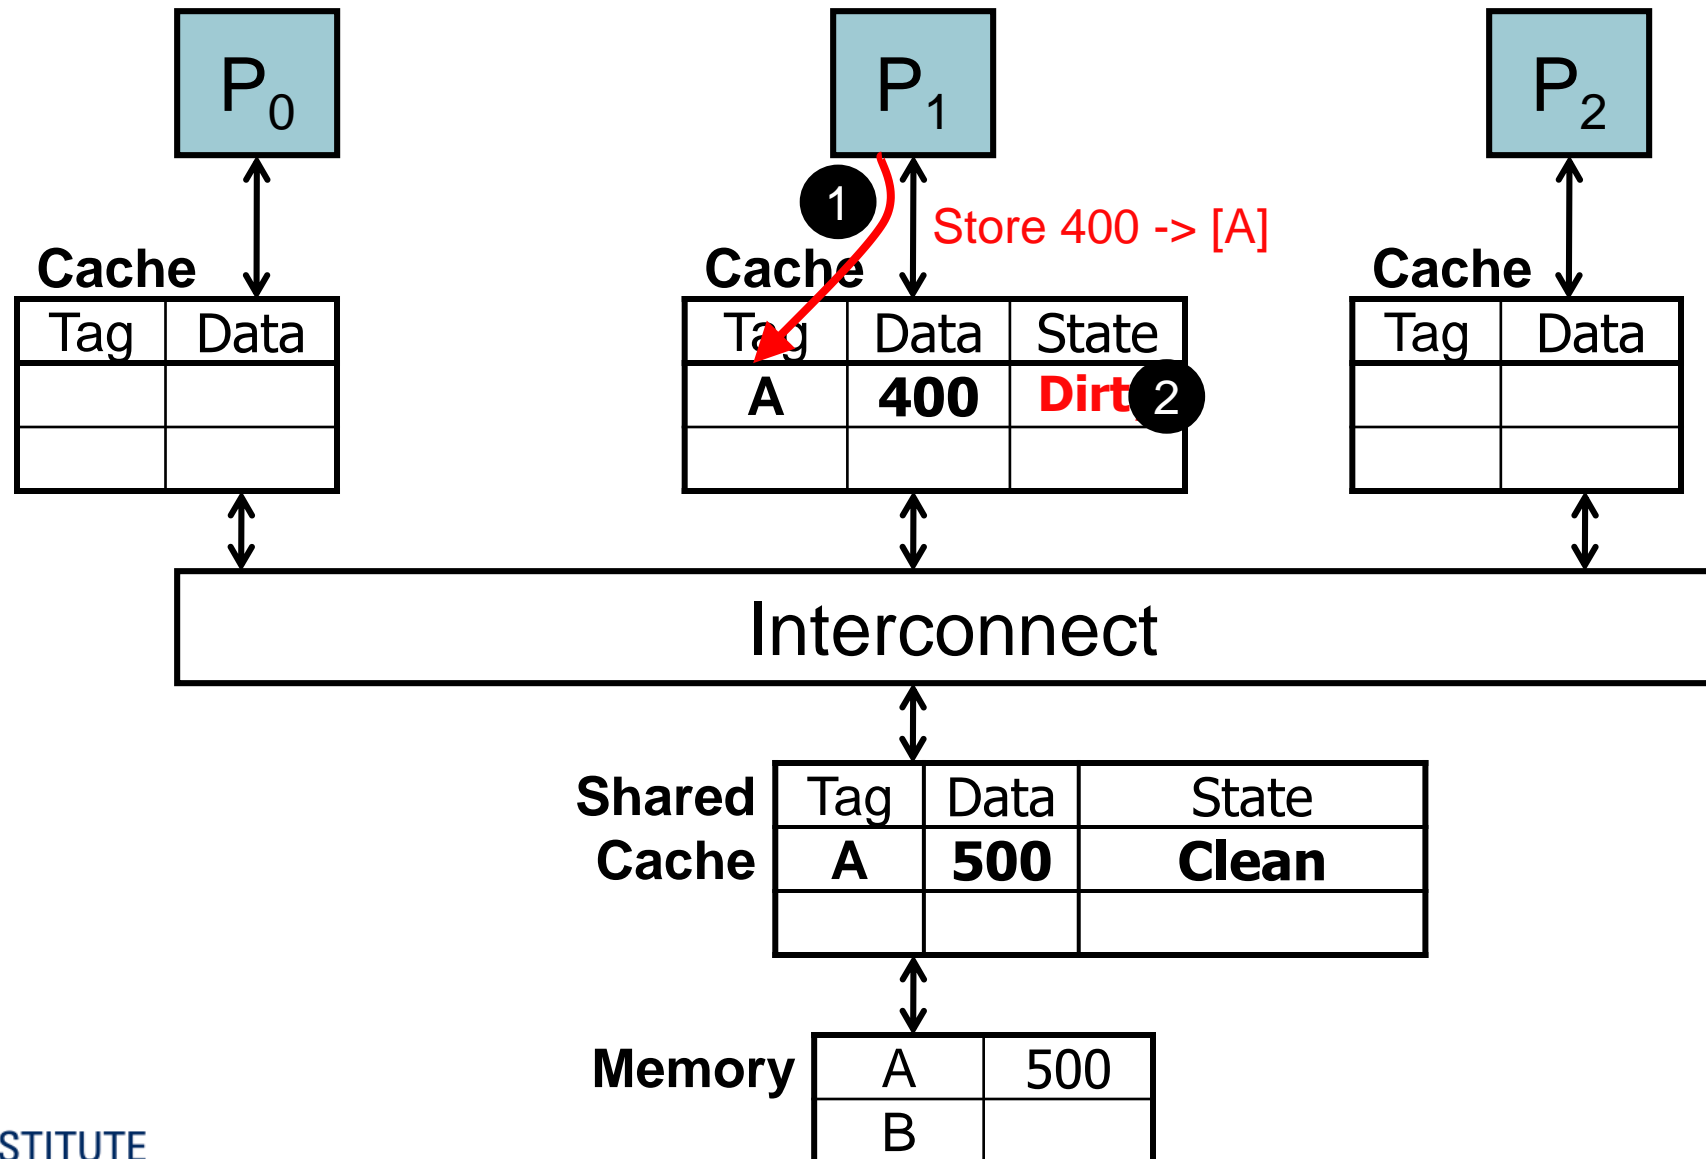
VI Protocol

- **So far: Option #1 “Snooping”: broadcast to all, whoever has it responds**
- **Option #2: “Directory”: track sharers with separate structure**

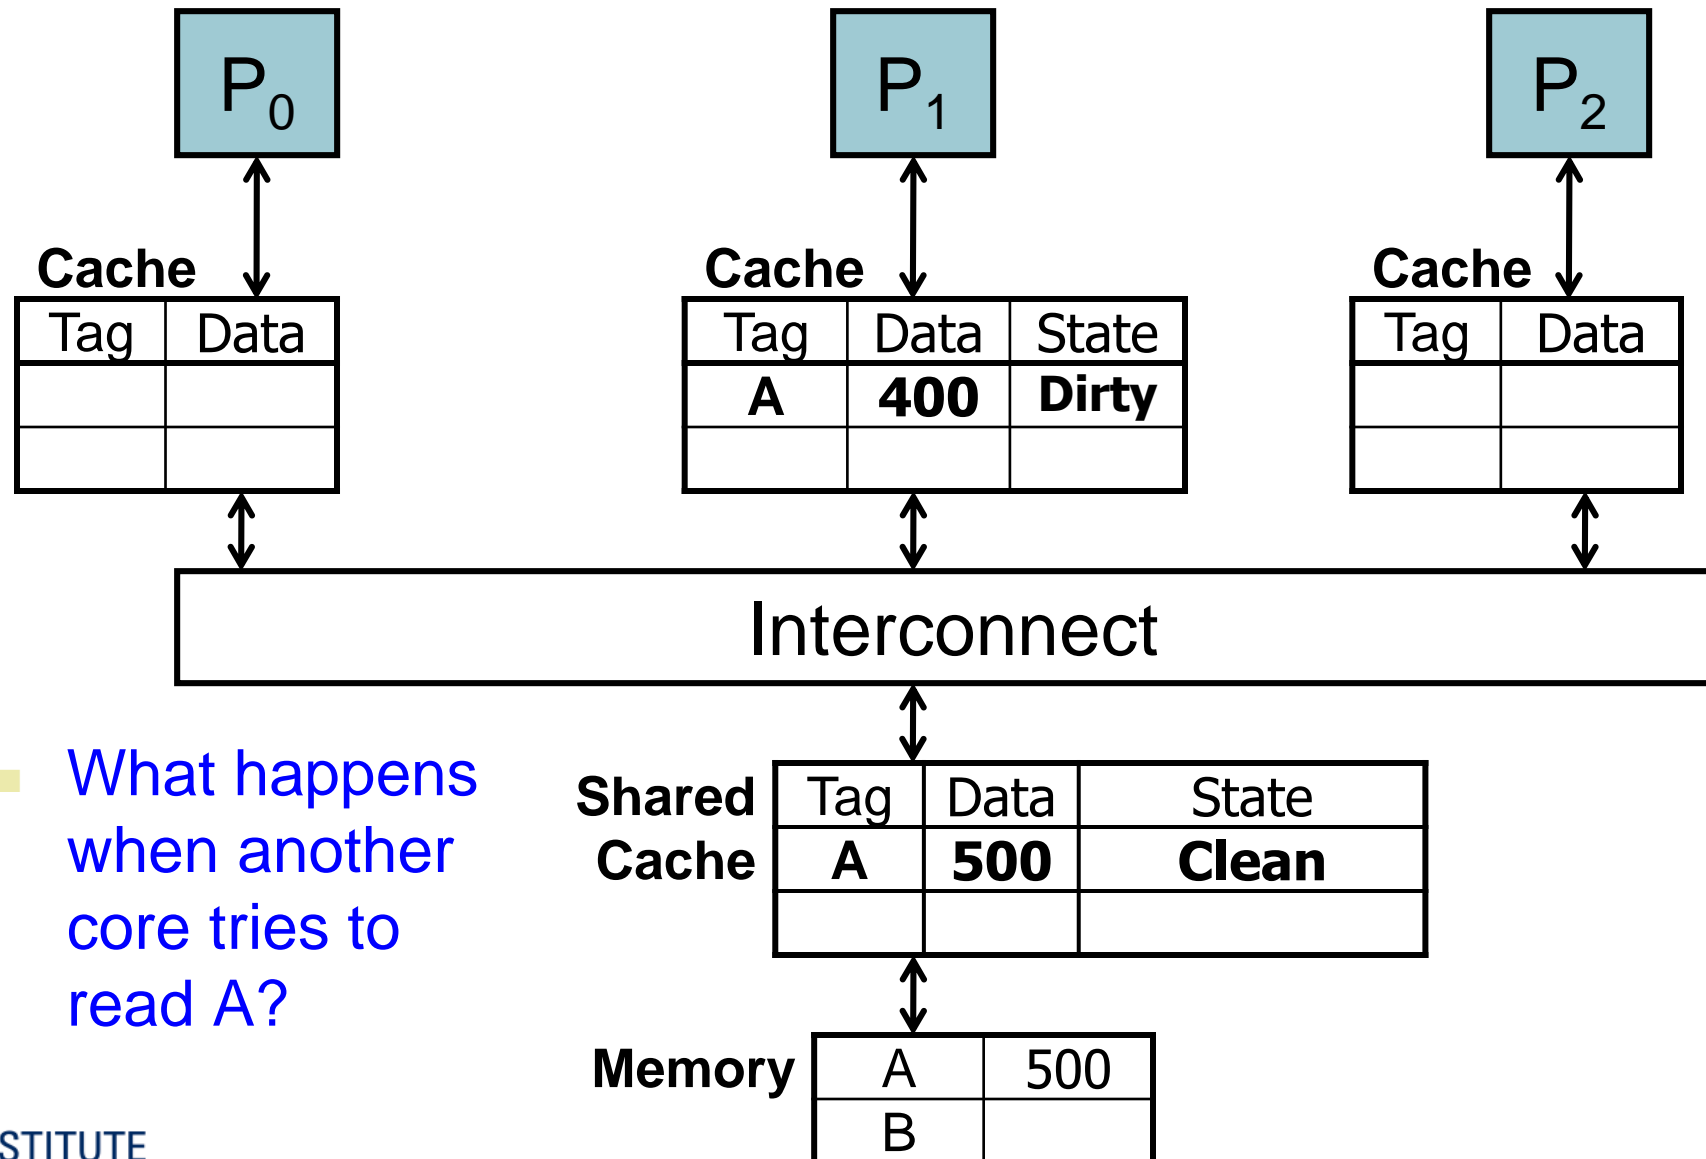
VI Protocol Example



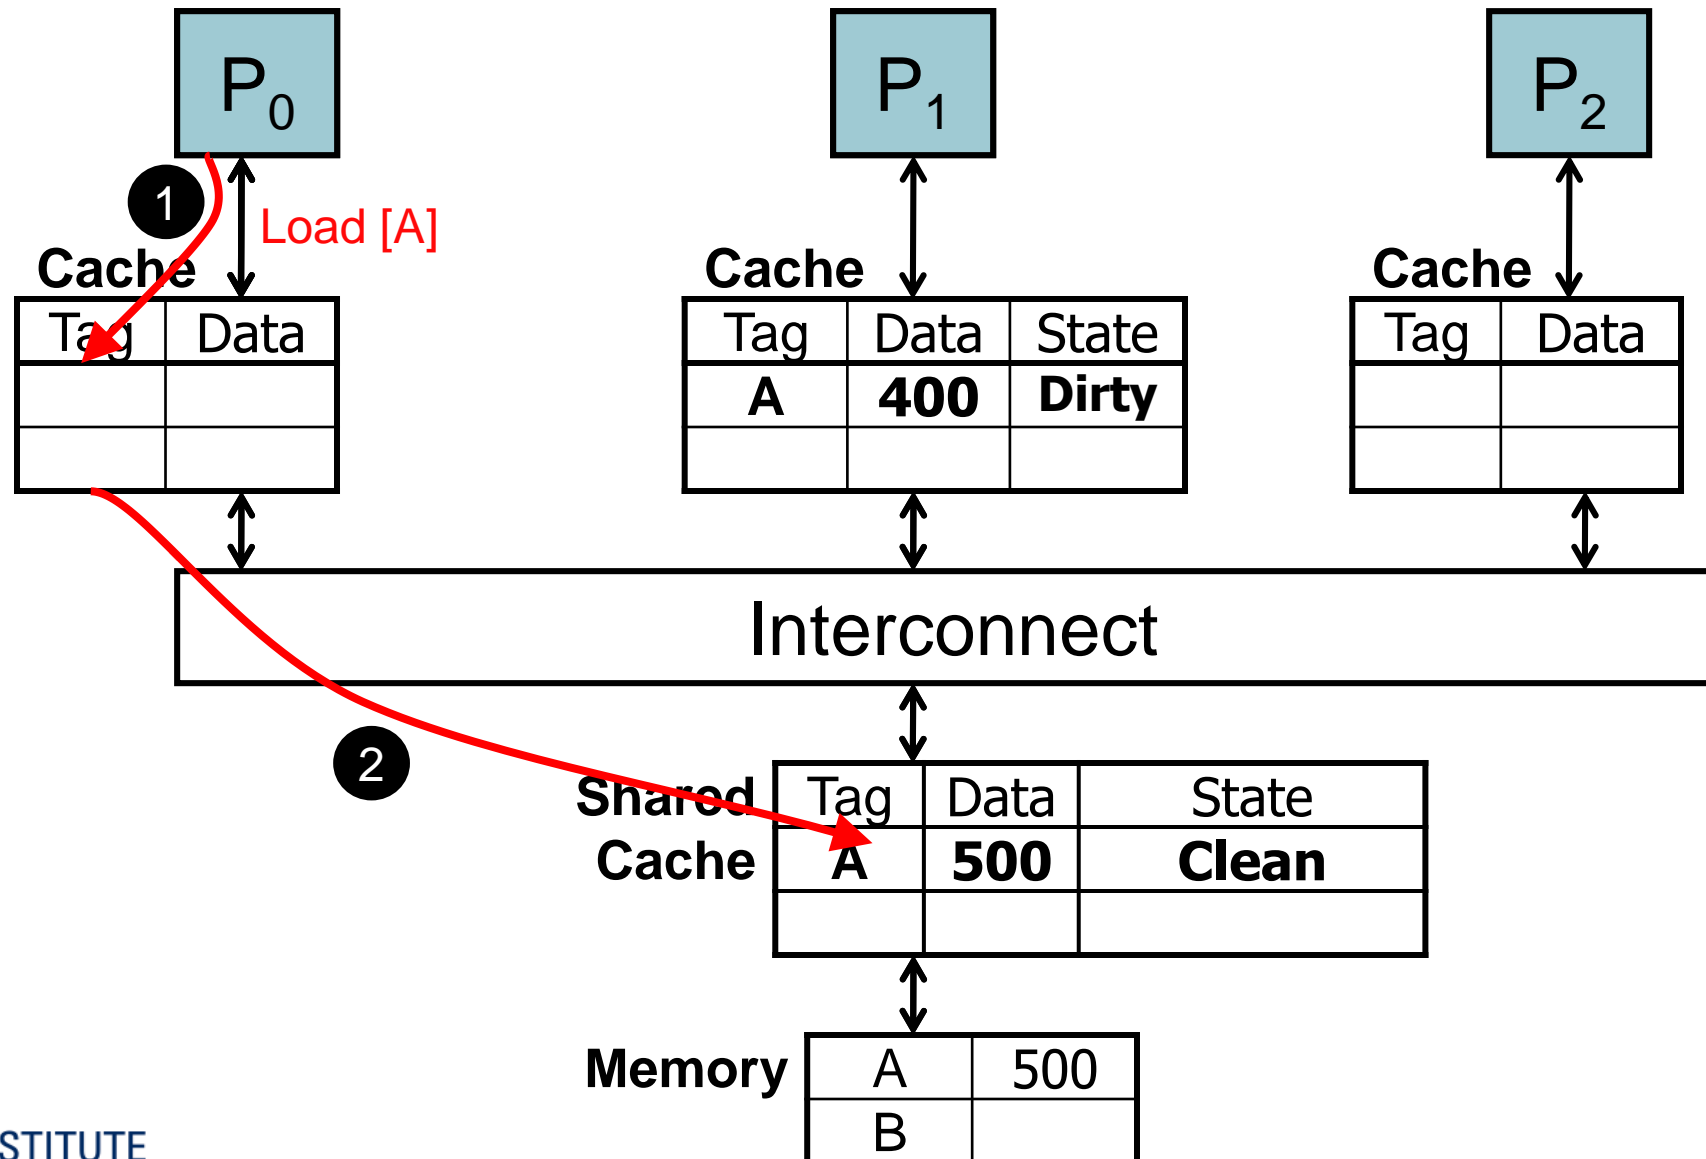
VI Protocol Example



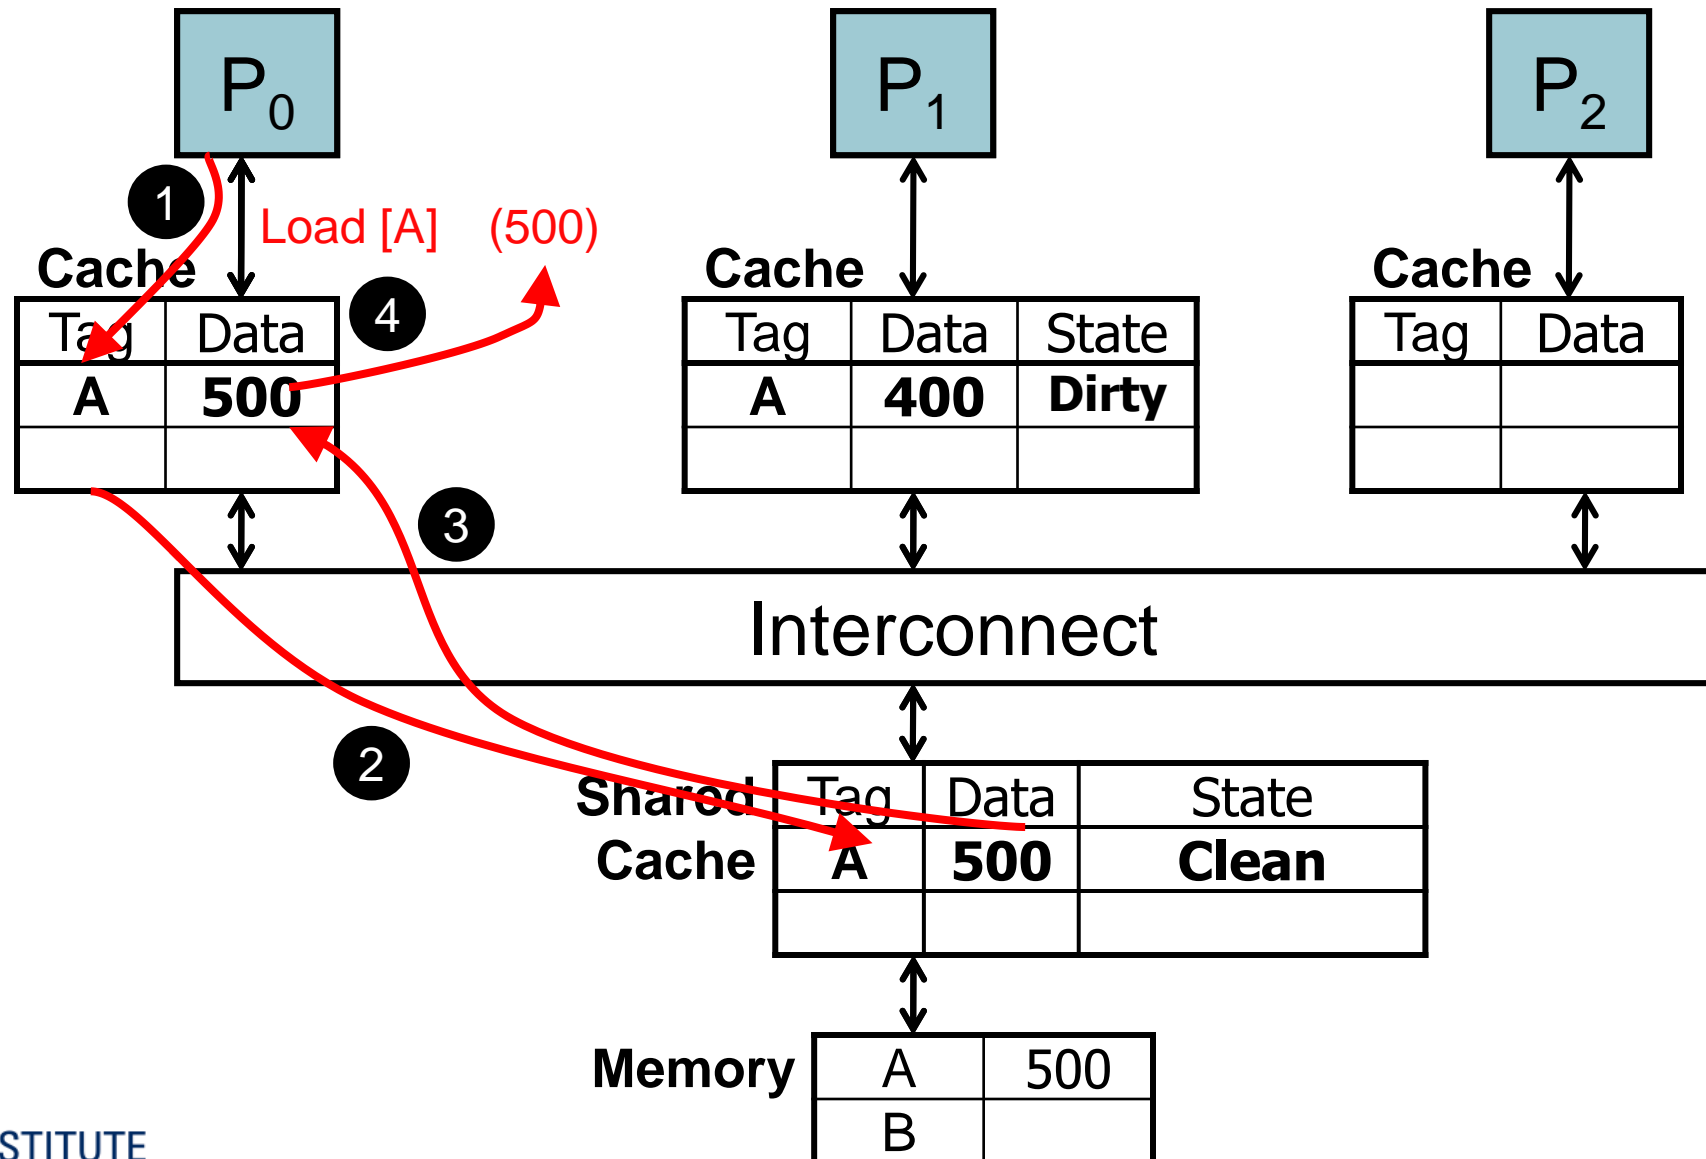
VI Protocol Example



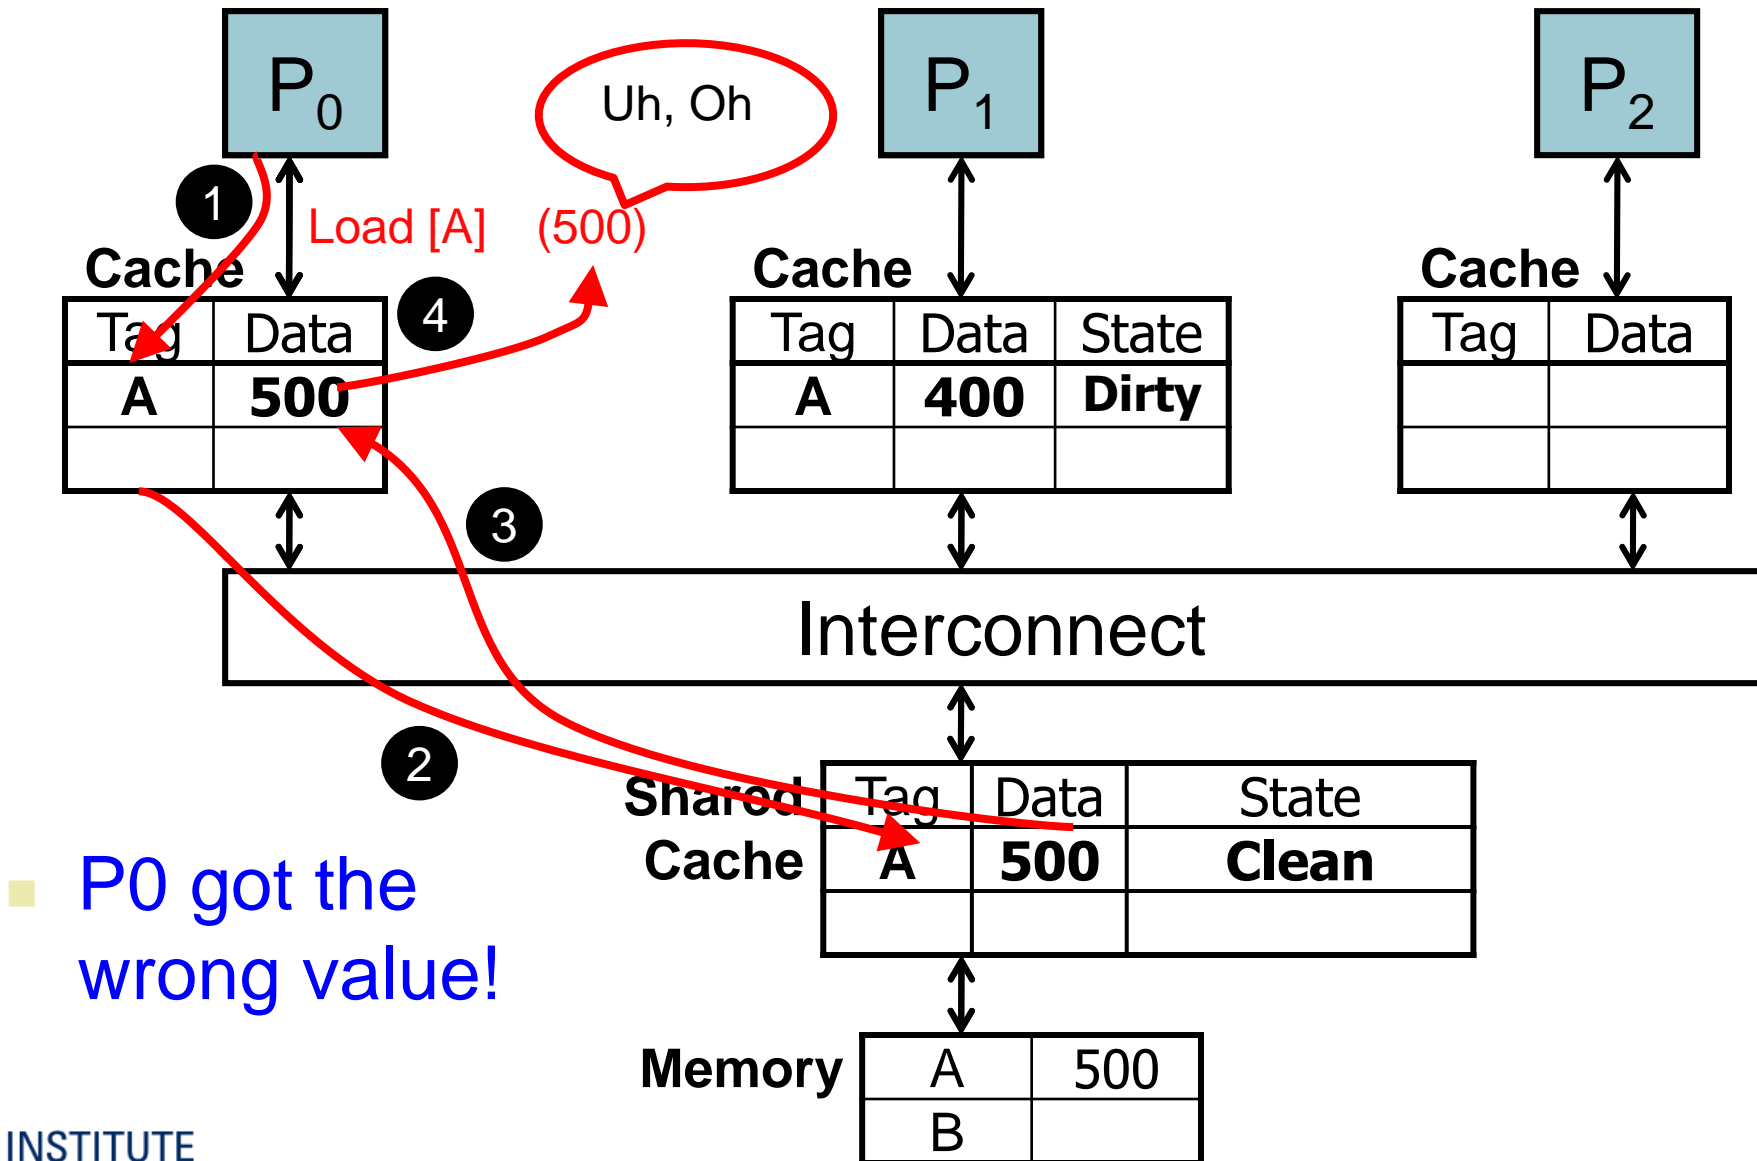
VI Protocol Example



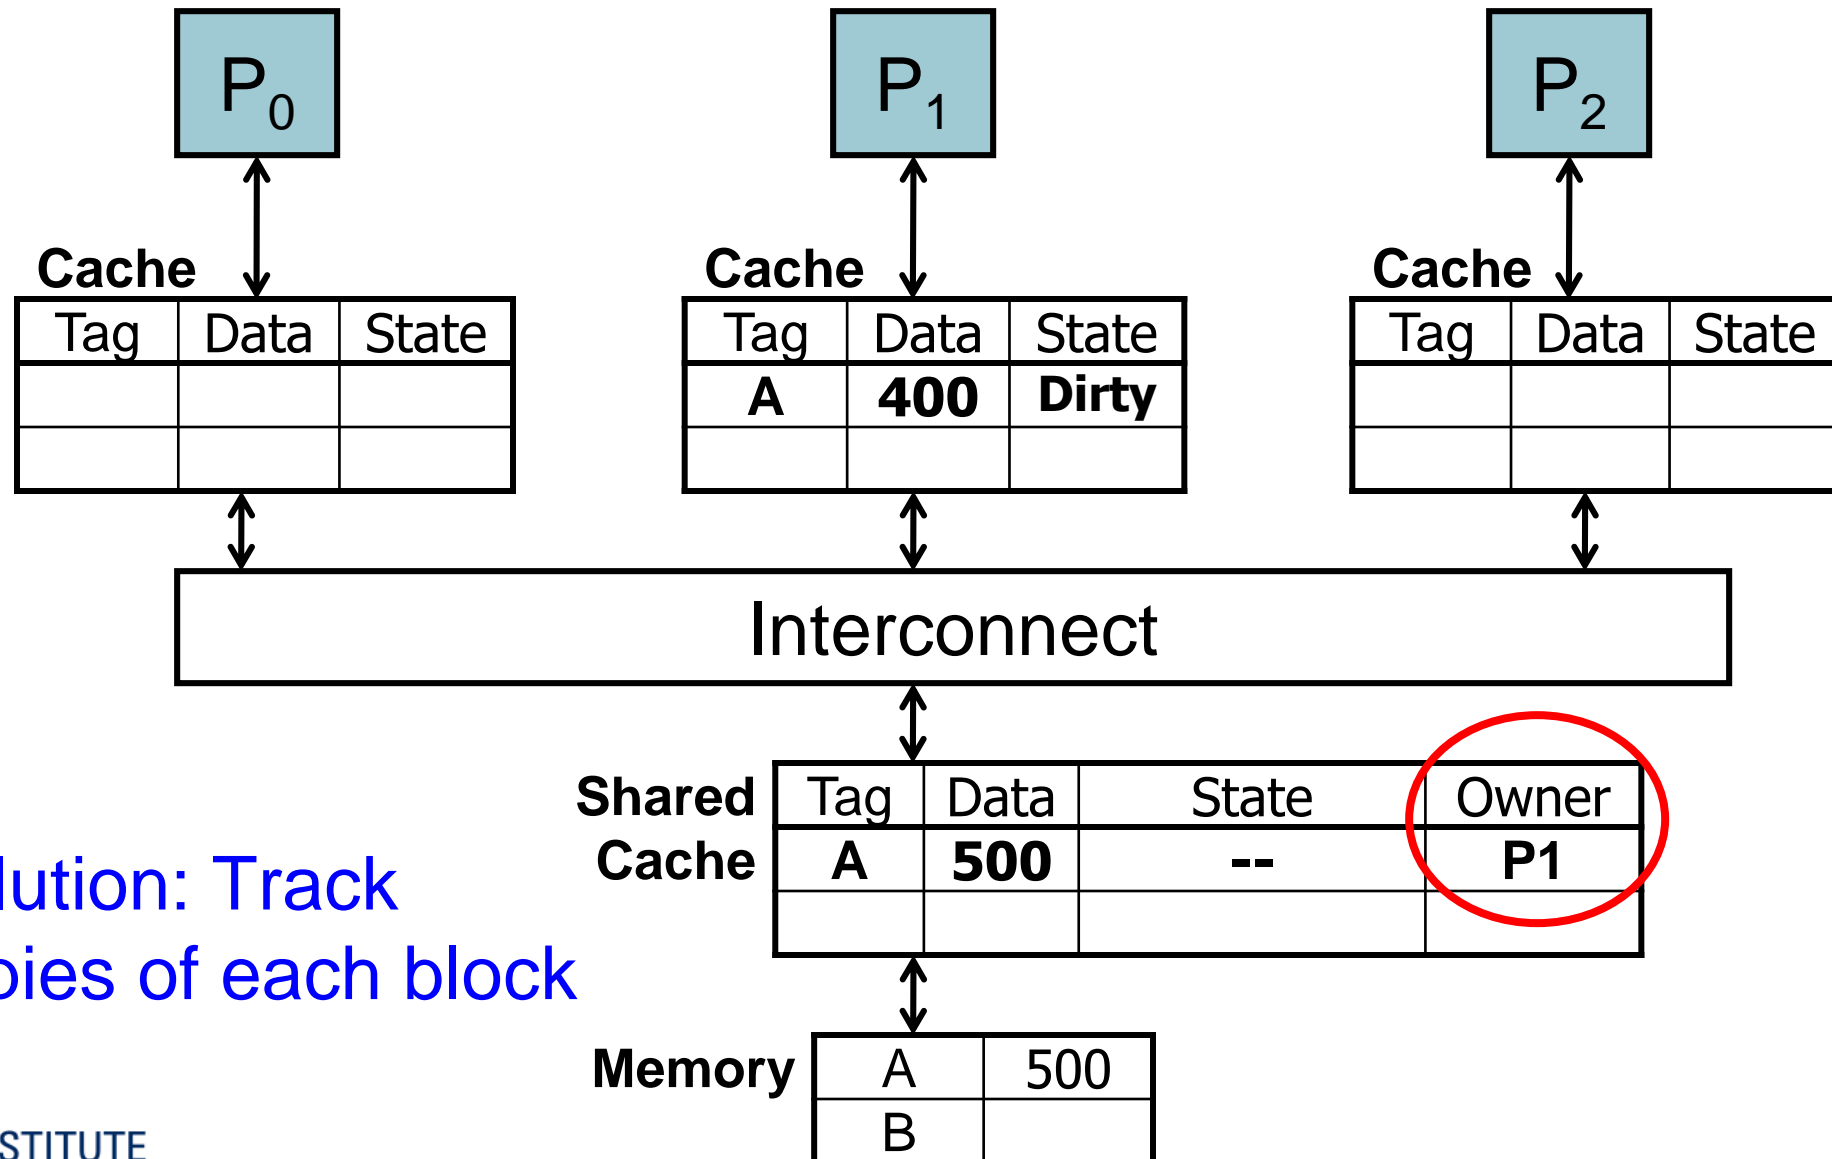
VI Protocol Example



VI Protocol Example

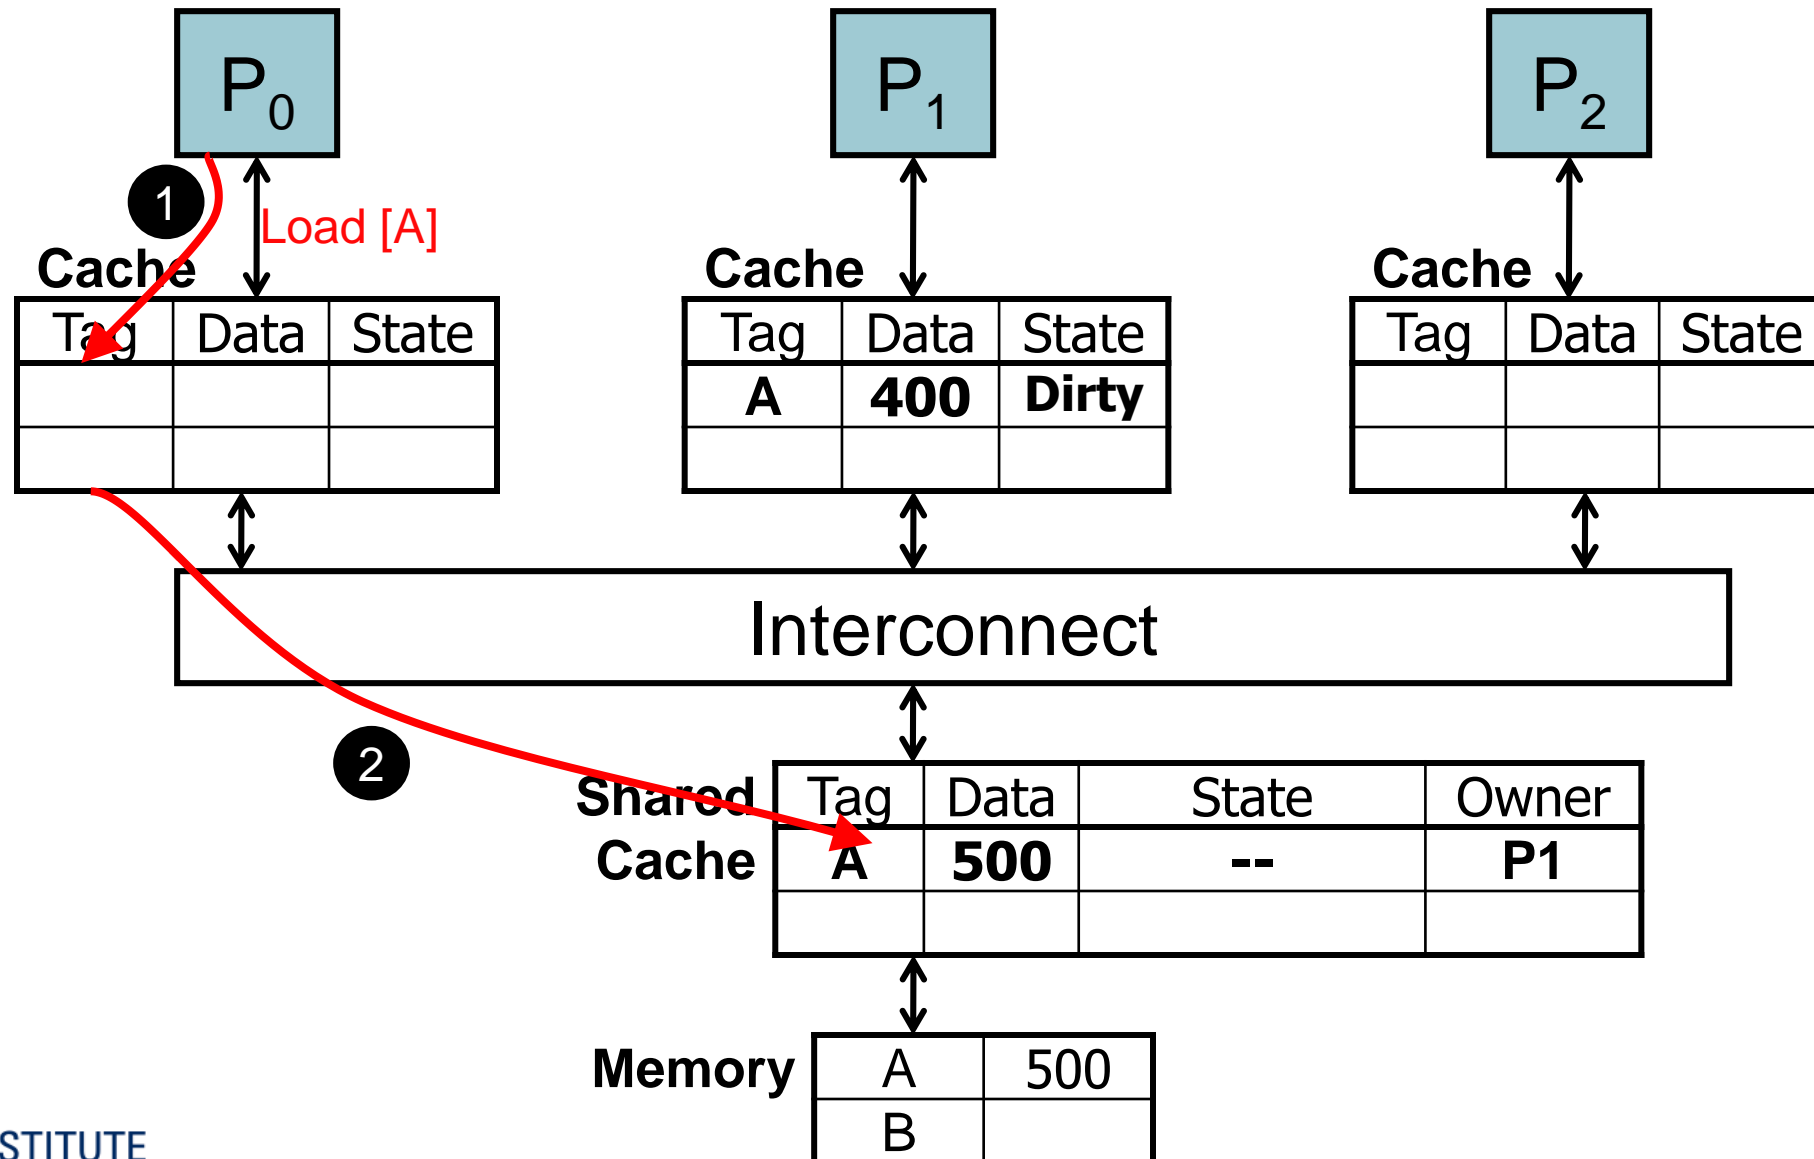


Rewind: Fix Problem by Tracking Sharers

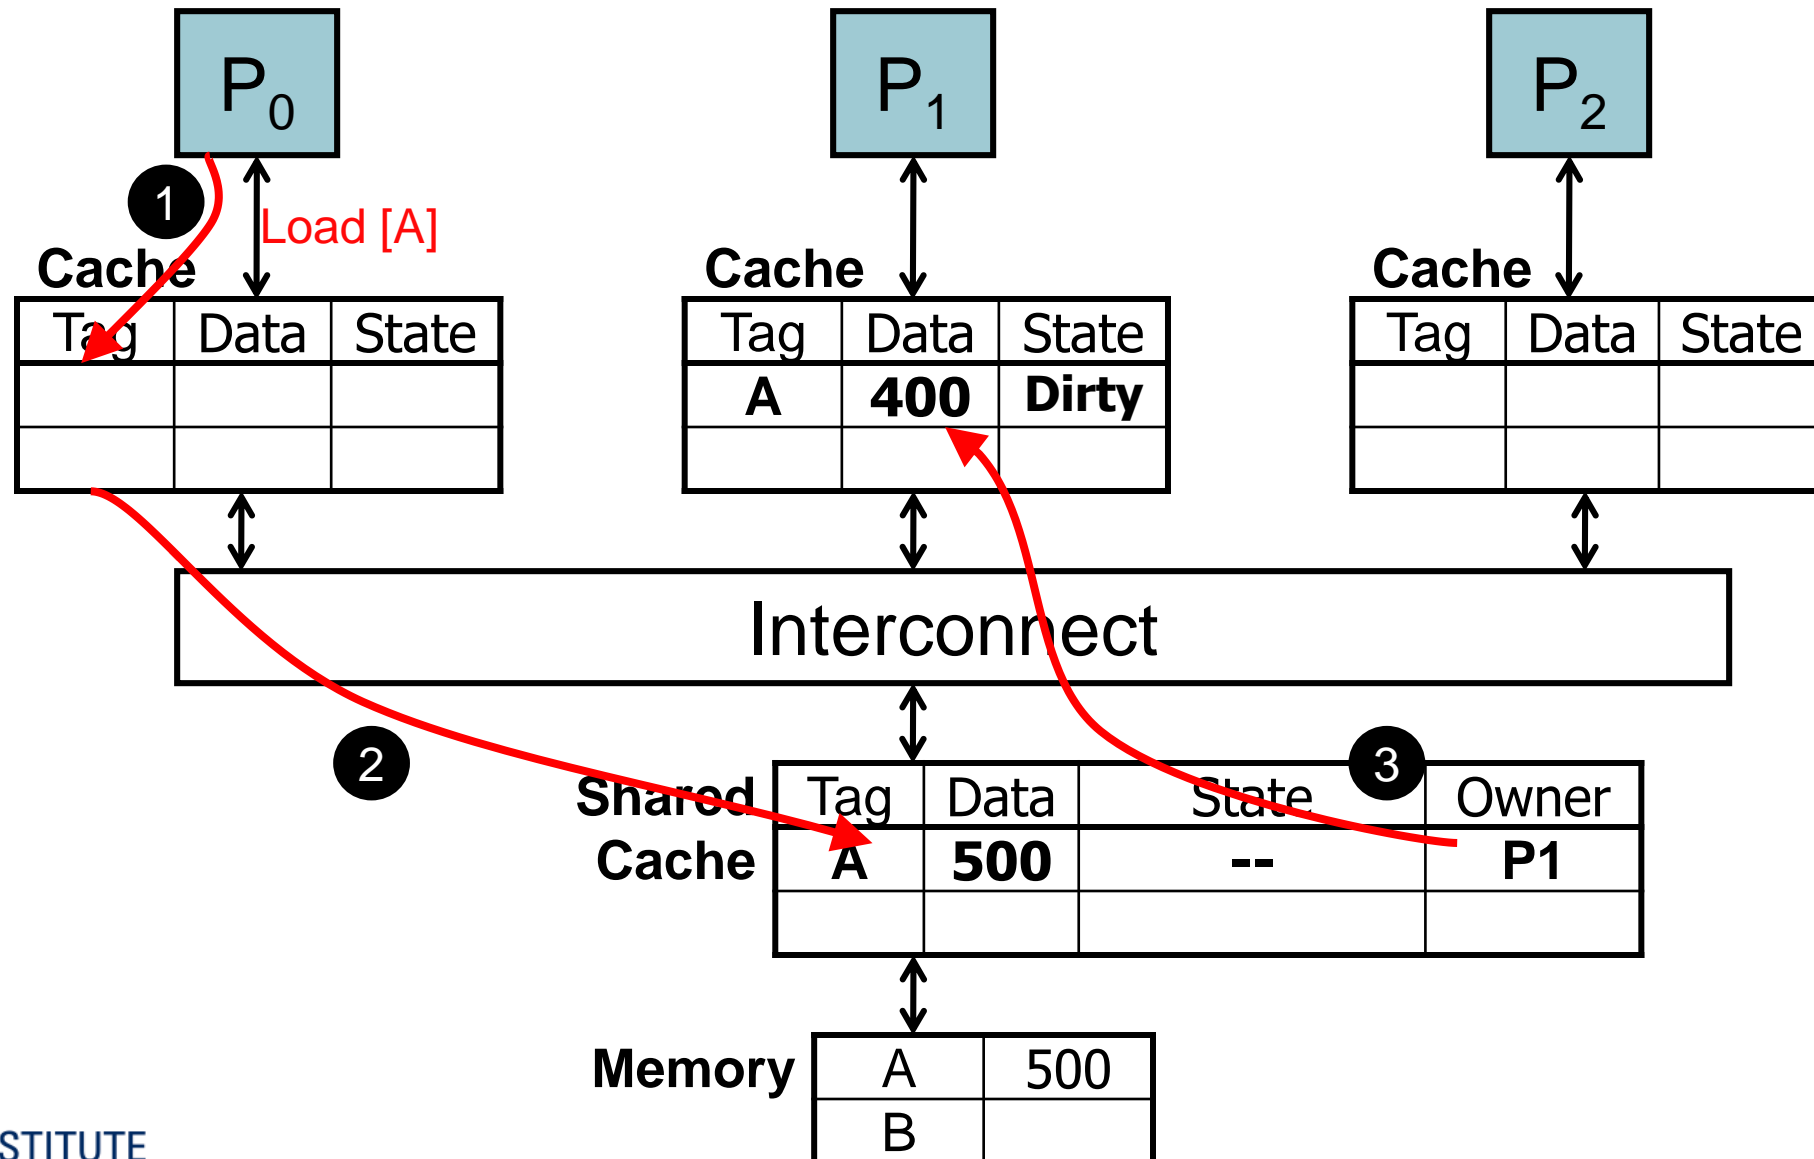


- Solution: Track copies of each block

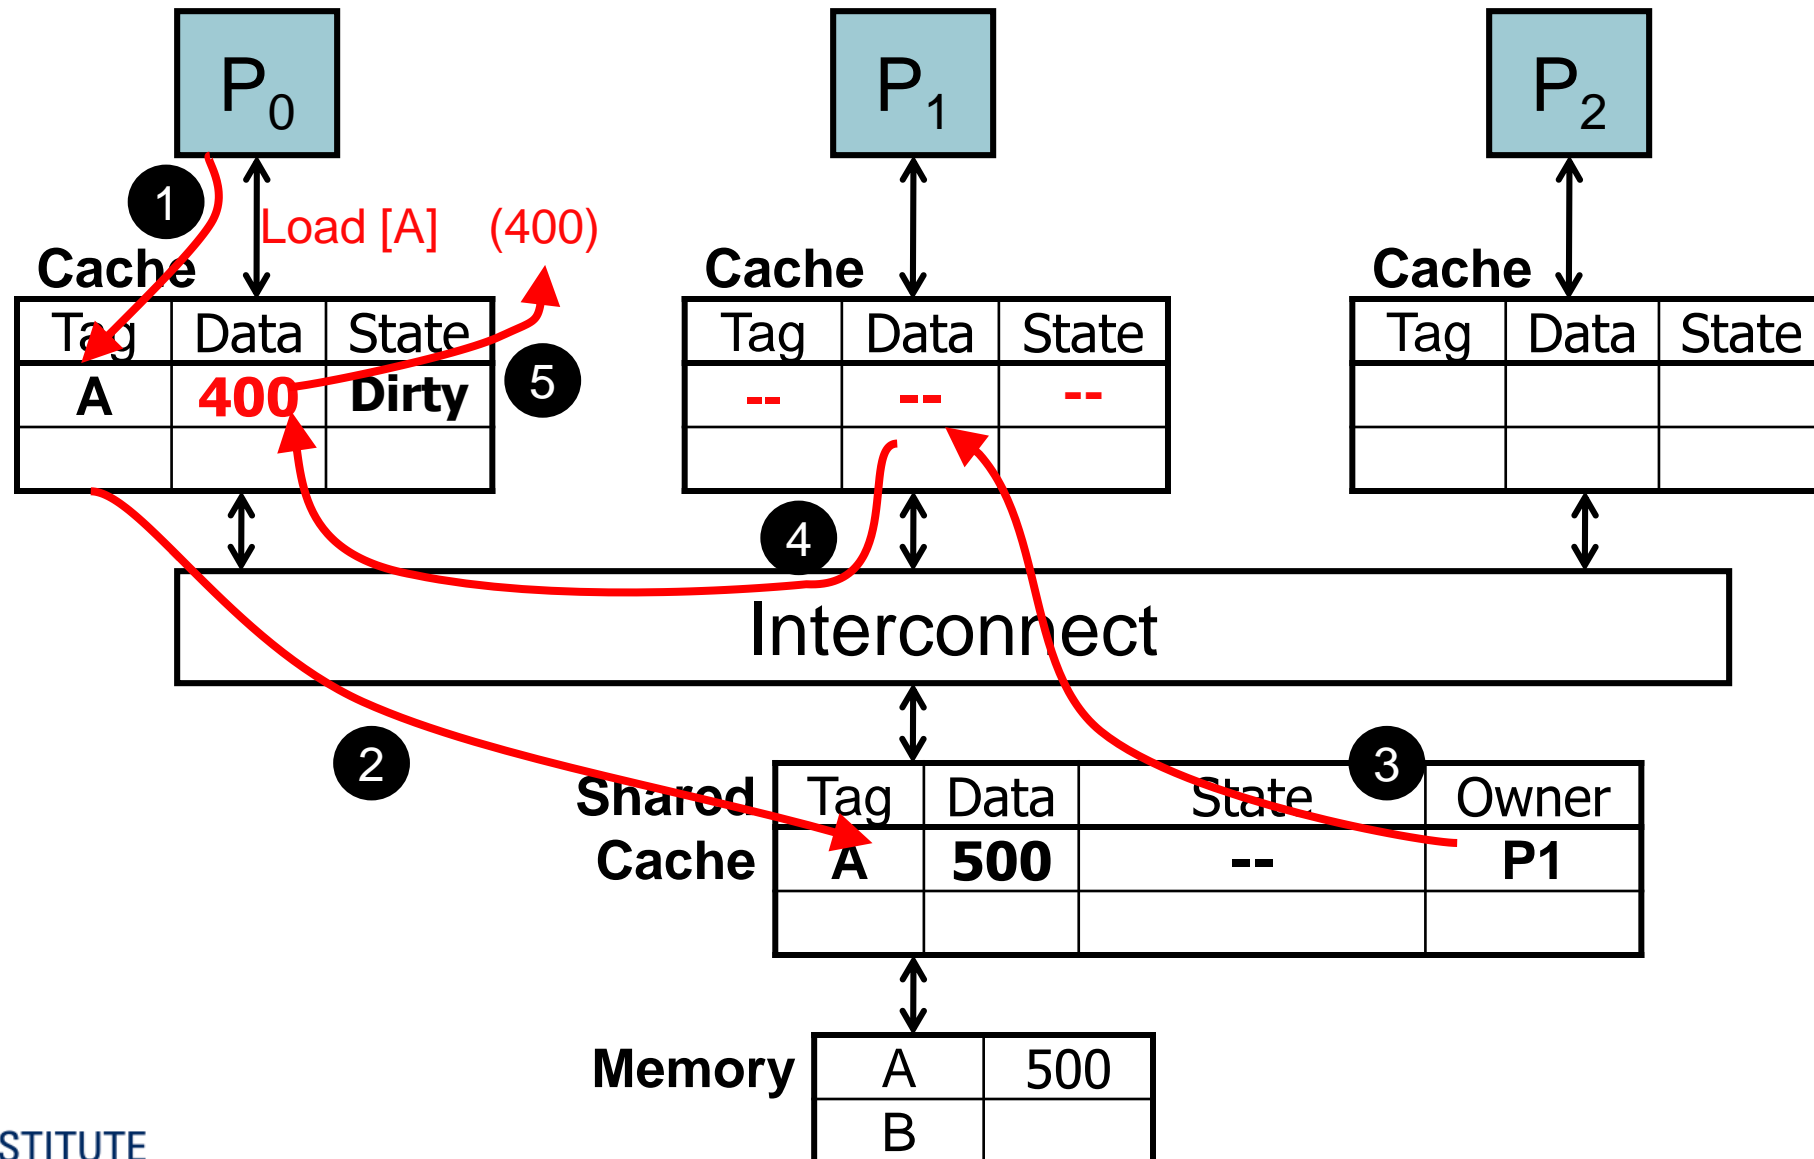
Use Tracking Information to “Invalidate”



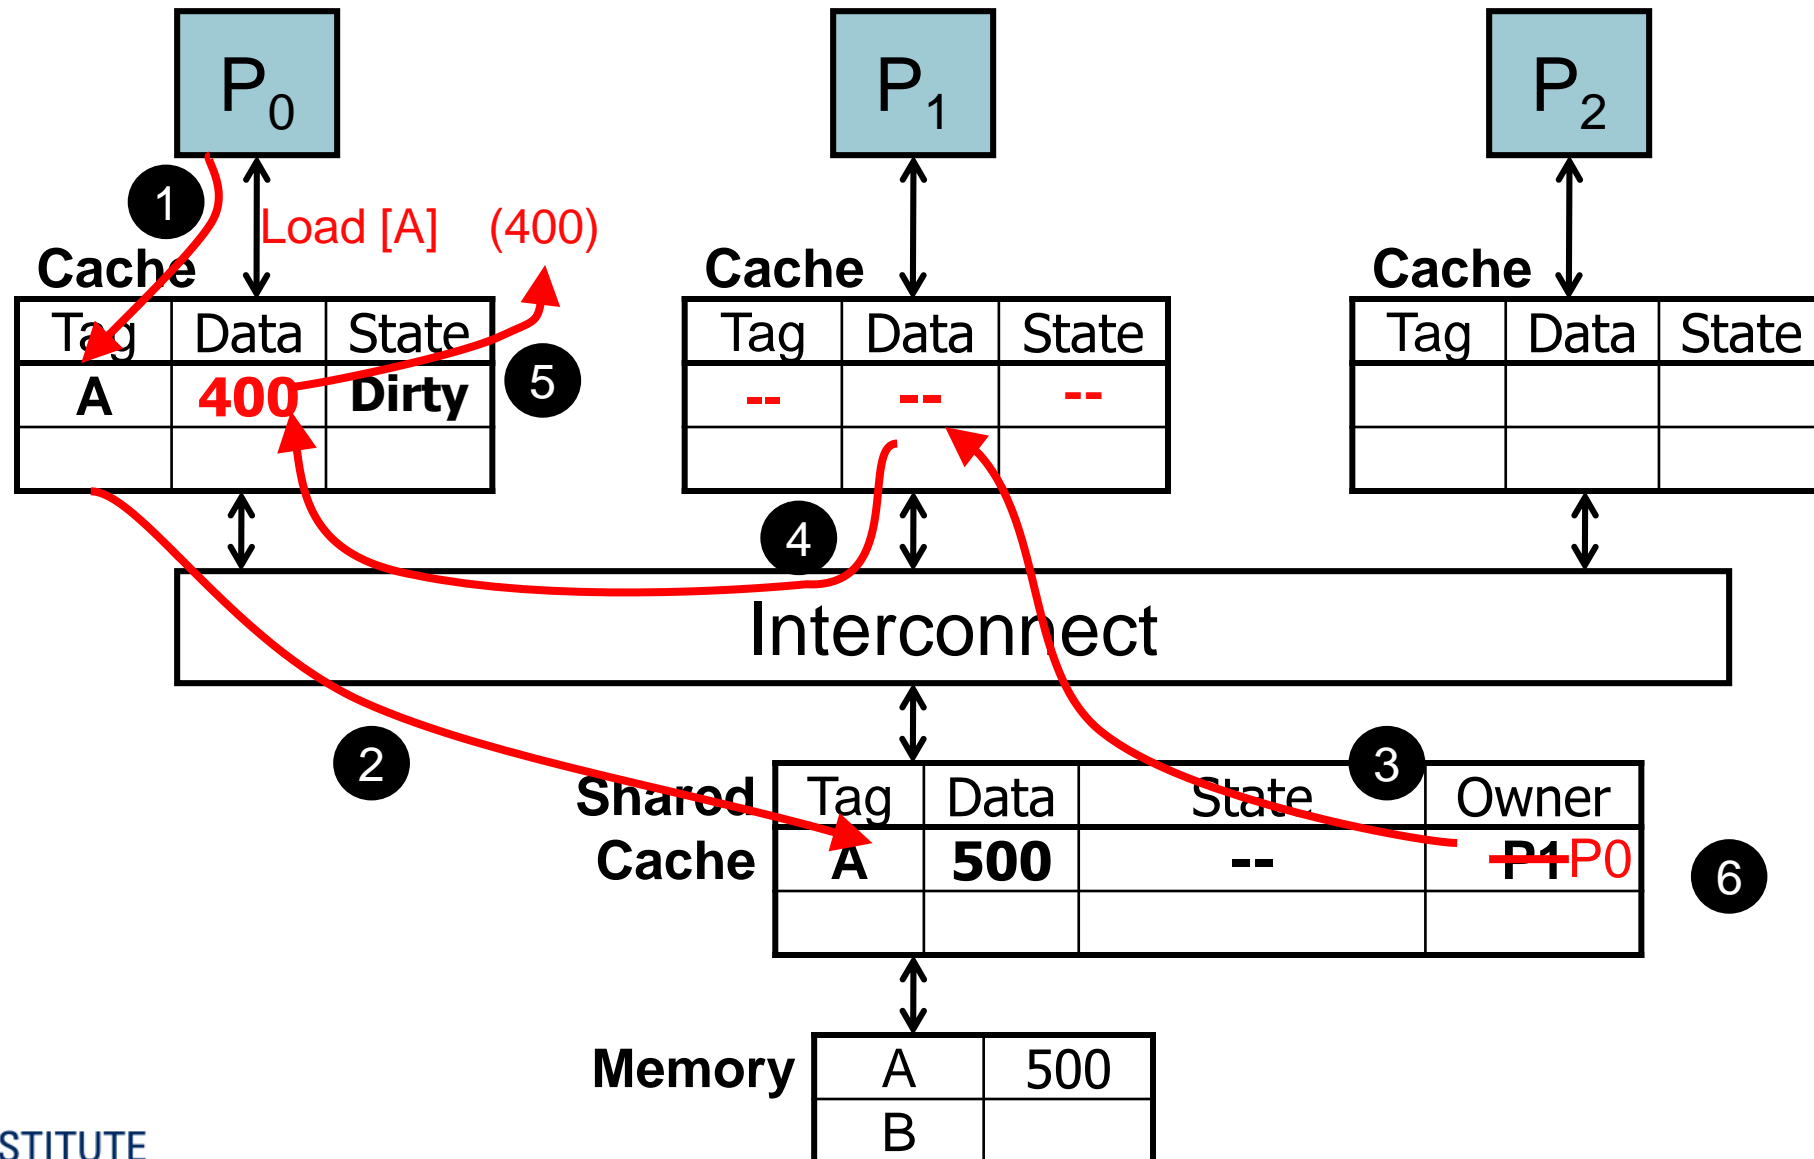
Use Tracking Information to “Invalidate”



Use Tracking Information to “Invalidate”



Use Tracking Information to “Invalidate”



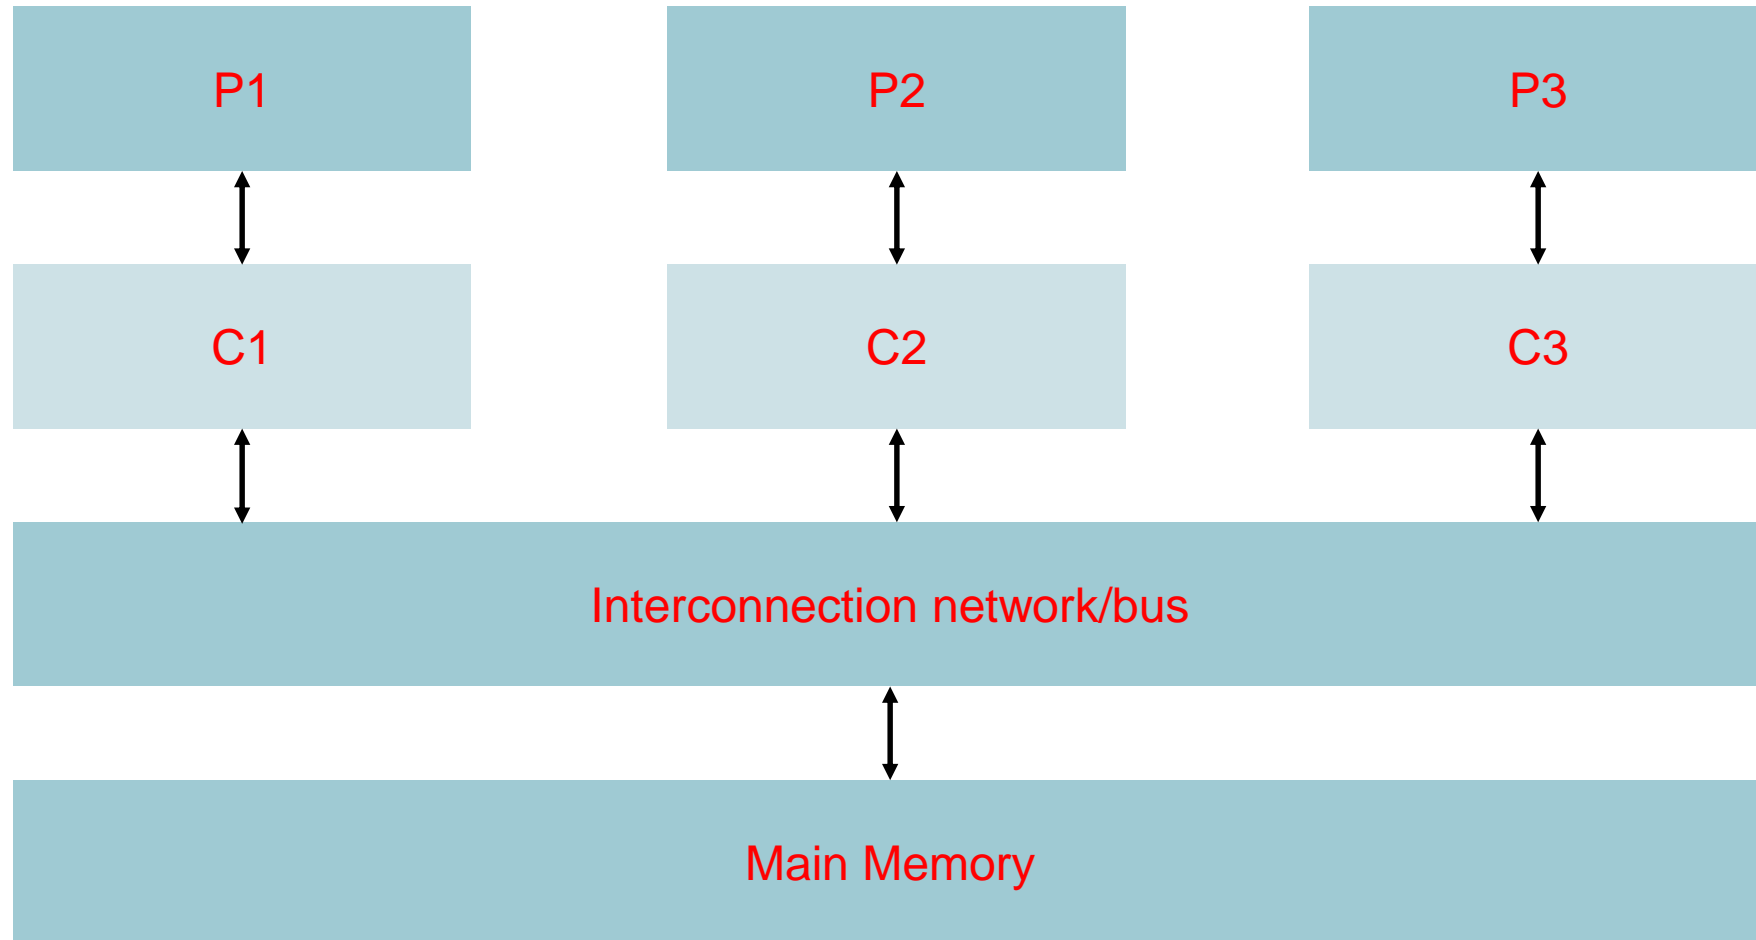
“Valid/Invalid” Cache Coherence Summary

- To enforce the shared memory invariant...
 - “Loads read the value written by the most recent store”
- Enforce the invariant...
 - “**At most one valid copy of the block**”
 - Simplest form is a **two-state “valid/invalid” protocol**
 - If a core wants a copy, must find and “invalidate” it
- On a cache miss, how is the valid copy found?
 - Option #1 “**Snooping**”: broadcast to all, whoever has it responds
 - Option #2: “**Directory**”: track sharers with separate structure
- **Problem**: multiple copies can’t exist, even if read-only
 - Consider mostly-read data structures, instructions, etc.

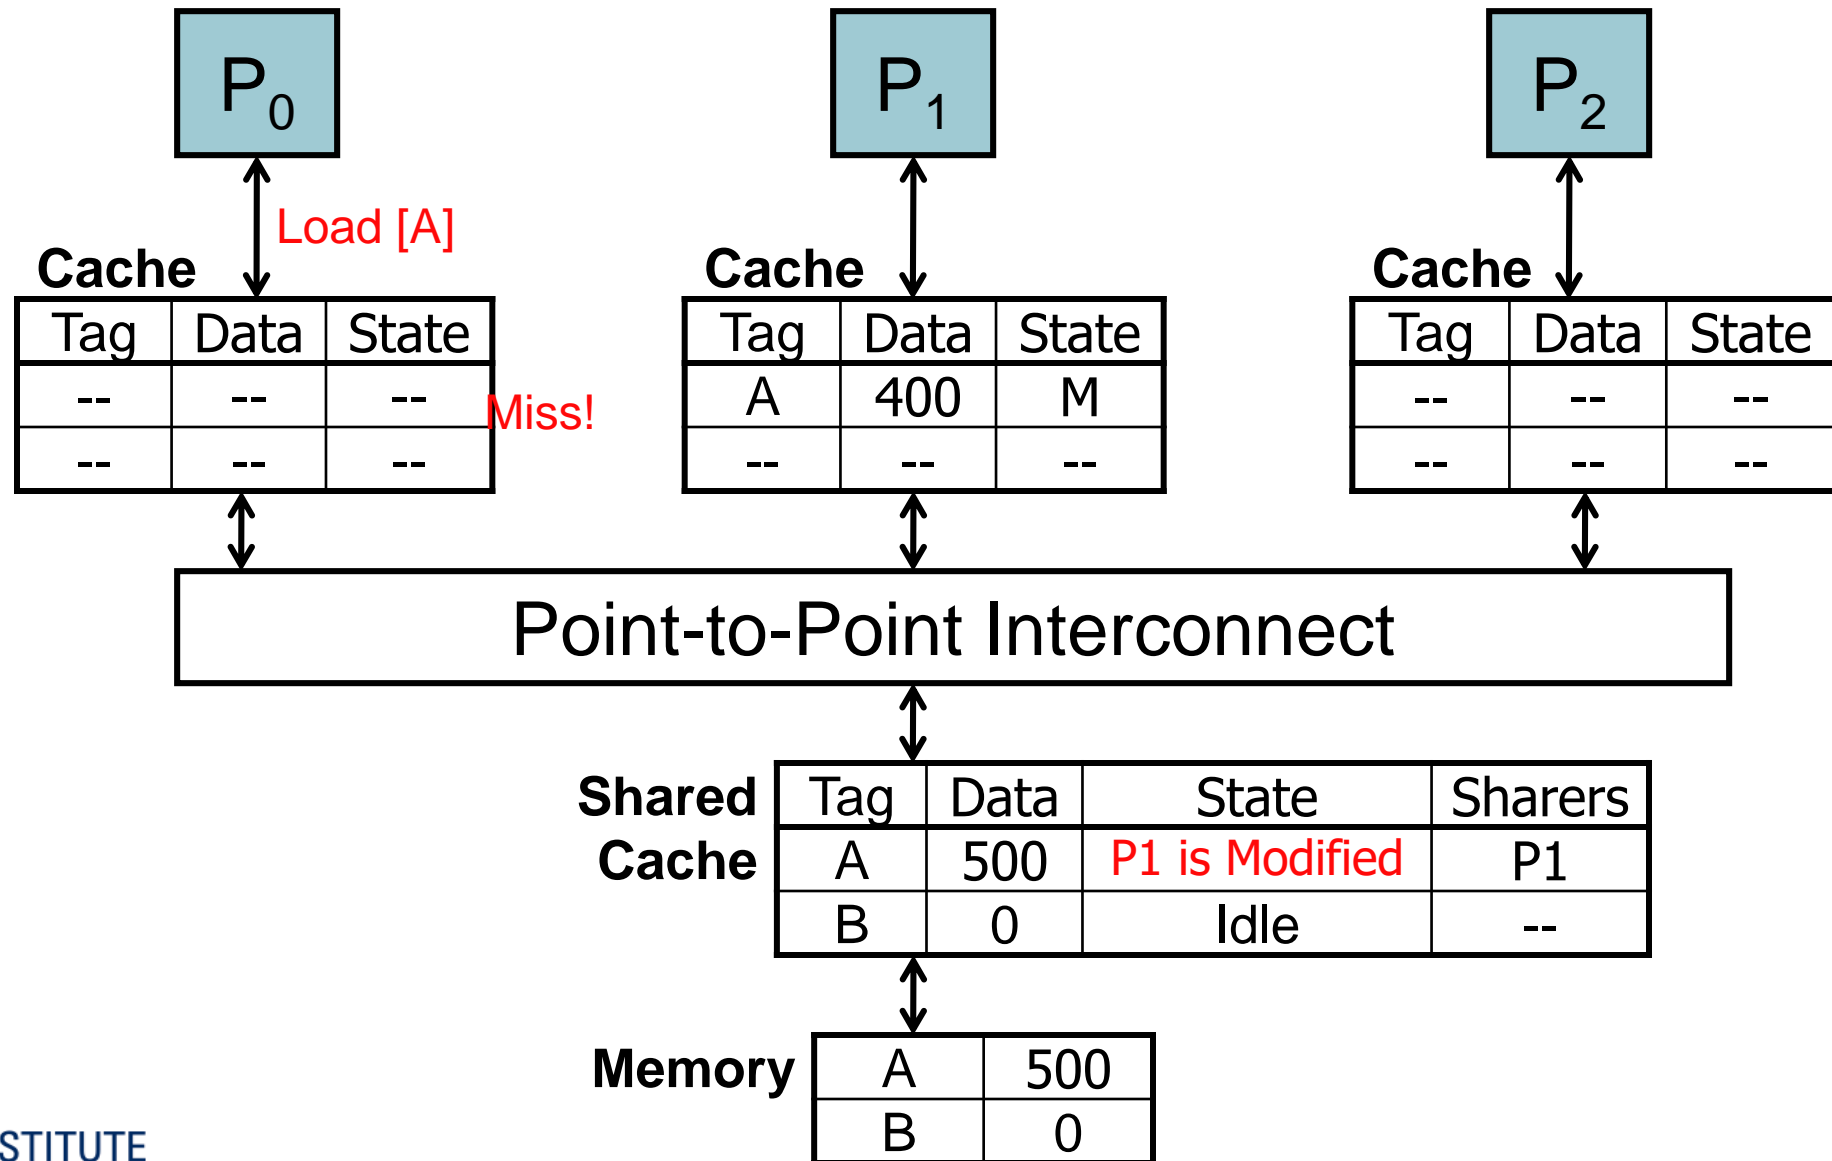
Modified-Shared-Invalid (MSI) Protocol

- Allows for many readers...
- Three states tracked per-block at each cache
 - **Modified** – cache has the only copy; writable; dirty (read/write permission)
 - Dirty == memory is out of date
 - **Shared** – cache has a read-only copy; clean (read-only permission)
 - Clean == memory is up to date
 - **Invalid** – cache does not have a copy (blocked, no permission)
- Three processor actions
 - Load, Store, Evict
- Five bus messages
 - BusRd, BusRdX, BusInv, BusWB, BusReply
 - Could combine some of these

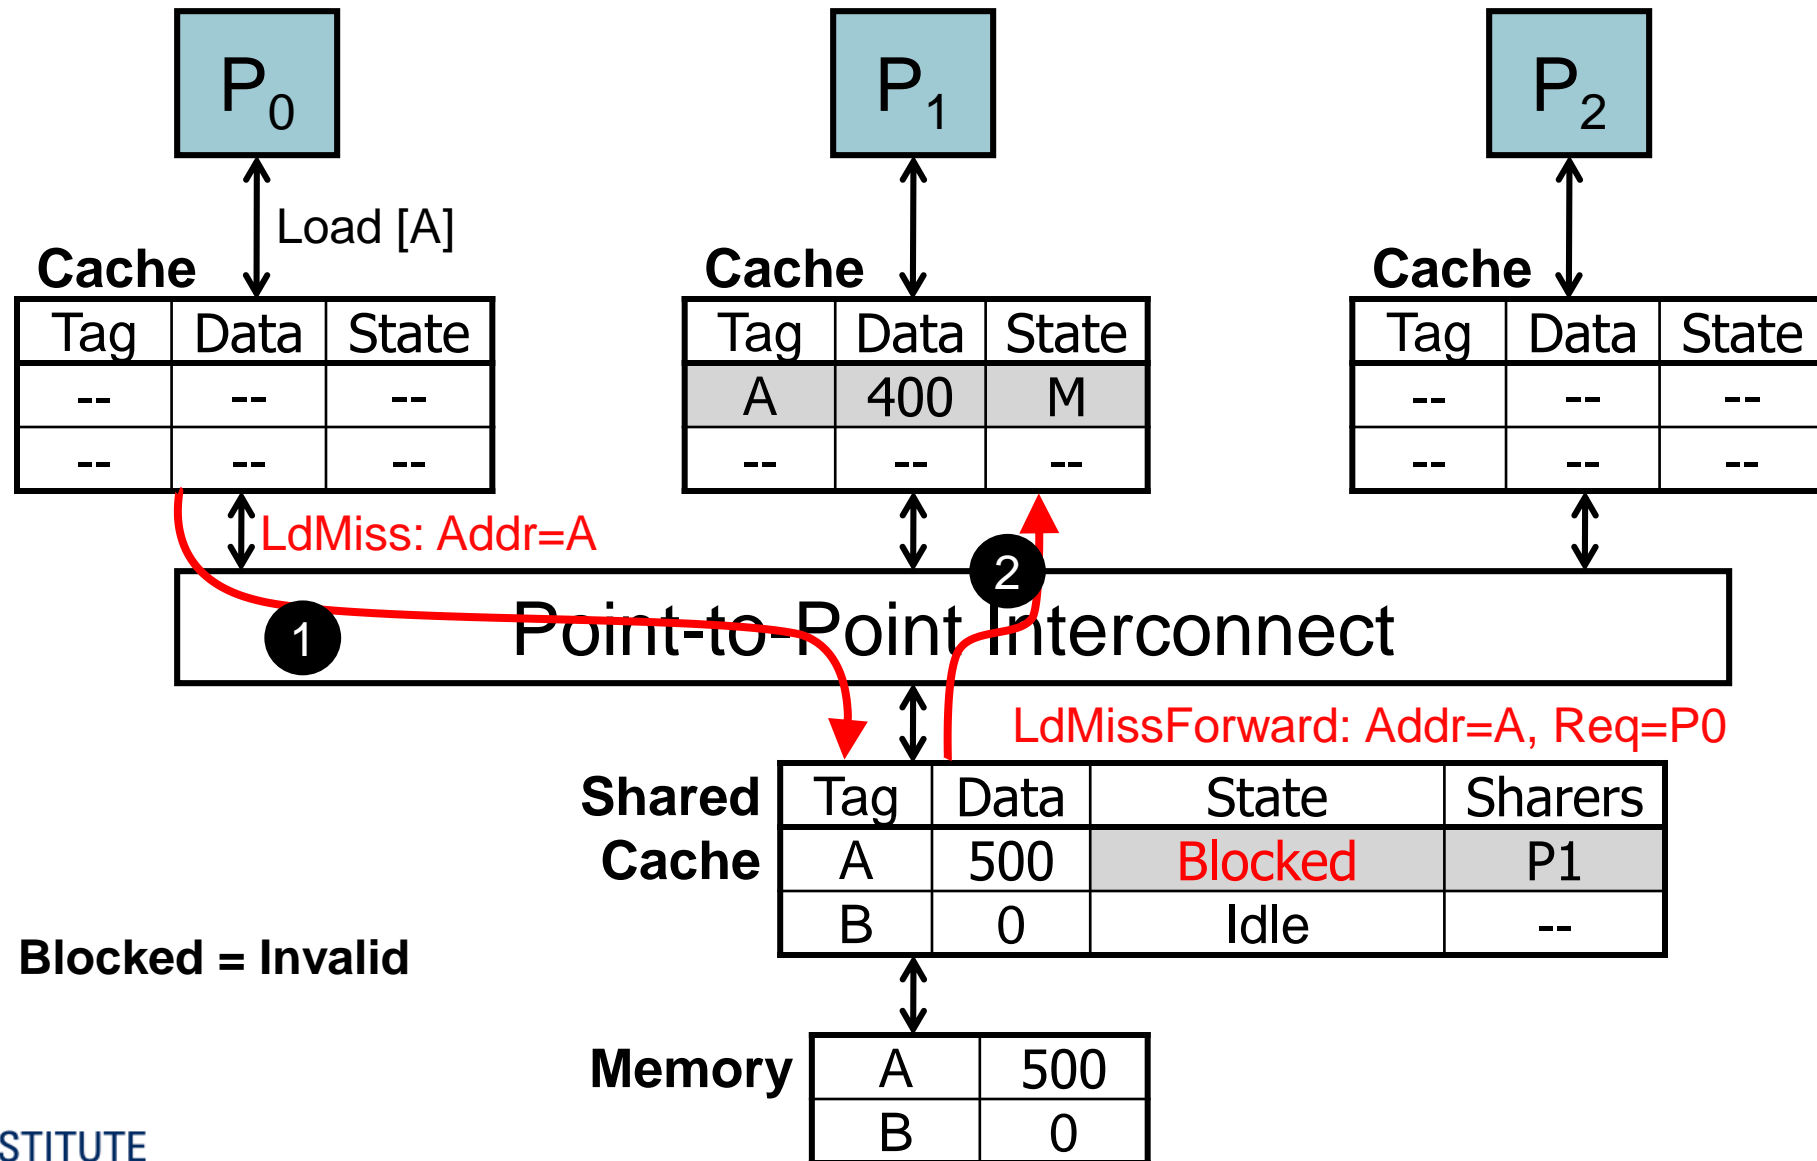
MSI (Snooping Based)



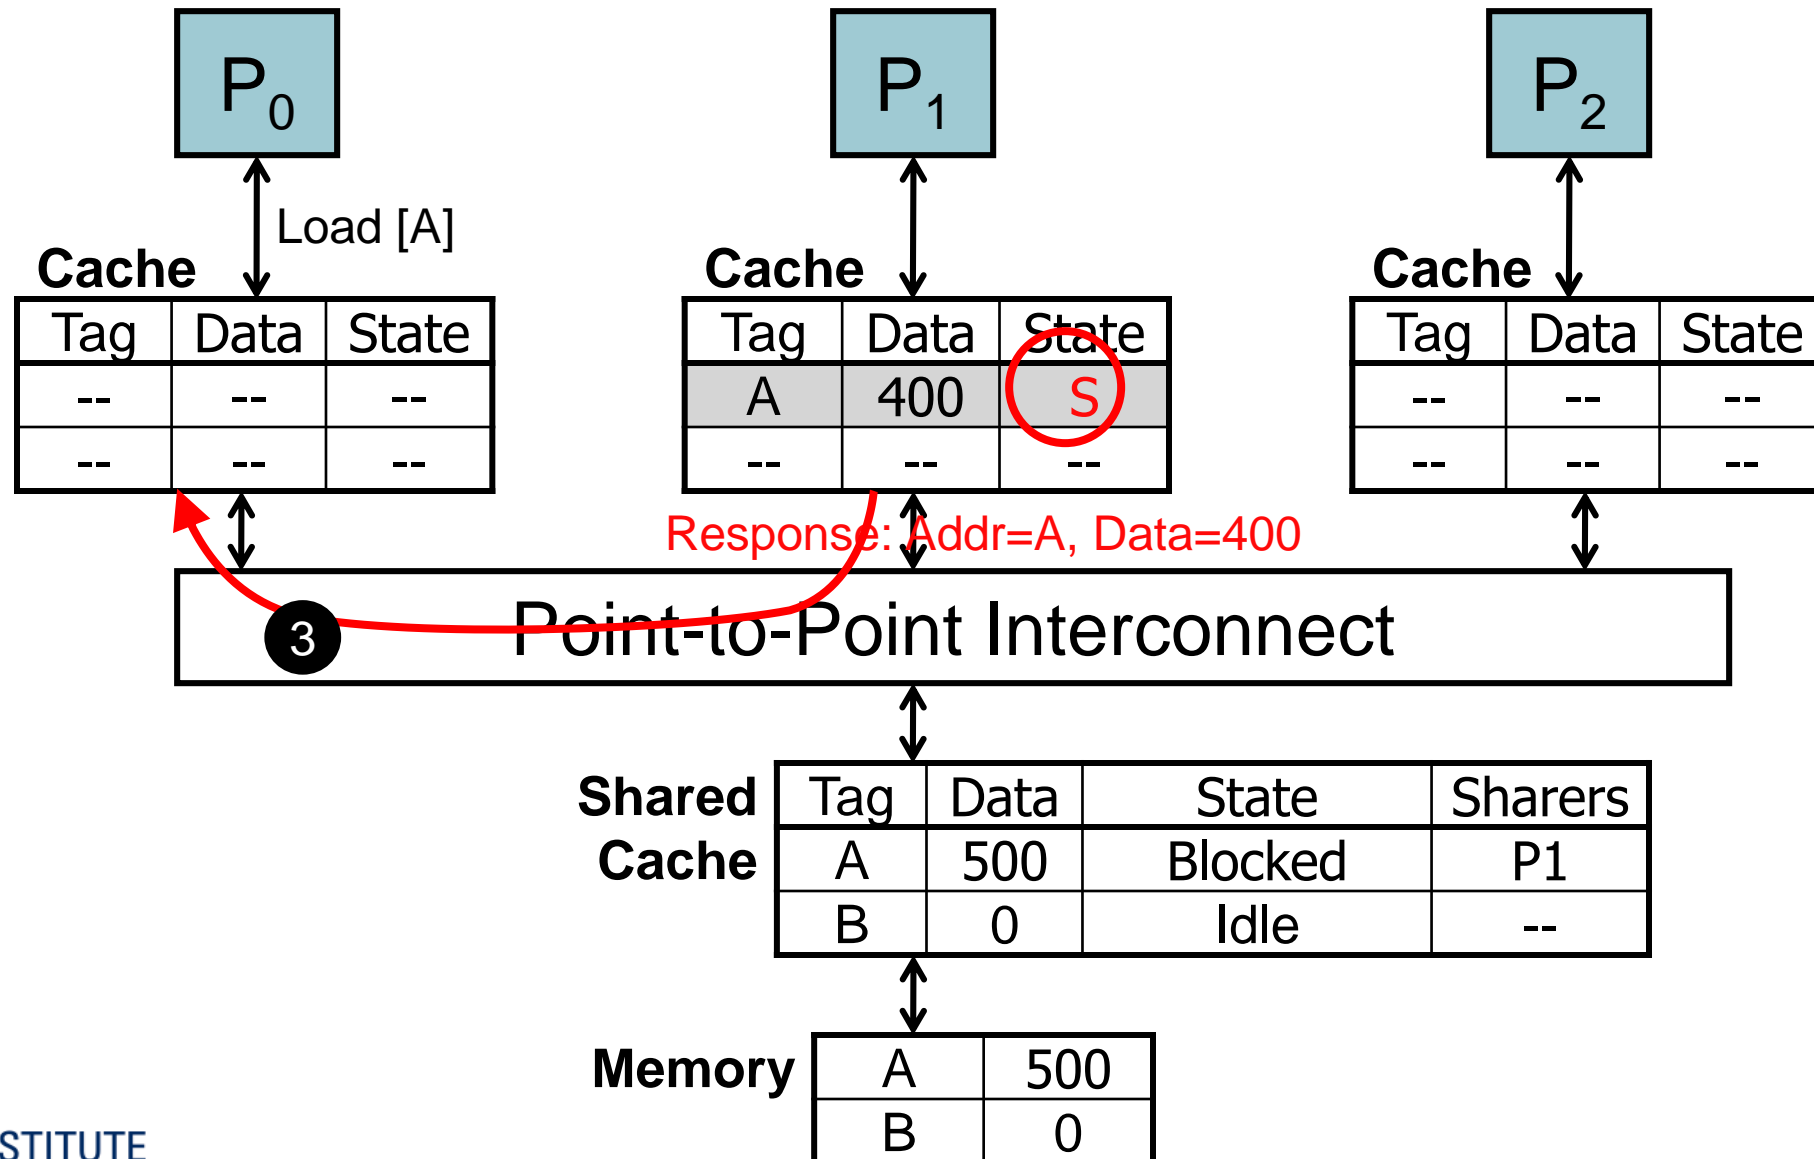
MSI Coherence Example: Step #1



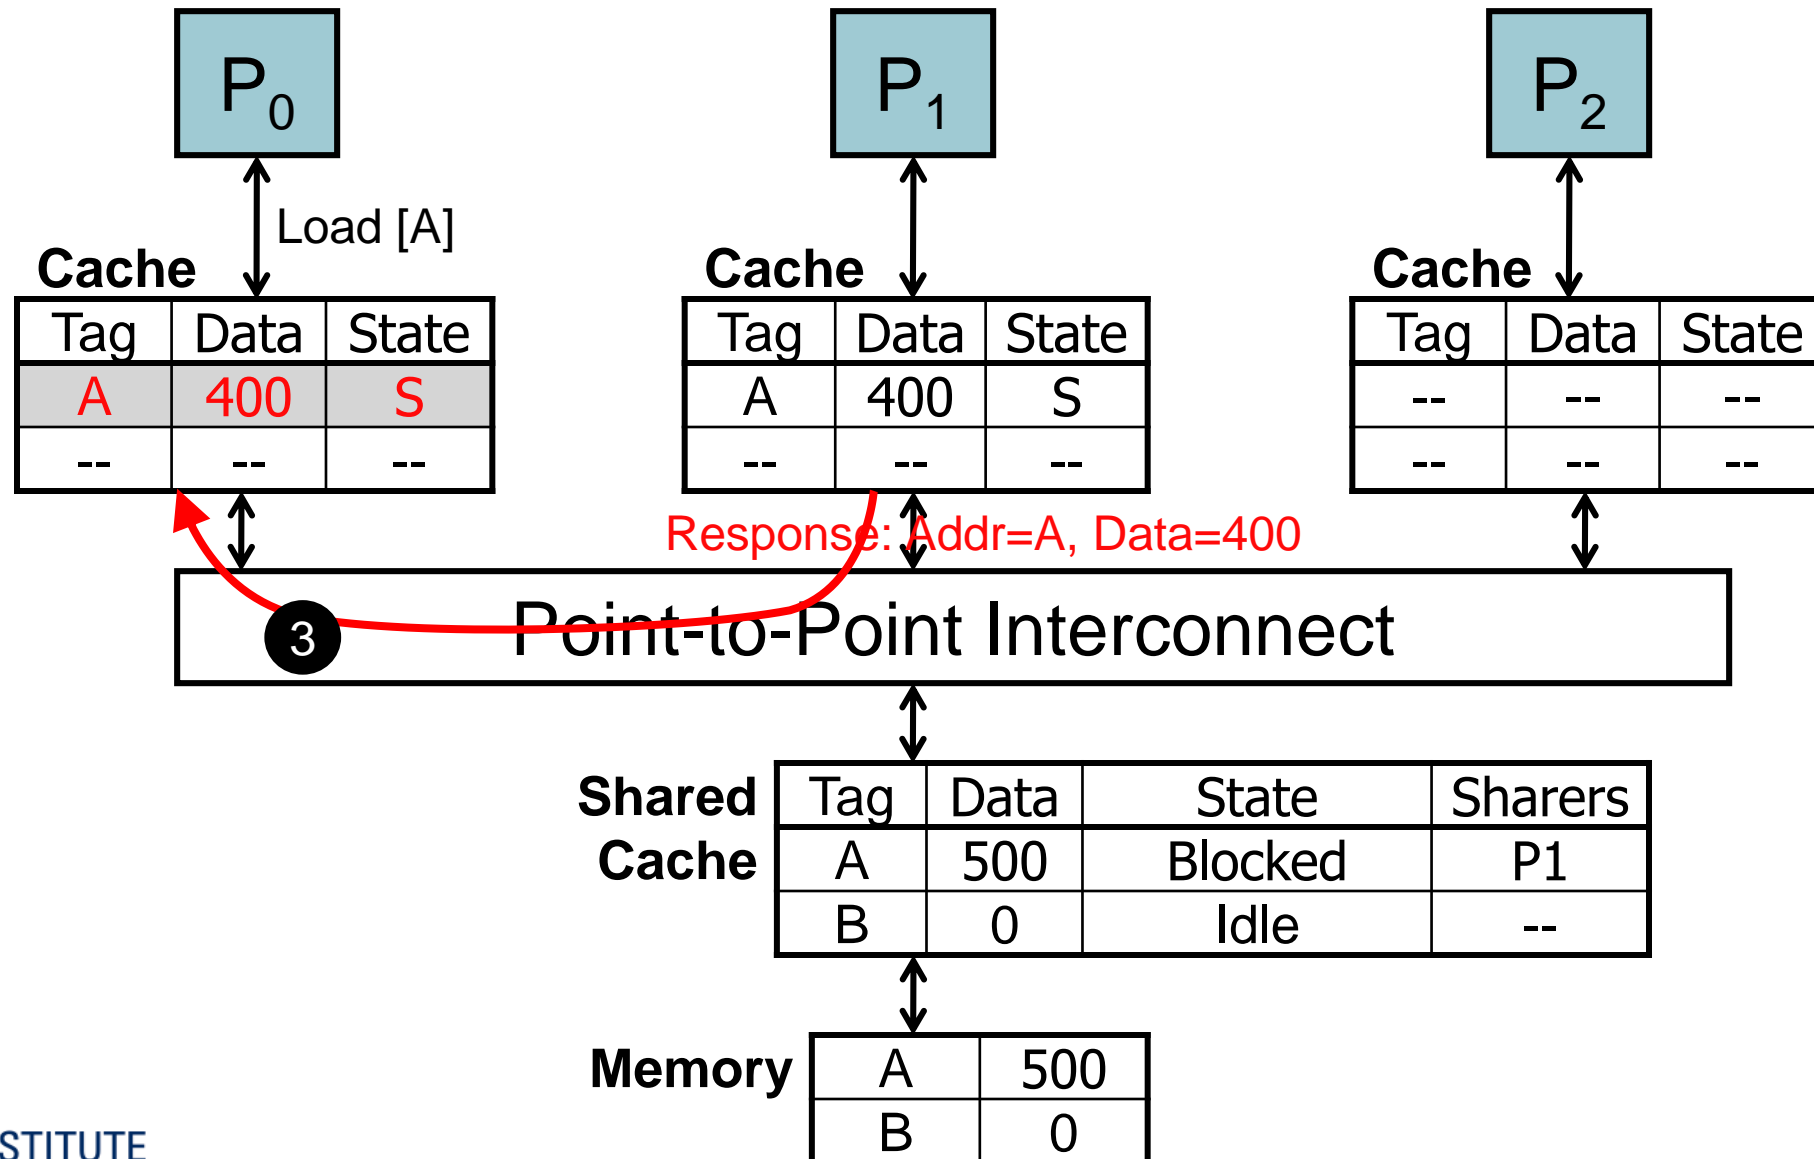
MSI Coherence Example: Step #2



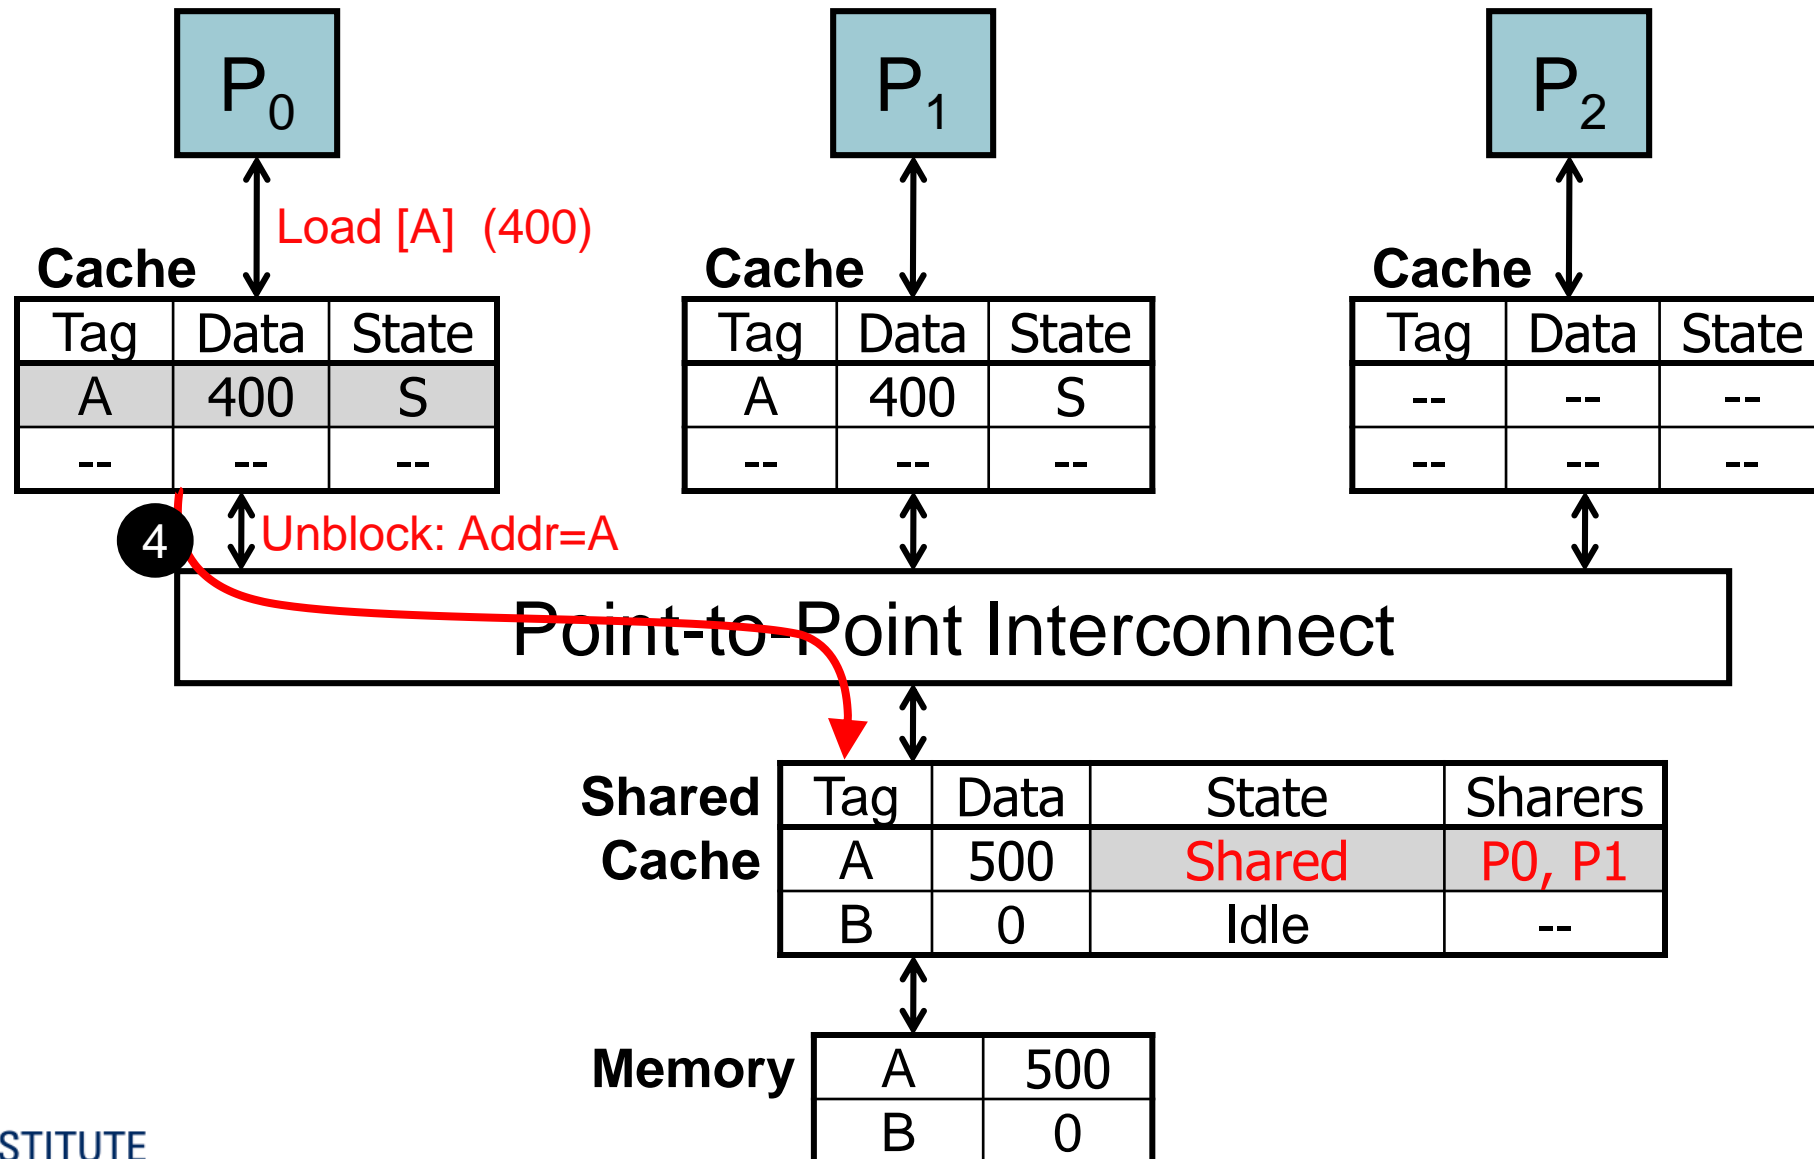
MSI Coherence Example: Step #3



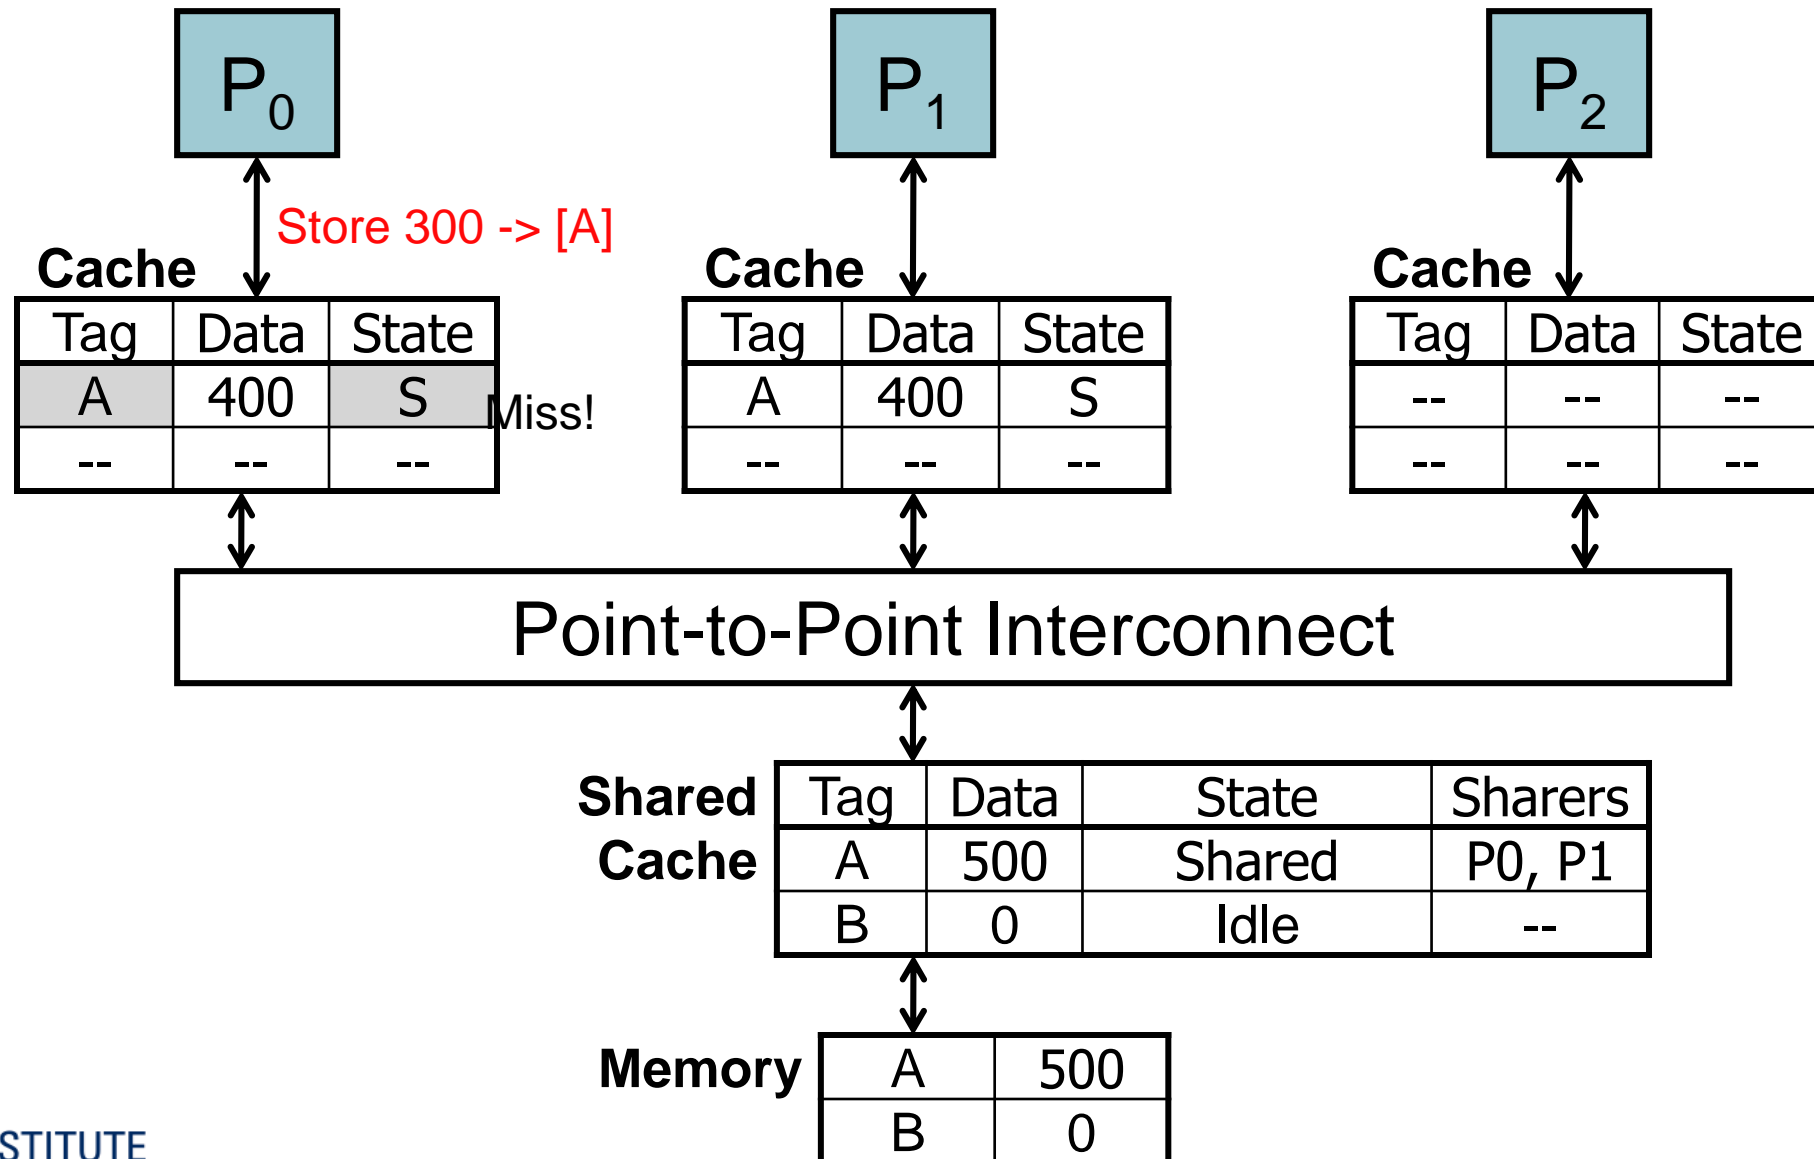
MSI Coherence Example: Step #4



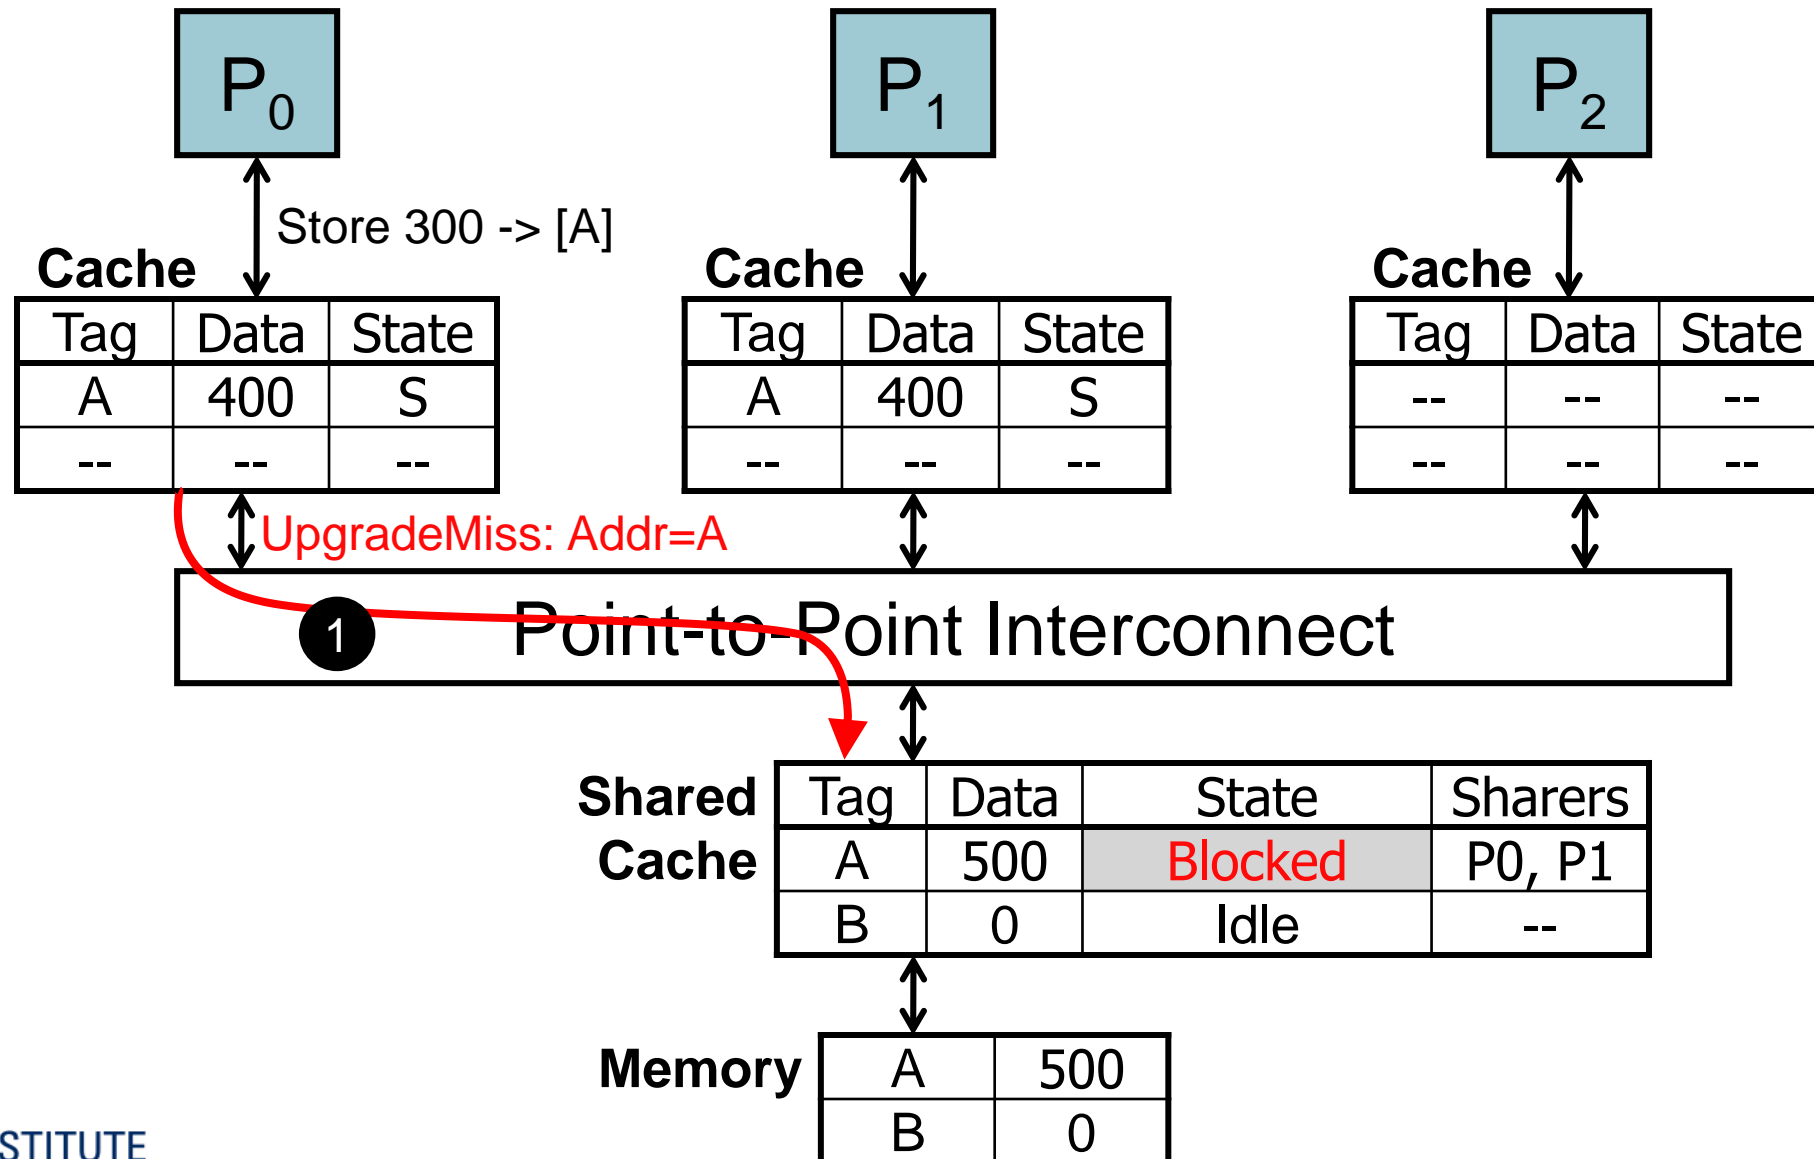
MSI Coherence Example: Step #5



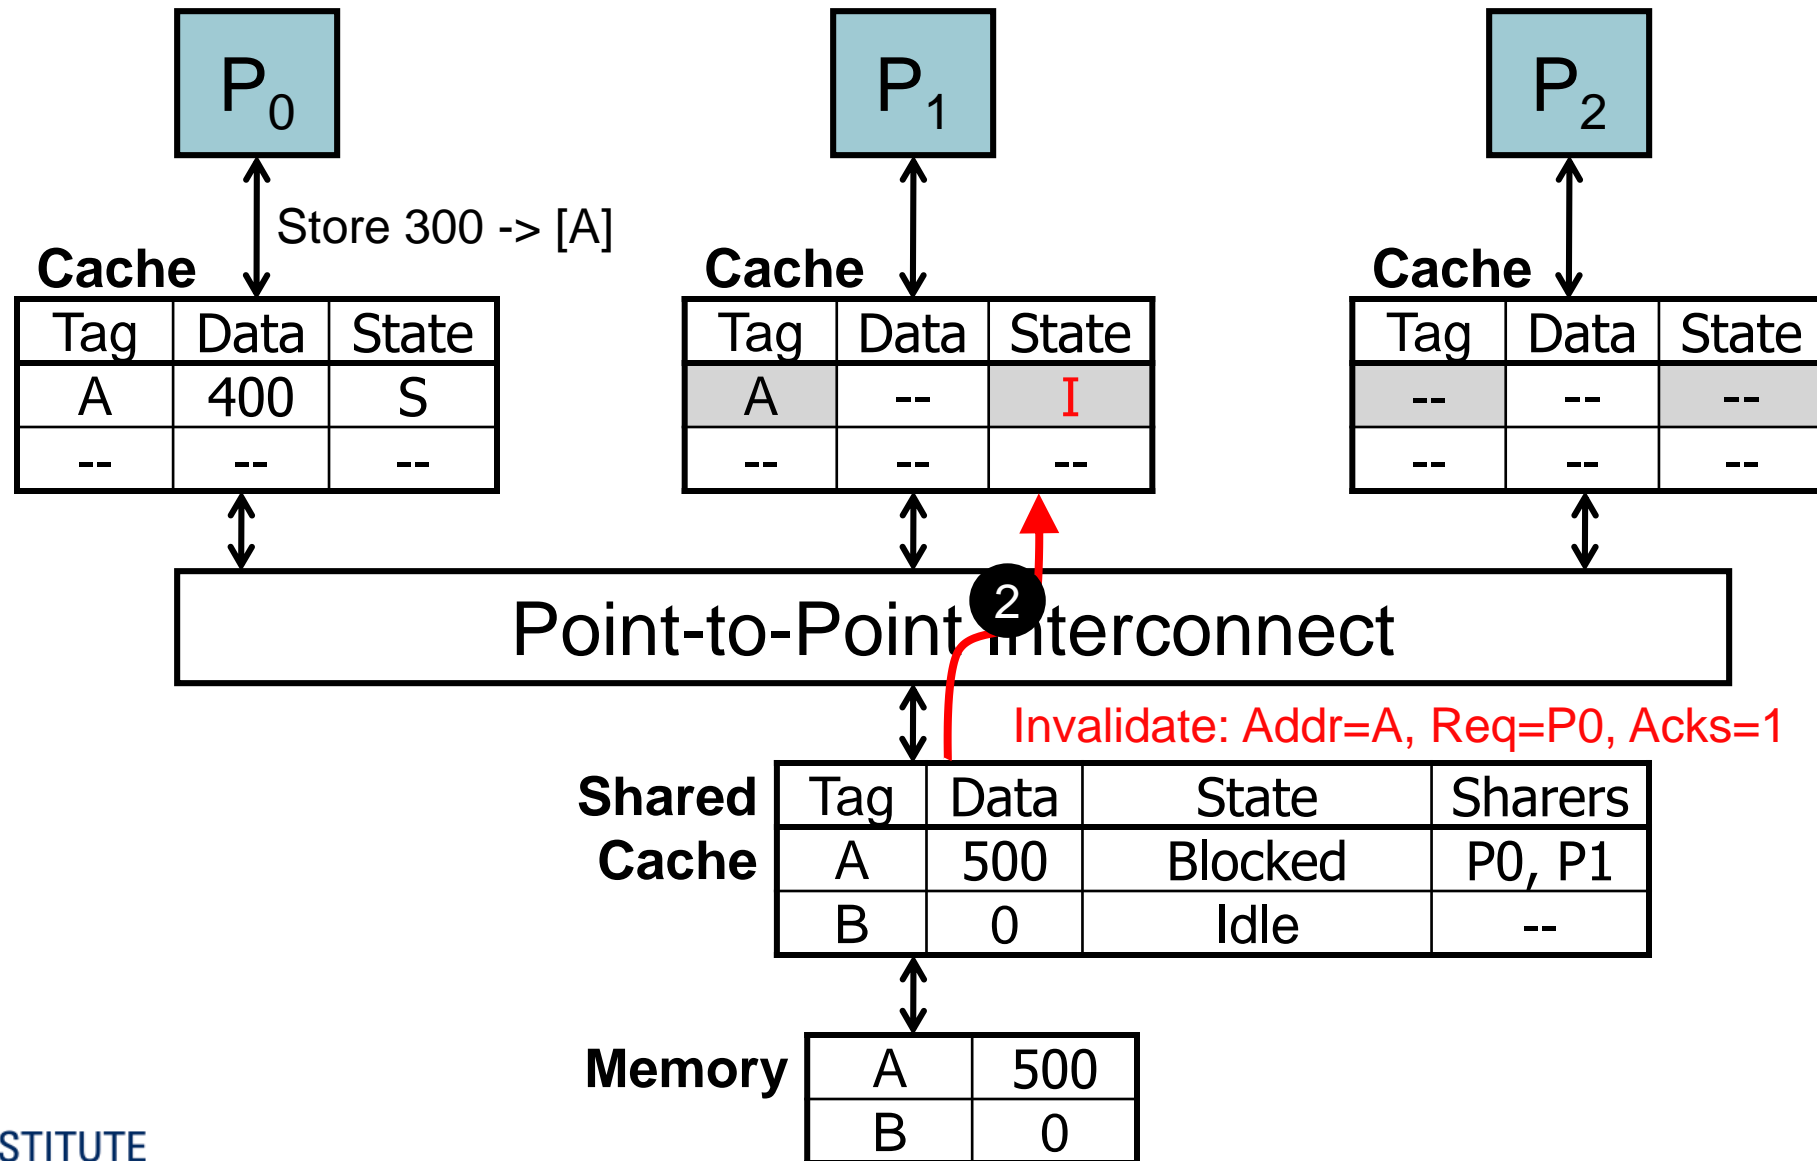
MSI Coherence Example: Step #6



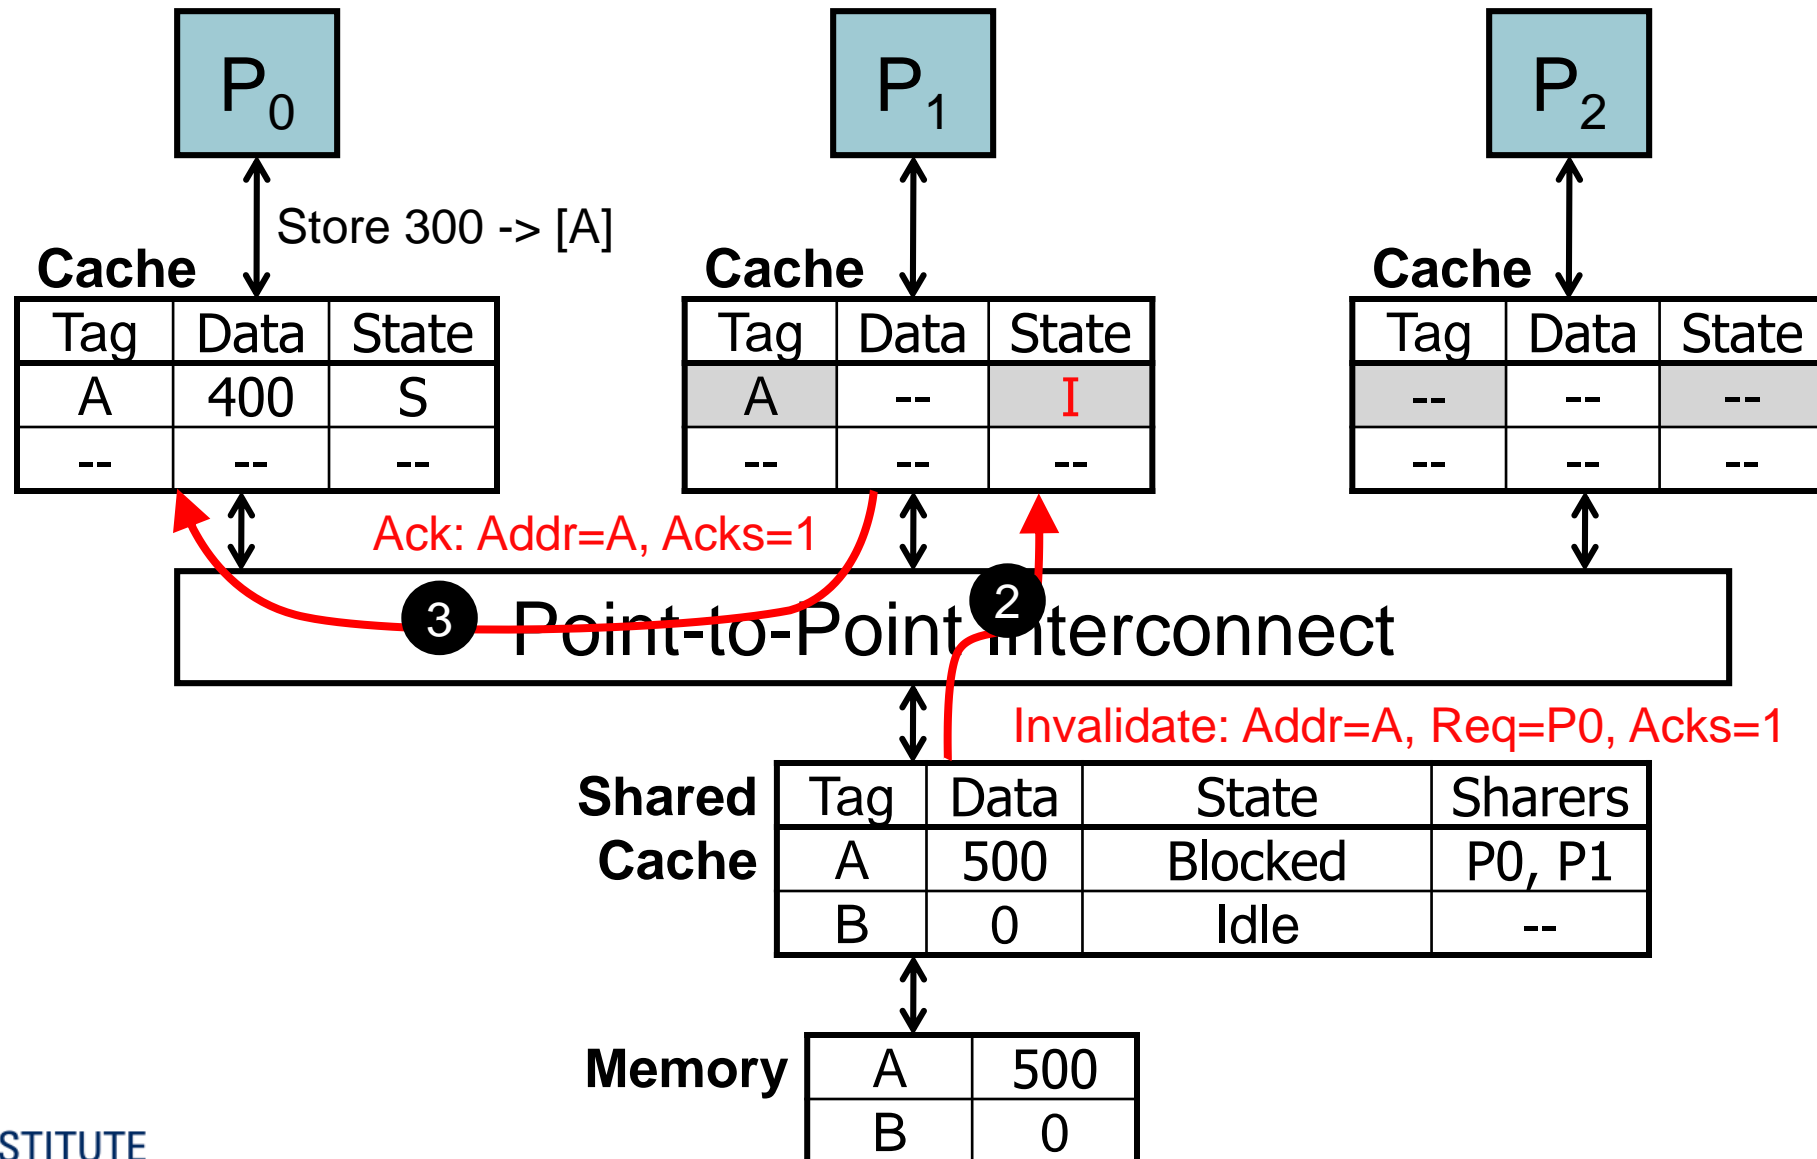
MSI Coherence Example: Step #7



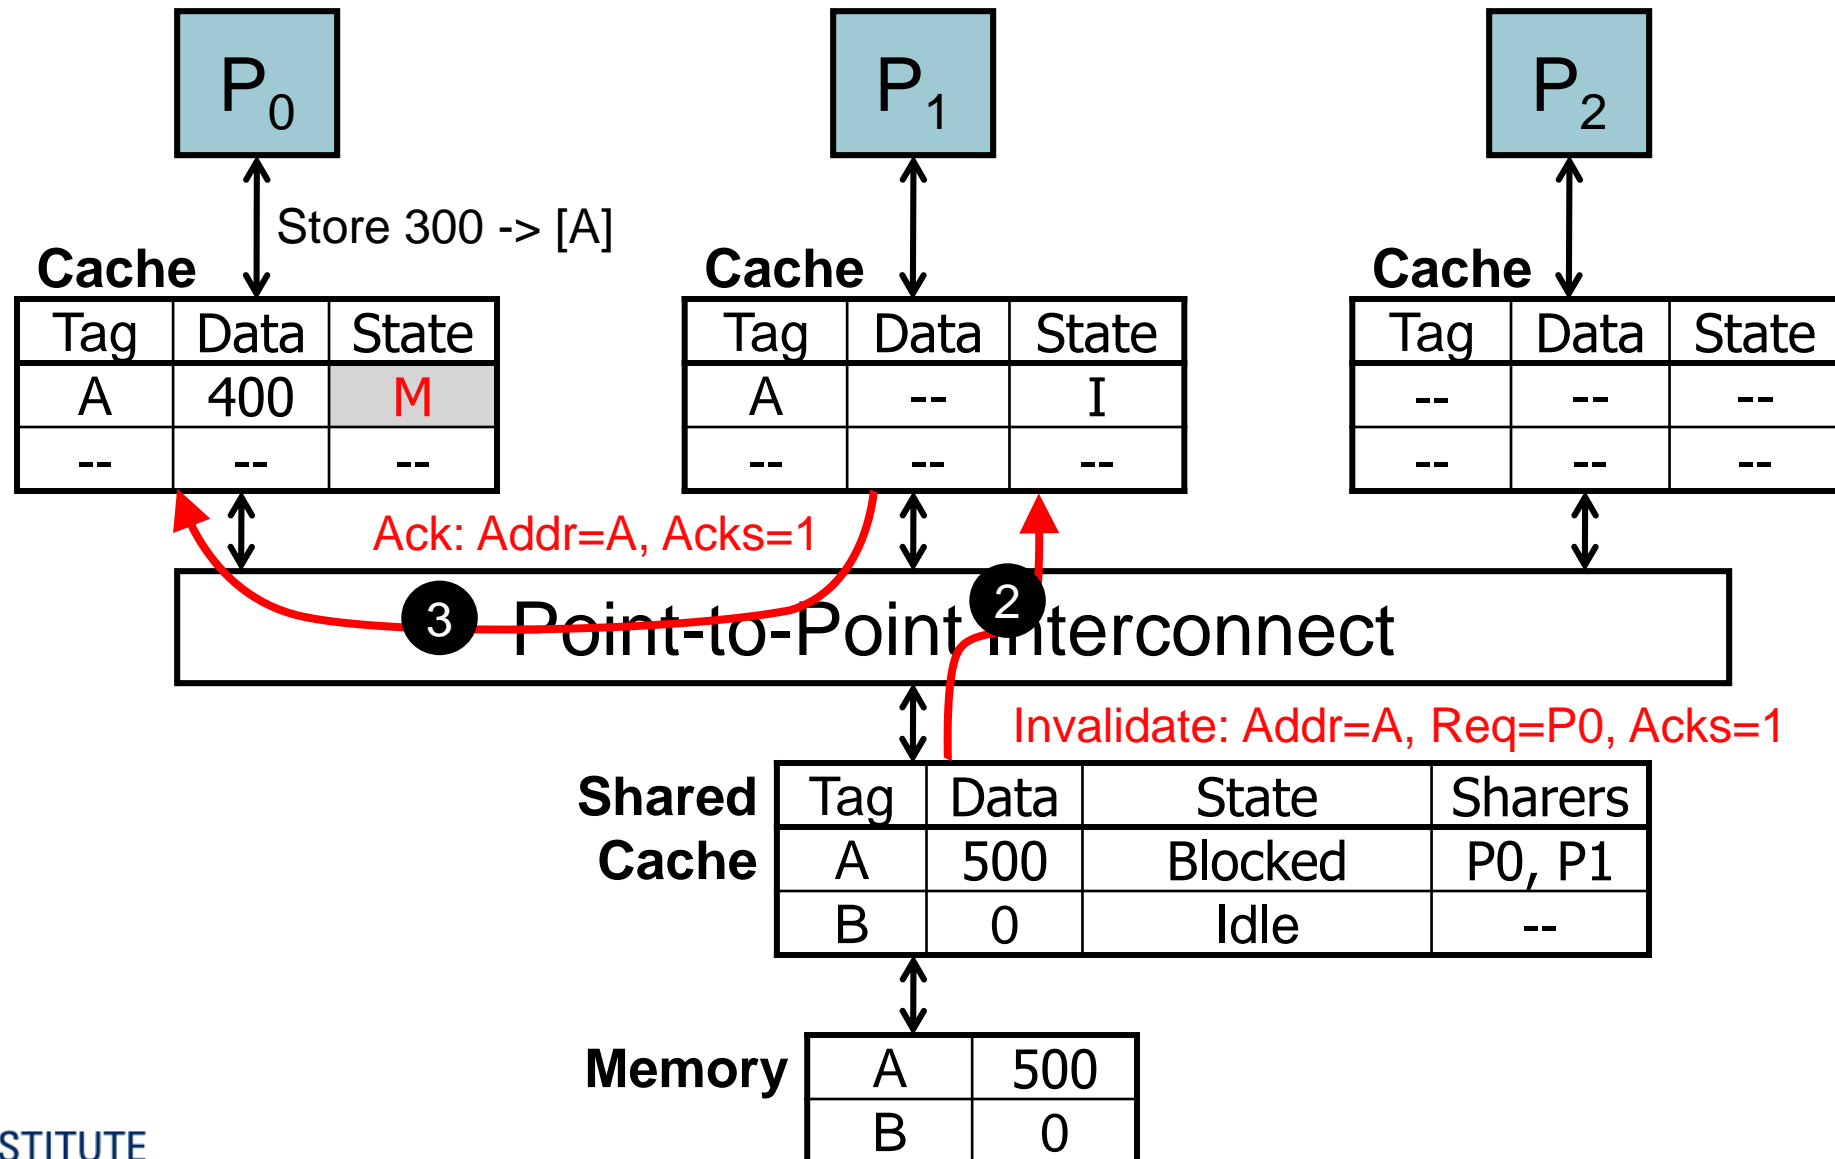
MSI Coherence Example: Step #8



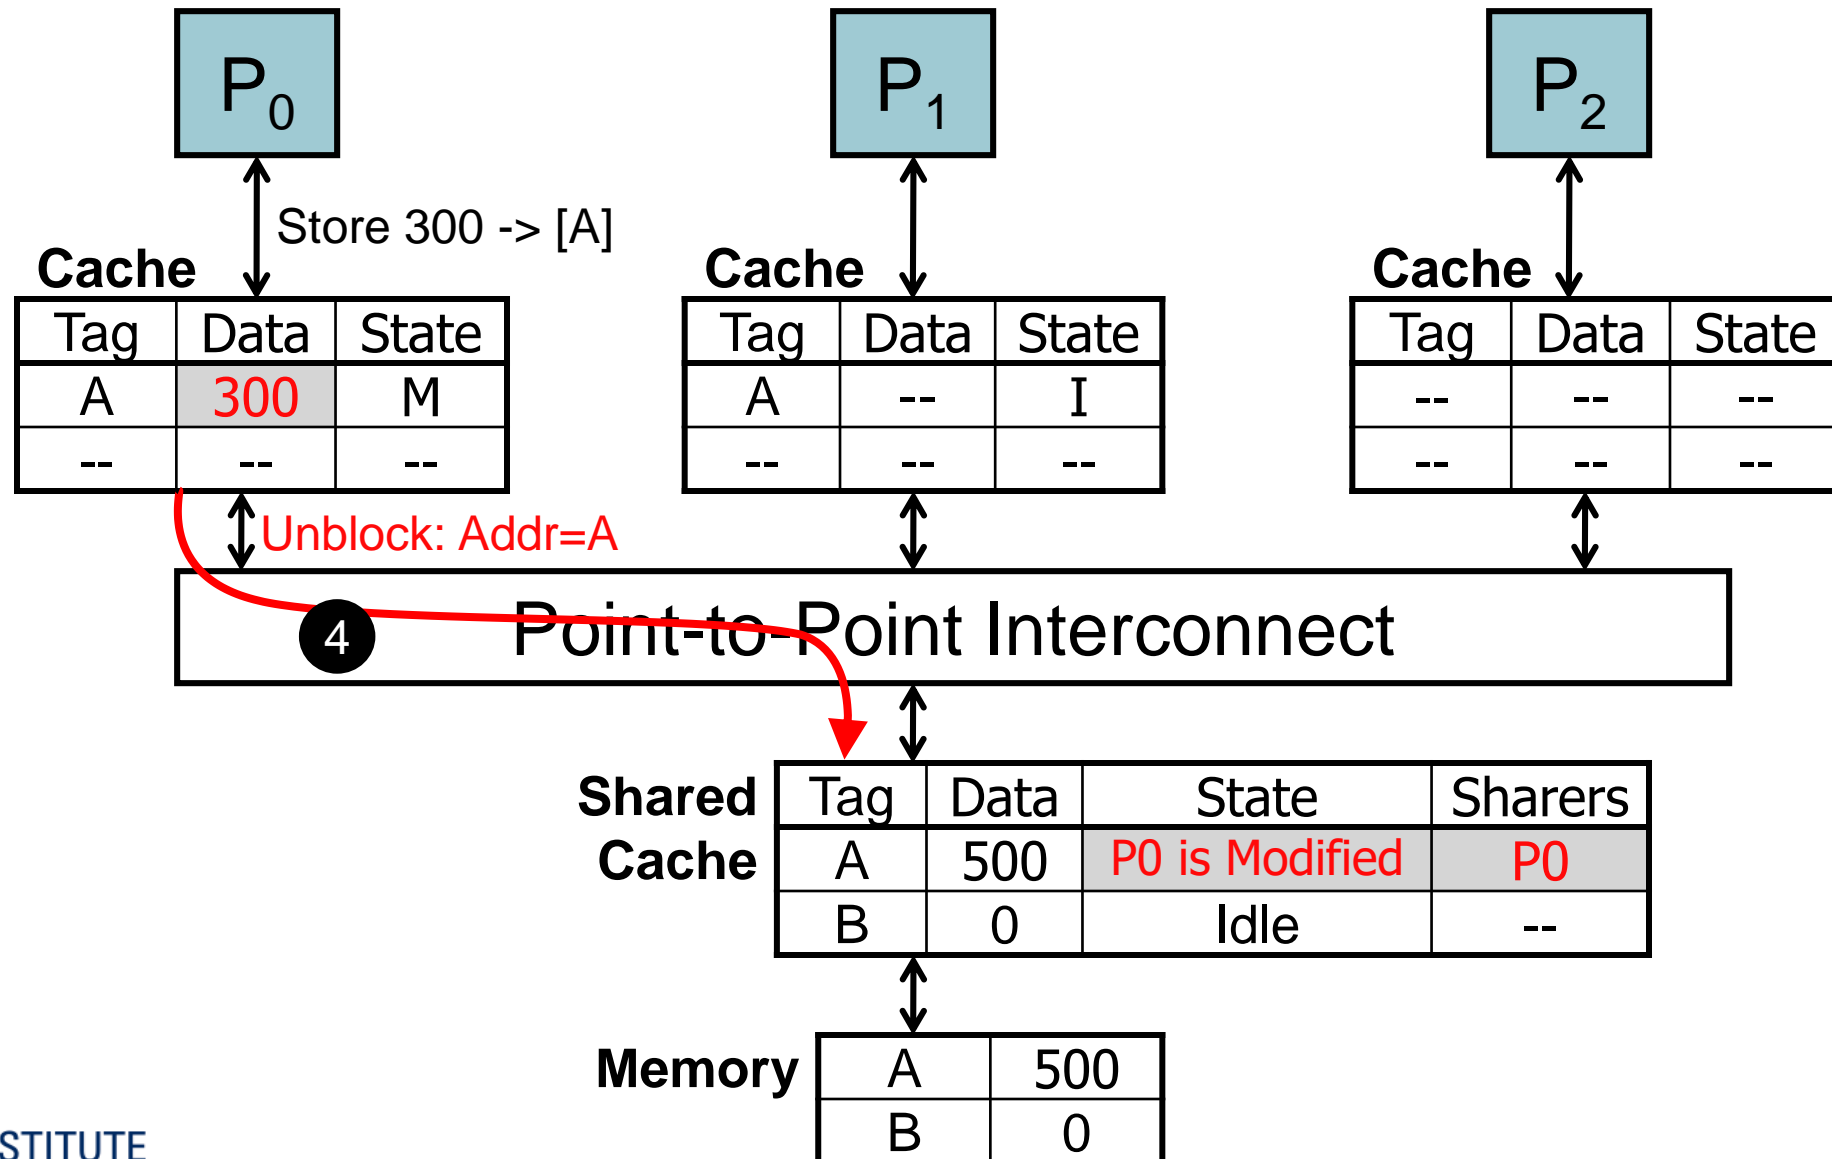
MSI Coherence Example: Step #9



MSI Coherence Example: Step #10

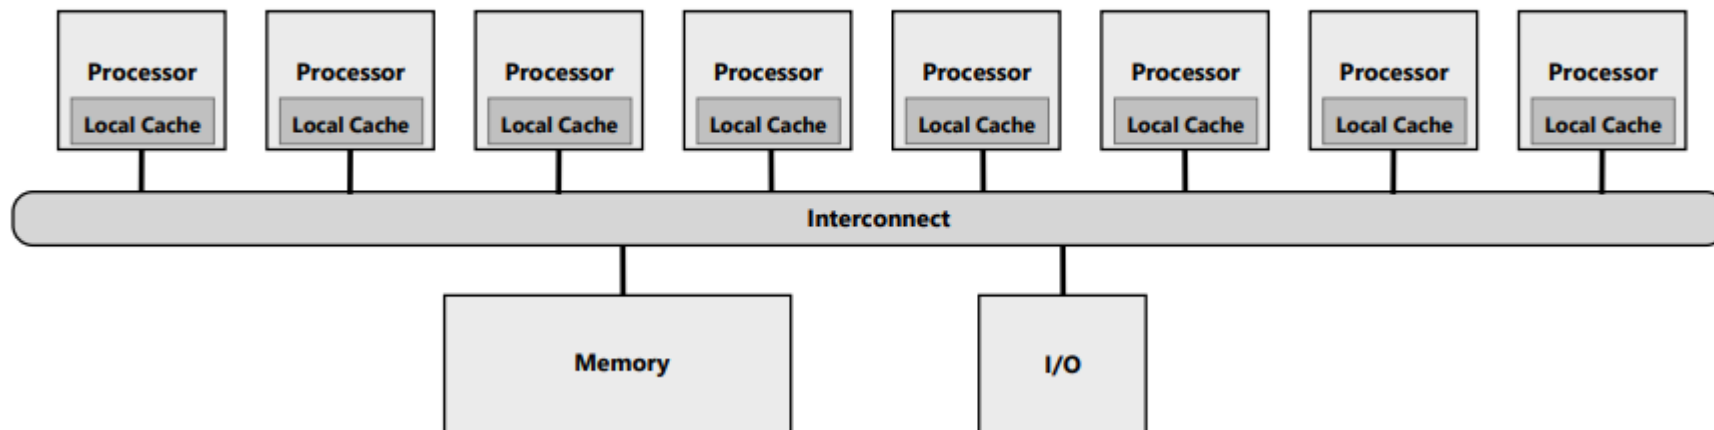


MSI Coherence Example: Step #11

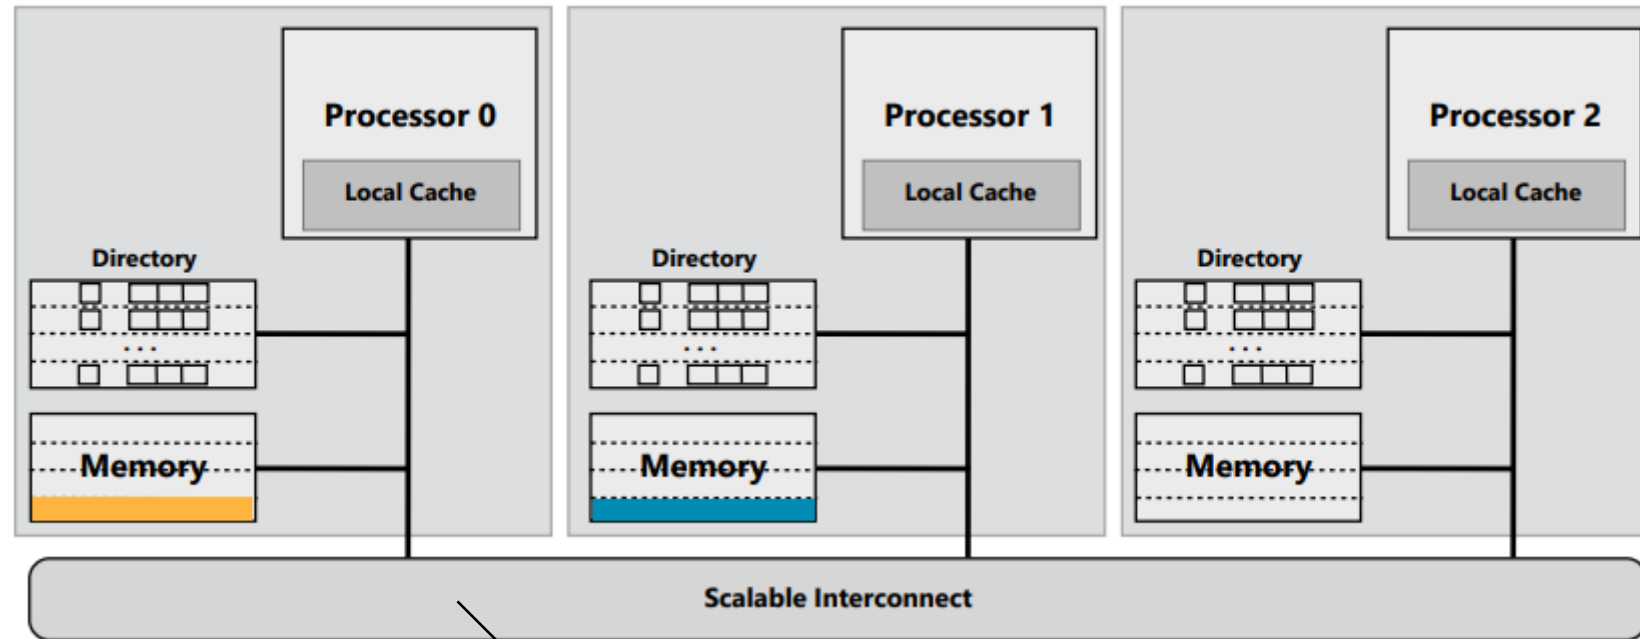


Snooping based protocol

- V-I, MSI, MESI (Read Textbook Page 388)
- Every time a cache miss occurred, the triggering cache communicated with all other caches
- Communication limited to one message at a time
- Scalability is a huge challenge



Directory Based Cache Coherence



Not a bus! A lot of point-to-point connection, can be a ring

Read Ch. 5.4!

Where are we Heading?

- T7: More multicore

Acknowledgement

Slides in this topic are inspired in part by material developed and copyright by:

- ARM Courseware
- Nvidia Courseware
- Prof. Onur Mutlu @ ETH
- Prof. Joe Devietti @ Upenn, CIS 571
- Prof. Hakim Weatherspoon @ Cornell, CS 3410
- Prof. Krste Asanovic @ UCB, CS252
- Xinfei Guo @ JI, VE370 2021 SU

Action Items

- Lab #6 due by 11:59pm July 15th, 2022 (Beijing Time)
- Final project
- Reading Materials
 - Ch. 5.1, 5.2, 5.6